

The Python 3 Tutorial Professional Manual

Compiled automatically.

The Python Tutorial

Tip

This tutorial is designed for *programmers* that are new to the Python language, **not** *beginners* who are new to programming.

Python is an easy to learn, powerful programming language. It has efficient high-level data structures and a simple but effective approach to object-oriented programming. Python's elegant syntax and dynamic typing, together with its interpreted nature, make it an ideal language for scripting and rapid application development in many areas on most platforms.

The Python interpreter and the extensive standard library are freely available in source or binary form for all major platforms from the Python website, <https://www.python.org/>, and may be freely distributed. The same site also contains distributions of and pointers to many free third party Python modules, programs and tools, and additional documentation.

The Python interpreter is easily extended with new functions and data types implemented in C or C++ (or other languages callable from C). Python is also suitable as an extension language for customizable applications.

This tutorial introduces the reader informally to the basic concepts and features of the Python language and system. Be aware that it expects you to have a basic understanding of programming in general. It helps to have a Python interpreter handy for hands-on experience, but all examples are self-contained, so the tutorial can be read off-line as well.

For a description of standard objects and modules, see The Python Standard Library. The Python Language Reference gives a more formal definition of the language. To write extensions in C or C++, read Extending and Embedding the Python Interpreter and Python/C API reference manual. There are also several books covering Python in depth.

This tutorial does not attempt to be comprehensive and cover every single feature, or even every commonly used feature. Instead, it introduces many of Python's most noteworthy features, and will give you a good idea of the language's flavor and style. After reading it, you will be able to read and write Python modules and programs, and you will be ready to learn more about the various Python library modules described in The Python Standard Library.

The Glossary is also worth going through.

- 1. Whetting Your Appetite
- 2. Using the Python Interpreter
 - 2.1. Invoking the Interpreter
 - 2.1.1. Argument Passing
 - 2.1.2. Interactive Mode
 - 2.2. The Interpreter and Its Environment

- 2.2.1. Source Code Encoding
- 3. An Informal Introduction to Python
 - 3.1. Using Python as a Calculator
 - 3.1.1. Numbers
 - 3.1.2. Text
 - 3.1.3. Lists
 - 3.2. First Steps Towards Programming
- 4. More Control Flow Tools
 - 4.1.

`if`

 Statements
 - 4.2.

`for`

 Statements
 - 4.3. The

`range()`

 Function
 - 4.4.

`break`

 and

`continue`

 Statements
 - 4.5.

`else`

 Clauses on Loops
 - 4.6.

`pass`

 Statements
 - 4.7.

`match`

 Statements
 - 4.8. Defining Functions
 - 4.9. More on Defining Functions
 - 4.9.1. Default Argument Values
 - 4.9.2. Keyword Arguments
 - 4.9.3. Special parameters
 - 4.9.3.1. Positional-or-Keyword Arguments

- 4.9.3.2. Positional-Only Parameters
 - 4.9.3.3. Keyword-Only Arguments
 - 4.9.3.4. Function Examples
 - 4.9.3.5. Recap
 - 4.9.4. Arbitrary Argument Lists
 - 4.9.5. Unpacking Argument Lists
 - 4.9.6. Lambda Expressions
 - 4.9.7. Documentation Strings
 - 4.9.8. Function Annotations
- 4.10. Intermezzo: Coding Style
- 5. Data Structures
 - 5.1. More on Lists
 - 5.1.1. Using Lists as Stacks
 - 5.1.2. Using Lists as Queues
 - 5.1.3. List Comprehensions
 - 5.1.4. Nested List Comprehensions
 - 5.2. The

`del`

 statement
 - 5.3. Tuples and Sequences
 - 5.4. Sets
 - 5.5. Dictionaries
 - 5.6. Looping Techniques
 - 5.7. More on Conditions
 - 5.8. Comparing Sequences and Other Types
- 6. Modules
 - 6.1. More on Modules
 - 6.1.1. Executing modules as scripts
 - 6.1.2. The Module Search Path
 - 6.1.3. “Compiled” Python files
 - 6.2. Standard Modules
 - 6.3. The

`dir()`

 Function
 - 6.4. Packages
 - 6.4.1. Importing * From a Package
 - 6.4.2. Intra-package References
 - 6.4.3. Packages in Multiple Directories
- 7. Input and Output
 - 7.1. Fancier Output Formatting
 - 7.1.1. Formatted String Literals
 - 7.1.2. The String format() Method

- 7.1.3. Manual String Formatting
 - 7.1.4. Old string formatting
- 7.2. Reading and Writing Files
 - 7.2.1. Methods of File Objects
 - 7.2.2. Saving structured data with

json
- 8. Errors and Exceptions
 - 8.1. Syntax Errors
 - 8.2. Exceptions
 - 8.3. Handling Exceptions
 - 8.4. Raising Exceptions
 - 8.5. Exception Chaining
 - 8.6. User-defined Exceptions
 - 8.7. Defining Clean-up Actions
 - 8.8. Predefined Clean-up Actions
 - 8.9. Raising and Handling Multiple Unrelated Exceptions
 - 8.10. Enriching Exceptions with Notes
- 9. Classes
 - 9.1. A Word About Names and Objects
 - 9.2. Python Scopes and Namespaces
 - 9.2.1. Scopes and Namespaces Example
 - 9.3. A First Look at Classes
 - 9.3.1. Class Definition Syntax
 - 9.3.2. Class Objects
 - 9.3.3. Instance Objects
 - 9.3.4. Method Objects
 - 9.3.5. Class and Instance Variables
 - 9.4. Random Remarks
 - 9.5. Inheritance
 - 9.5.1. Multiple Inheritance
 - 9.6. Private Variables
 - 9.7. Odds and Ends
 - 9.8. Iterators
 - 9.9. Generators
 - 9.10. Generator Expressions
- 10. Brief Tour of the Standard Library
 - 10.1. Operating System Interface
 - 10.2. File Wildcards
 - 10.3. Command Line Arguments
 - 10.4. Error Output Redirection and Program Termination
 - 10.5. String Pattern Matching
 - 10.6. Mathematics
 - 10.7. Internet Access

- 10.8. Dates and Times
 - 10.9. Data Compression
 - 10.10. Performance Measurement
 - 10.11. Quality Control
 - 10.12. Batteries Included
- 11. Brief Tour of the Standard Library — Part II
 - 11.1. Output Formatting
 - 11.2. Templating
 - 11.3. Working with Binary Data Record Layouts
 - 11.4. Multi-threading
 - 11.5. Logging
 - 11.6. Weak References
 - 11.7. Tools for Working with Lists
 - 11.8. Decimal Floating-Point Arithmetic
- 12. Virtual Environments and Packages
 - 12.1. Introduction
 - 12.2. Creating Virtual Environments
 - 12.3. Managing Packages with pip
- 13. What Now?
- 14. Interactive Input Editing and History Substitution
 - 14.1. Tab Completion and History Editing
 - 14.2. Alternatives to the Interactive Interpreter
- 15. Floating-Point Arithmetic: Issues and Limitations
 - 15.1. Representation Error
- 16. Appendix
 - 16.1. Interactive Mode
 - 16.1.1. Error Handling
 - 16.1.2. Executable Python Scripts
 - 16.1.3. The Interactive Startup File
 - 16.1.4. The Customization Modules

1. Whetting Your Appetite¶

If you do much work on computers, eventually you find that there's some task you'd like to automate. For example, you may wish to perform a search-and-replace over a large number of text files, or rename and rearrange a bunch of photo files in a complicated way. Perhaps you'd like to write a small custom database, or a specialized GUI application, or a simple game.

If you're a professional software developer, you may have to work with several C/C++/Java libraries but find the usual write/compile/test/re-compile cycle is too slow. Perhaps you're writing a test suite for such a library and find writing the testing code a tedious task. Or maybe you've written a program that could use an extension language, and you don't want to design and implement a whole new language for your application.

Python is just the language for you.

You could write a Unix shell script or Windows batch files for some of these tasks, but shell scripts are best at moving around files and changing text data, not well-suited for GUI applications or games. You could write a C/C++/Java program, but it can take a lot of development time to get even a first-draft program. Python is simpler to use, available on Windows, macOS, and Unix operating systems, and will help you get the job done more quickly.

Python is simple to use, but it is a real programming language, offering much more structure and support for large programs than shell scripts or batch files can offer. On the other hand, Python also offers much more error checking than C, and, being a *very-high-level language*, it has high-level data types built in, such as flexible arrays and dictionaries. Because of its more general data types Python is applicable to a much larger problem domain than Awk or even Perl, yet many things are at least as easy in Python as in those languages.

Python allows you to split your program into modules that can be reused in other Python programs. It comes with a large collection of standard modules that you can use as the basis of your programs — or as examples to start learning to program in Python. Some of these modules provide things like file I/O, system calls, sockets, and even interfaces to graphical user interface toolkits like Tk.

Python is an interpreted language, which can save you considerable time during program development because no compilation and linking is necessary. The interpreter can be used interactively, which makes it easy to experiment with features of the language, to write throw-away programs, or to test functions during bottom-up program development. It is also a handy desk calculator.

Python enables programs to be written compactly and readably. Programs written in Python are typically much shorter than equivalent C, C++, or Java programs, for several reasons:

- the high-level data types allow you to express complex operations in a single statement;
- statement grouping is done by indentation instead of beginning and ending brackets;

- no variable or argument declarations are necessary.

Python is *extensible*: if you know how to program in C it is easy to add a new built-in function or module to the interpreter, either to perform critical operations at maximum speed, or to link Python programs to libraries that may only be available in binary form (such as a vendor-specific graphics library). Once you are really hooked, you can link the Python interpreter into an application written in C and use it as an extension or command language for that application.

By the way, the language is named after the BBC show “Monty Python’s Flying Circus” and has nothing to do with reptiles. Making references to Monty Python skits in documentation is not only allowed, it is encouraged!

Now that you are all excited about Python, you’ll want to examine it in some more detail. Since the best way to learn a language is to use it, the tutorial invites you to play with the Python interpreter as you read.

In the next chapter, the mechanics of using the interpreter are explained. This is rather mundane information, but essential for trying out the examples shown later.

The rest of the tutorial introduces various features of the Python language and system through examples, beginning with simple expressions, statements and data types, through functions and modules, and finally touching upon advanced concepts like exceptions and user-defined classes.

2. Using the Python Interpreter¶

2.1. Invoking the Interpreter¶

The Python interpreter is usually installed as

```
/usr/local/bin/python3.14
```

on those machines where it is available; putting

```
/usr/local/bin
```

in your Unix shell's search path makes it possible to start it by typing the command:

```
python3.14
```

to the shell. [1] Since the choice of the directory where the interpreter lives is an installation option, other places are possible; check with your local Python guru or system administrator. (E.g.,

```
/usr/local/python
```

is a popular alternative location.)

On Windows machines where you have installed Python from the Microsoft Store, the

```
python3.14
```

command will be available. If you have the py.exe launcher installed, you can use the

```
py
```

command. See Python install manager for other ways to launch Python.

Typing an end-of-file character (Control-D on Unix, Control-Z on Windows) at the primary prompt causes the interpreter to exit with a zero exit status. If that doesn't work, you can exit the interpreter by typing the following command:

```
quit()
```

.

The interpreter's line-editing features include interactive editing, history substitution and code completion on most systems. Perhaps the quickest check to see whether command line editing is supported is typing a word in on the Python prompt, then pressing Left arrow (or Control-b). If the cursor moves, you have command line editing; see Appendix Interactive Input Editing and History Substitution for an introduction to the keys. If nothing appears to happen, or if a sequence like

```
^[[D
```

or

```
^B
```

appears, command line editing isn't available; you'll only be able to use backspace to remove characters from the current line.

The interpreter operates somewhat like the Unix shell: when called with standard input connected to a tty device, it reads and executes commands interactively; when called with a file name argument or with a file as standard input, it reads and executes a *script* from that file.

A second way of starting the interpreter is

```
python -c command [arg] ...
```

, which executes the statement(s) in *command*, analogous to the shell's

```
-c
```

option. Since Python statements often contain spaces or other characters that are special to the shell, it is usually advised to quote *command* in its entirety.

Some Python modules are also useful as scripts. These can be invoked using

```
python -m module [arg] ...
```

, which executes the source file for *module* as if you had spelled out its full name on the command line.

When a script file is used, it is sometimes useful to be able to run the script and enter interactive mode afterwards. This can be done by passing

```
-i
```

before the script.

All command line options are described in Command line and environment.

2.1.1. Argument Passing¶

When known to the interpreter, the script name and additional arguments thereafter are turned into a list of strings and assigned to the

```
argv
```

variable in the

```
sys
```

module. You can access this list by executing

```
import  
sys
```

. The length of the list is at least one; when no script and no arguments are given,

```
sys.argv[0]
```

is an empty string. When the script name is given as

```
'_ '
```

(meaning standard input),

```
sys.argv[0]
```

is set to

```
'_ '
```

. When

```
-c
```

command is used,

```
sys.argv[0]
```

is set to

```
'-c '
```

. When

```
-m
```

module is used,

```
sys.argv[0]
```

is set to the full name of the located module. Options found after

```
-c
```

command or

```
-m
```

module are not consumed by the Python interpreter's option processing but left in

```
sys.argv
```

for the command or module to handle.

2.1.2. Interactive Mode¶

When commands are read from a tty, the interpreter is said to be in *interactive mode*. In this mode it prompts for the next command with the *primary prompt*, usually three greater-than signs (

```
>>>
```

); for continuation lines it prompts with the *secondary prompt*, by default three dots (

...

). The interpreter prints a welcome message stating its version number and a copyright notice before printing the first prompt:

Copy

```
$ python3.14
Python 3.14 (default, April 4 2024, 09:25:04)
[GCC 10.2.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
>>>
```

Continuation lines are needed when entering a multi-line construct. As an example, take a look at this

```
if
```

statement:

Copy

```
>>> the_world_is_flat = True
>>> if the_world_is_flat:
...     print("Be careful not to fall off!")
...
Be careful not to fall off!
```

For more on interactive mode, see [Interactive Mode](#).

2.2. The Interpreter and Its Environment¶

2.2.1. Source Code Encoding¶

By default, Python source files are treated as encoded in UTF-8. In that encoding, characters of most languages in the world can be used simultaneously in string literals, identifiers and comments — although the standard library only uses ASCII characters for identifiers, a convention that any portable code should follow. To display all these characters properly, your editor must recognize that the file is UTF-8, and it must use a font that supports all the characters in the file.

To declare an encoding other than the default one, a special comment line should be added as the *first* line of the file. The syntax is as follows:

Copy

```
# -*- coding: encoding -*-
```

where *encoding* is one of the valid

```
codecs
```

supported by Python.

For example, to declare that Windows-1252 encoding is to be used, the first line of your source code file should be:

Copy

```
# -*- coding: cp1252 -*-
```

One exception to the *first line* rule is when the source code starts with a UNIX “shebang” line. In this case, the encoding declaration should be added as the second line of the file. For example:

Copy

```
#!/usr/bin/env python3  
# -*- coding: cp1252 -*-
```

Footnotes

[\[1\]](#)

On Unix, the Python 3.x interpreter is by default not installed with the executable named

```
python
```

, so that it does not conflict with a simultaneously installed Python 2.x executable.

3. An Informal Introduction to Python¶

In the following examples, input and output are distinguished by the presence or absence of prompts (`>>>` and `...`): to repeat the example, you must type everything after the prompt, when the prompt appears; lines that do not begin with a prompt are output from the interpreter. Note that a secondary prompt on a line by itself in an example means you must type a blank line; this is used to end a multi-line command.

You can use the “Copy” button (it appears in the upper-right corner when hovering over or tapping a code example), which strips prompts and omits output, to copy and paste the input lines into your interpreter.

Many of the examples in this manual, even those entered at the interactive prompt, include comments. Comments in Python start with the hash character,

```
#
```

, and extend to the end of the physical line. A comment may appear at the start of a line or following whitespace or code, but not within a string literal. A hash character within a string literal is just a hash character. Since comments are to clarify code and are not interpreted by Python, they may be omitted when typing in examples.

Some examples:

Copy

```
# this is the first comment
spam = 1 # and this is the second comment
        # ... and now a third!
text = "# This is not a comment because it's inside quotes."
```

3.1. Using Python as a Calculator¶

Let’s try some simple Python commands. Start the interpreter and wait for the primary prompt,

```
>>>
```

. (It shouldn’t take long.)

3.1.1. Numbers¶

The interpreter acts as a simple calculator: you can type an expression into it and it will write the value. Expression syntax is straightforward: the operators

```
+
```

,

-

,

*

and

/

can be used to perform arithmetic; parentheses (

()

) can be used for grouping. For example:

Copy

```
>>> 2 + 2
4
>>> 50 - 5*6
20
>>> (50 - 5*6) / 4
5.0
>>> 8 / 5 # division always returns a floating-point number
1.6
```

The integer numbers (e.g.

2

,

4

,

20

) have type

int

, the ones with a fractional part (e.g.

5.0

,

1.6

) have type

```
float
```

. We will see more about numeric types later in the tutorial.

Division (

```
/
```

) always returns a float. To do floor division and get an integer result you can use the

```
//
```

operator; to calculate the remainder you can use

```
%
```

:

Copy

```
>>> 17 / 3 # classic division returns a float
5.666666666666667
>>>
>>> 17 // 3 # floor division discards the fractional part
5
>>> 17 % 3 # the % operator returns the remainder of the division
2
>>> 5 * 3 + 2 # floored quotient * divisor + remainder
17
```

With Python, it is possible to use the

```
**
```

operator to calculate powers [\[1\]](#):

Copy

```
>>> 5 ** 2 # 5 squared
25
>>> 2 ** 7 # 2 to the power of 7
128
```

The equal sign (

```
=
```

) is used to assign a value to a variable. Afterwards, no result is displayed before the next interactive prompt:

Copy

```
>>> width = 20
>>> height = 5 * 9
>>> width * height
900
```

If a variable is not “defined” (assigned a value), trying to use it will give you an error:

Copy

```
>>> n # try to access an undefined variable
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
NameError: name 'n' is not defined
```

There is full support for floating point; operators with mixed type operands convert the integer operand to floating point:

Copy

```
>>> 4 * 3.75 - 1
14.0
```

In interactive mode, the last printed expression is assigned to the variable

—

. This means that when you are using Python as a desk calculator, it is somewhat easier to continue calculations, for example:

Copy

```
>>> tax = 12.5 / 100
>>> price = 100.50
>>> price * tax
12.5625
>>> price + _
113.0625
>>> round(_, 2)
113.06
```

This variable should be treated as read-only by the user. Don’t explicitly assign a value to it — you would create an independent local variable with the same name masking the built-in variable with its magic behavior.

In addition to

```
int
```

and

```
float
```

, Python supports other types of numbers, such as

```
Decimal
```

and

```
Fraction
```

. Python also has built-in support for complex numbers, and uses the

```
j
```

or

```
J
```

suffix to indicate the imaginary part (e.g.

```
3+5j
```

).

3.1.2. Text

Python can manipulate text (represented by type

```
str
```

, so-called “strings”) as well as numbers. This includes characters “

```
!
```

”, words “

```
rabbit
```

”, names “

```
Paris
```

”, sentences “

```
Got your back.
```

”, etc. “

```
Yay! :)
```

”. They can be enclosed in single quotes (

```
'...'
```

) or double quotes (

```
"..."
```

) with the same result [\[2\]](#).

Copy

```
>>> 'spam eggs' # single quotes
'spam eggs'
>>> "Paris rabbit got your back :)! Yay!" # double quotes
'Paris rabbit got your back :)! Yay!'
>>> '1975' # digits and numerals enclosed in quotes are also strings
'1975'
```

To quote a quote, we need to “escape” it, by preceding it with

```
\
```

. Alternatively, we can use the other type of quotation marks:

Copy

```
>>> 'doesn\'t' # use \' to escape the single quote...
"doesn't"
>>> "doesn't" # ...or use double quotes instead
"doesn't"
>>> '"Yes," they said.'
'"Yes," they said.'
>>> "\"Yes,\" they said."
'"Yes," they said.'
>>> '"Isn\'t," they said.'
'"Isn\'t," they said.'
```

In the Python shell, the string definition and output string can look different. The

```
print()
```

function produces a more readable output, by omitting the enclosing quotes and by printing escaped and special characters:

Copy

```
>>> s = 'First line.\nSecond line.' # \n means newline
>>> s # without print(), special characters are included in the string
'First line.\nSecond line.'
>>> print(s) # with print(), special characters are interpreted, so \n produces new line
First line.
Second line.
```

If you don't want characters prefaced by

```
\
```

to be interpreted as special characters, you can use *raw strings* by adding an

```
r
```

before the first quote:

Copy

```
>>> print('C:\some\name') # here \n means newline!
C:\some
ame
>>> print(r'C:\some\name') # note the r before the quote
C:\some\name
```

There is one subtle aspect to raw strings: a raw string may not end in an odd number of

```
\
```

characters; see the FAQ entry for more information and workarounds.

String literals can span multiple lines. One way is using triple-quotes:

```
"""..."""
```

or

```
'''...'''
```

. End-of-line characters are automatically included in the string, but it's possible to prevent this by adding a

```
\
```

at the end of the line. In the following example, the initial newline is not included:

Copy

```
>>> print("""\
... Usage: thingy [OPTIONS]
...     -h                        Display this usage message
...     -H hostname                Hostname to connect to
... """)
Usage: thingy [OPTIONS]
      -h                        Display this usage message
      -H hostname                Hostname to connect to

>>>
```

Strings can be concatenated (glued together) with the

```
+
```

operator, and repeated with

```
*
```

```
:
```

Copy

```
>>> # 3 times 'un', followed by 'ium'
>>> 3 * 'un' + 'ium'
'ununinium'
```

Two or more *string literals* (i.e. the ones enclosed between quotes) next to each other are automatically concatenated.

Copy

```
>>> 'Py' 'thon'
'Python'
```

This feature is particularly useful when you want to break long strings:

Copy

```
>>> text = ('Put several strings within parentheses '
...         'to have them joined together.')
>>> text
'Put several strings within parentheses to have them joined together.'
```

This only works with two literals though, not with variables or expressions:

Copy

```
>>> prefix = 'Py'
>>> prefix 'thon' # can't concatenate a variable and a string literal
File "<stdin>", line 1
    prefix 'thon'
        ^^^^^^
SyntaxError: invalid syntax
>>> ('un' * 3) 'ium'
File "<stdin>", line 1
    ('un' * 3) 'ium'
        ^^^^^
SyntaxError: invalid syntax
```

If you want to concatenate variables or a variable and a literal, use

+

:

Copy

```
>>> prefix + 'thon'
'Python'
```

Strings can be *indexed* (subscripted), with the first character having index 0. There is no separate character type; a character is simply a string of size one:

Copy

```
>>> word = 'Python'
>>> word[0] # character in position 0
'p'
>>> word[5] # character in position 5
'n'
```

Indices may also be negative numbers, to start counting from the right:

Copy

```
>>> word[-1] # last character
'n'
>>> word[-2] # second-last character
'o'
>>> word[-6]
'p'
```

Note that since -0 is the same as 0, negative indices start from -1.

In addition to indexing, *slicing* is also supported. While indexing is used to obtain individual characters, *slicing* allows you to obtain a substring:

Copy

```
>>> word[0:2] # characters from position 0 (included) to 2 (excluded)
'Py'
>>> word[2:5] # characters from position 2 (included) to 5 (excluded)
'tho'
```

Slice indices have useful defaults; an omitted first index defaults to zero, an omitted second index defaults to the size of the string being sliced.

Copy

```
>>> word[:2] # character from the beginning to position 2 (excluded)
'Py'
>>> word[4:] # characters from position 4 (included) to the end
'on'
>>> word[-2:] # characters from the second-last (included) to the end
'on'
```

Note how the start is always included, and the end always excluded. This makes sure that

```
s[:i] + s[i:]
```

is always equal to

```
s
```

:

Copy

```
>>> word[:2] + word[2:]
'Python'
>>> word[:4] + word[4:]
'Python'
```

One way to remember how slices work is to think of the indices as pointing *between* characters, with the left edge of the first character numbered 0. Then the right edge of the last character of a string of n characters has index n , for example:

Copy

```
+---+---+---+---+---+---+
| P | y | t | h | o | n |
+---+---+---+---+---+---+
0   1   2   3   4   5   6
-6  -5  -4  -3  -2  -1
```

The first row of numbers gives the position of the indices 0...6 in the string; the second row gives the corresponding negative indices. The slice from i to j consists of all characters between the edges labeled i and j , respectively.

For non-negative indices, the length of a slice is the difference of the indices, if both are within bounds. For example, the length of

```
word[1:3]
```

is 2.

Attempting to use an index that is too large will result in an error:

Copy

```
>>> word[42] # the word only has 6 characters
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
IndexError: string index out of range
```

However, out of range slice indexes are handled gracefully when used for slicing:

Copy

```
>>> word[4:42]
'on'
>>> word[42:]
''
```


Python strings cannot be changed — they are immutable. Therefore, assigning to an indexed position in the string results in an error:

Copy

```
>>> word[0] = 'J'
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: 'str' object does not support item assignment
>>> word[2:] = 'py'
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: 'str' object does not support item assignment
```

If you need a different string, you should create a new one:

Copy

```
>>> 'J' + word[1:]
'Jython'
>>> word[:2] + 'py'
'Pypy'
```

The built-in function

```
len()
```

returns the length of a string:

Copy

```
>>> s = 'supercalifragilisticexpialidocious'
>>> len(s)
34
```

See also

Text Sequence Type — str

Strings are examples of *sequence types*, and support the common operations supported by such types.

String Methods

Strings support a large number of methods for basic transformations and searching.

f-strings

String literals that have embedded expressions.

Format String Syntax

Information about string formatting with

```
str.format()
```

.

printf-style String Formatting

The old formatting operations invoked when strings are the left operand of the

```
%
```

operator are described in more detail here.

3.1.3. Lists¶

Python knows a number of *compound* data types, used to group together other values. The most versatile is the *list*, which can be written as a list of comma-separated values (items) between square brackets. Lists might contain items of different types, but usually the items all have the same type.

Copy

```
>>> squares = [1, 4, 9, 16, 25]
>>> squares
[1, 4, 9, 16, 25]
```

Like strings (and all other built-in sequence types), lists can be indexed and sliced:

Copy

```
>>> squares[0] # indexing returns the item
1
>>> squares[-1]
25
>>> squares[-3:] # slicing returns a new list
[9, 16, 25]
```

Lists also support operations like concatenation:

Copy

```
>>> squares + [36, 49, 64, 81, 100]
[1, 4, 9, 16, 25, 36, 49, 64, 81, 100]
```

Unlike strings, which are immutable, lists are a mutable type, i.e. it is possible to change their content:

Copy

```
>>> cubes = [1, 8, 27, 65, 125] # something's wrong here
>>> 4 ** 3 # the cube of 4 is 64, not 65!
64
>>> cubes[3] = 64 # replace the wrong value
>>> cubes
[1, 8, 27, 64, 125]
```

You can also add new items at the end of the list, by using the

```
list.append()
```

method (we will see more about methods later):

Copy

```
>>> cubes.append(216) # add the cube of 6
>>> cubes.append(7 ** 3) # and the cube of 7
>>> cubes
[1, 8, 27, 64, 125, 216, 343]
```

Simple assignment in Python never copies data. When you assign a list to a variable, the variable refers to the *existing list*. Any changes you make to the list through one variable will be seen through all other variables that refer to it.:

Copy

```
>>> rgb = ["Red", "Green", "Blue"]
>>> rgba = rgb
>>> id(rgb) == id(rgba) # they reference the same object
True
>>> rgba.append("Alph")
>>> rgb
["Red", "Green", "Blue", "Alph"]
```

All slice operations return a new list containing the requested elements. This means that the following slice returns a shallow copy of the list:

Copy

```
>>> correct_rgba = rgba[:]
>>> correct_rgba[-1] = "Alpha"
>>> correct_rgba
["Red", "Green", "Blue", "Alpha"]
>>> rgba
["Red", "Green", "Blue", "Alph"]
```

Assignment to slices is also possible, and this can even change the size of the list or clear it entirely:

Copy

```
>>> letters = ['a', 'b', 'c', 'd', 'e', 'f', 'g']
>>> letters
['a', 'b', 'c', 'd', 'e', 'f', 'g']
>>> # replace some values
>>> letters[2:5] = ['C', 'D', 'E']
>>> letters
['a', 'b', 'C', 'D', 'E', 'f', 'g']
>>> # now remove them
>>> letters[2:5] = []
>>> letters
['a', 'b', 'f', 'g']
>>> # clear the list by replacing all the elements with an empty list
>>> letters[:] = []
>>> letters
[]
```

The built-in function

`len()`

also applies to lists:

Copy

```
>>> letters = ['a', 'b', 'c', 'd']
>>> len(letters)
4
```

It is possible to nest lists (create lists containing other lists), for example:

Copy

```
>>> a = ['a', 'b', 'c']
>>> n = [1, 2, 3]
>>> x = [a, n]
>>> x
[['a', 'b', 'c'], [1, 2, 3]]
>>> x[0]
['a', 'b', 'c']
>>> x[0][1]
'b'
```

3.2. First Steps Towards Programming¶

Of course, we can use Python for more complicated tasks than adding two and two together. For instance, we can write an initial sub-sequence of the [Fibonacci series](#) as follows:

Copy

```
>>> # Fibonacci series:
>>> # the sum of two elements defines the next
>>> a, b = 0, 1
>>> while a < 10:
...     print(a)
...     a, b = b, a+b
...
0
1
1
2
3
5
8
```

This example introduces several new features.

- The first line contains a *multiple assignment*: the variables

a

and

b

simultaneously get the new values 0 and 1. On the last line this is used again, demonstrating that the expressions on the right-hand side are all evaluated first before any of the

assignments take place. The right-hand side expressions are evaluated from the left to the right.

- The

```
while
```

loop executes as long as the condition (here:

```
a < 10
```

) remains true. In Python, like in C, any non-zero integer value is true; zero is false. The condition may also be a string or list value, in fact any sequence; anything with a non-zero length is true, empty sequences are false. The test used in the example is a simple comparison. The standard comparison operators are written the same as in C:

```
<
```

(less than),

```
>
```

(greater than),

```
==
```

(equal to),

```
<=
```

(less than or equal to),

```
>=
```

(greater than or equal to) and

```
!=
```

(not equal to).

- The *body* of the loop is *indented*: indentation is Python's way of grouping statements. At the interactive prompt, you have to type a tab or space(s) for each indented line. In practice you will prepare more complicated input for Python with a text editor; all decent text editors have an auto-indent facility. When a compound statement is entered interactively, it must be followed by a blank line to indicate completion (since the parser cannot guess when you have typed the last line). Note that each line within a basic block must be indented by the same amount.
- The

```
print()
```

function writes the value of the argument(s) it is given. It differs from just writing the expression you want to write (as we did earlier in the calculator examples) in the way it

handles multiple arguments, floating-point quantities, and strings. Strings are printed without quotes, and a space is inserted between items, so you can format things nicely, like this:

Copy

```
>>> i = 256*256
>>> print('The value of i is', i)
The value of i is 65536
```

The keyword argument *end* can be used to avoid the newline after the output, or end the output with a different string:

Copy

```
>>> a, b = 0, 1
>>> while a < 1000:
...     print(a, end=',')
...     a, b = b, a+b
...
0,1,1,2,3,5,8,13,21,34,55,89,144,233,377,610,987,
```

Footnotes

[1]

Since

```
**
```

has higher precedence than

```
-
```

```
,
```

```
-3**2
```

will be interpreted as

```
-(3**2)
```

and thus result in

```
-9
```

. To avoid this and get

```
9
```

, you can use

```
(-3)**2
```

.

[\[2\]](#)

Unlike other languages, special characters such as

```
\n
```

have the same meaning with both single (

```
'...'
```

) and double (

```
"..."
```

) quotes. The only difference between the two is that within single quotes you don't need to escape

```
"
```

(but you have to escape

```
\'
```

) and vice versa.

4. More Control Flow Tools¶

As well as the

```
while
```

statement just introduced, Python uses a few more that we will encounter in this chapter.

4.1.

```
if
```

Statements¶

Perhaps the most well-known statement type is the

```
if
```

statement. For example:

Copy

```
>>> x = int(input("Please enter an integer: "))
Please enter an integer: 42
>>> if x < 0:
...     x = 0
...     print('Negative changed to zero')
... elif x == 0:
...     print('Zero')
... elif x == 1:
...     print('Single')
... else:
...     print('More')
...
More
```

There can be zero or more

```
elif
```

parts, and the

```
else
```

part is optional. The keyword ‘

```
elif
```

' is short for 'else if', and is useful to avoid excessive indentation. An

```
if
```

```
...
```

```
elif
```

```
...
```

```
elif
```

... sequence is a substitute for the

```
switch
```

or

```
case
```

statements found in other languages.

If you're comparing the same value to several constants, or checking for specific types or attributes, you may also find the

```
match
```

statement useful. For more details see [match Statements](#).

4.2.

```
for
```

Statements¶

The

```
for
```

statement in Python differs a bit from what you may be used to in C or Pascal. Rather than always iterating over an arithmetic progression of numbers (like in Pascal), or giving the user the ability to define both the iteration step and halting condition (as C), Python's

```
for
```

statement iterates over the items of any sequence (a list or a string), in the order that they appear in the sequence. For example (no pun intended):

Copy

```
>>> # Measure some strings:
>>> words = ['cat', 'window', 'defenestrate']
>>> for w in words:
...     print(w, len(w))
...
cat 3
window 6
defenestrate 12
```

Code that modifies a collection while iterating over that same collection can be tricky to get right. Instead, it is usually more straight-forward to loop over a copy of the collection or to create a new collection:

Copy

```
# Create a sample collection
users = {'Hans': 'active', 'Éléonore': 'inactive', '景太郎': 'active'}

# Strategy: Iterate over a copy
for user, status in users.copy().items():
    if status == 'inactive':
        del users[user]

# Strategy: Create a new collection
active_users = {}
for user, status in users.items():
    if status == 'active':
        active_users[user] = status
```

4.3. The

range()

Function¶

If you do need to iterate over a sequence of numbers, the built-in function

range()

comes in handy. It generates arithmetic progressions:

Copy

```
>>> for i in range(5):
...     print(i)
...
0
1
2
3
4
```

The given end point is never part of the generated sequence;

```
range(10)
```

generates 10 values, the legal indices for items of a sequence of length 10. It is possible to let the range start at another number, or to specify a different increment (even negative; sometimes this is called the 'step'):

Copy

```
>>> list(range(5, 10))
[5, 6, 7, 8, 9]

>>> list(range(0, 10, 3))
[0, 3, 6, 9]

>>> list(range(-10, -100, -30))
[-10, -40, -70]
```

To iterate over the indices of a sequence, you can combine

```
range()
```

and

```
len()
```

as follows:

Copy

```
>>> a = ['Mary', 'had', 'a', 'little', 'lamb']
>>> for i in range(len(a)):
...     print(i, a[i])
...
0 Mary
1 had
2 a
3 little
4 lamb
```

In most such cases, however, it is convenient to use the

```
enumerate()
```

function, see Looping Techniques.

A strange thing happens if you just print a range:

Copy

```
>>> range(10)
range(0, 10)
```

In many ways the object returned by

```
range()
```

behaves as if it is a list, but in fact it isn't. It is an object which returns the successive items of the desired sequence when you iterate over it, but it doesn't really make the list, thus saving space.

We say such an object is iterable, that is, suitable as a target for functions and constructs that expect something from which they can obtain successive items until the supply is exhausted. We have seen that the

```
for
```

statement is such a construct, while an example of a function that takes an iterable is

```
sum()
```

:

Copy

```
>>> sum(range(4)) # 0 + 1 + 2 + 3
6
```

Later we will see more functions that return iterables and take iterables as arguments. In chapter Data Structures, we will discuss

```
list()
```

in more detail.

4.4.

break

and

continue

Statements¶

The

```
break
```

statement breaks out of the innermost enclosing

```
for
```

or

```
while
```

loop:

Copy

```
>>> for n in range(2, 10):
...     for x in range(2, n):
...         if n % x == 0:
...             print(f"{n} equals {x} * {n//x}")
...             break
...
4 equals 2 * 2
6 equals 2 * 3
8 equals 2 * 4
9 equals 3 * 3
```

The

```
continue
```

statement continues with the next iteration of the loop:

Copy

```
>>> for num in range(2, 10):  
...     if num % 2 == 0:  
...         print(f"Found an even number {num}")  
...         continue  
...     print(f"Found an odd number {num}")  
...  
Found an even number 2  
Found an odd number 3  
Found an even number 4  
Found an odd number 5  
Found an even number 6  
Found an odd number 7  
Found an even number 8  
Found an odd number 9
```

4.5.

else

Clauses on Loops¶

In a

for

or

while

loop the

break

statement may be paired with an

else

clause. If the loop finishes without executing the

break

, the

else

clause executes.

In a

```
for
```

loop, the

```
else
```

clause is executed after the loop finishes its final iteration, that is, if no break occurred.

In a

```
while
```

loop, it's executed after the loop's condition becomes false.

In either kind of loop, the

```
else
```

clause is **not** executed if the loop was terminated by a

```
break
```

. Of course, other ways of ending the loop early, such as a

```
return
```

or a raised exception, will also skip execution of the

```
else
```

clause.

This is exemplified in the following

```
for
```

loop, which searches for prime numbers:

Copy


```
>>> for n in range(2, 10):
...     for x in range(2, n):
...         if n % x == 0:
...             print(n, 'equals', x, '*', n//x)
...             break
...         else:
...             # loop fell through without finding a factor
...             print(n, 'is a prime number')
...
2 is a prime number
3 is a prime number
4 equals 2 * 2
5 is a prime number
6 equals 2 * 3
7 is a prime number
8 equals 2 * 4
9 equals 3 * 3
```

(Yes, this is the correct code. Look closely: the

```
else
```

clause belongs to the

```
for
```

loop, **not** the

```
if
```

statement.)

One way to think of the else clause is to imagine it paired with the

```
if
```

inside the loop. As the loop executes, it will run a sequence like if/if/if/else. The

```
if
```

is inside the loop, encountered a number of times. If the condition is ever true, a

```
break
```

will happen. If the condition is never true, the

```
else
```

clause outside the loop will execute.

When used with a loop, the

```
else
```

clause has more in common with the

```
else
```

clause of a

```
try
```

statement than it does with that of

```
if
```

statements: a

```
try
```

statement's

```
else
```

clause runs when no exception occurs, and a loop's

```
else
```

clause runs when no

```
break
```

occurs. For more on the

```
try
```

statement and exceptions, see Handling Exceptions.

4.6.

```
pass
```

Statements¶

The

```
pass
```

statement does nothing. It can be used when a statement is required syntactically but the program requires no action. For example:

Copy

```
>>> while True:
...     pass # Busy-wait for keyboard interrupt (Ctrl+C)
...
```

This is commonly used for creating minimal classes:

Copy

```
>>> class MyEmptyClass:
...     pass
...
```

Another place

```
pass
```

can be used is as a place-holder for a function or conditional body when you are working on new code, allowing you to keep thinking at a more abstract level. The

```
pass
```

is silently ignored:

Copy

```
>>> def initlog(*args):
...     pass # Remember to implement this!
...
```

For this last case, many people use the ellipsis literal

```
...
```

instead of

```
pass
```

. This use has no special meaning to Python, and is not part of the language definition (you could use any constant expression here), but

```
...
```

is used conventionally as a placeholder body as well. See [The Ellipsis Object](#).

4.7.

match

Statements¶

A

```
match
```

statement takes an expression and compares its value to successive patterns given as one or more case blocks. This is superficially similar to a switch statement in C, Java or JavaScript (and many other languages), but it's more similar to pattern matching in languages like Rust or Haskell. Only the first pattern that matches gets executed and it can also extract components (sequence elements or object attributes) from the value into variables. If no case matches, none of the branches is executed.

The simplest form compares a subject value against one or more literals:

Copy

```
def http_error(status):  
    match status:  
        case 400:  
            return "Bad request"  
        case 404:  
            return "Not found"  
        case 418:  
            return "I'm a teapot"  
        case _:  
            return "Something's wrong with the internet"
```

Note the last block: the “variable name”

```
—
```

acts as a *wildcard* and never fails to match.

You can combine several literals in a single pattern using

```
|
```

(“or”):

Copy

```
case 401 | 403 | 404:  
    return "Not allowed"
```

Patterns can look like unpacking assignments, and can be used to bind variables:

Copy

```
# point is an (x, y) tuple
match point:
    case (0, 0):
        print("Origin")
    case (0, y):
        print(f"Y={y}")
    case (x, 0):
        print(f"X={x}")
    case (x, y):
        print(f"X={x}, Y={y}")
    case _:
        raise ValueError("Not a point")
```

Study that one carefully! The first pattern has two literals, and can be thought of as an extension of the literal pattern shown above. But the next two patterns combine a literal and a variable, and the variable *binds* a value from the subject (

```
point
```

). The fourth pattern captures two values, which makes it conceptually similar to the unpacking assignment

```
(x, y) = point
```

If you are using classes to structure your data you can use the class name followed by an argument list resembling a constructor, but with the ability to capture attributes into variables:

Copy

```

class Point:
    def __init__(self, x, y):
        self.x = x
        self.y = y

def where_is(point):
    match point:
        case Point(x=0, y=0):
            print("Origin")
        case Point(x=0, y=y):
            print(f"Y={y}")
        case Point(x=x, y=0):
            print(f"X={x}")
        case Point():
            print("Somewhere else")
        case _:
            print("Not a point")

```

You can use positional parameters with some builtin classes that provide an ordering for their attributes (e.g. dataclasses). You can also define a specific position for attributes in patterns by setting the

```
__match_args__
```

special attribute in your classes. If it's set to ("x", "y"), the following patterns are all equivalent (and all bind the

```
y
```

attribute to the

```
var
```

variable):

Copy

```

Point(1, var)
Point(1, y=var)
Point(x=1, y=var)
Point(y=var, x=1)

```

A recommended way to read patterns is to look at them as an extended form of what you would put on the left of an assignment, to understand which variables would be set to what. Only the standalone names (like

```
var
```

above) are assigned to by a match statement. Dotted names (like

```
foo.bar
```

), attribute names (the

```
x=
```

and

```
y=
```

above) or class names (recognized by the “(...)” next to them like

```
Point
```

above) are never assigned to.

Patterns can be arbitrarily nested. For example, if we have a short list of Points, with

```
__match_args__
```

added, we could match it like this:

Copy

```
class Point:
    __match_args__ = ('x', 'y')
    def __init__(self, x, y):
        self.x = x
        self.y = y

match points:
    case []:
        print("No points")
    case [Point(0, 0)]:
        print("The origin")
    case [Point(x, y)]:
        print(f"Single point {x}, {y}")
    case [Point(0, y1), Point(0, y2)]:
        print(f"Two on the Y axis at {y1}, {y2}")
    case _:
        print("Something else")
```

We can add an

```
if
```

clause to a pattern, known as a “guard”. If the guard is false,

```
match
```

goes on to try the next case block. Note that value capture happens before the guard is evaluated:

Copy

```
match point:
    case Point(x, y) if x == y:
        print(f"Y=X at {x}")
    case Point(x, y):
        print(f"Not on the diagonal")
```

Several other key features of this statement:

- Like unpacking assignments, tuple and list patterns have exactly the same meaning and actually match arbitrary sequences. An important exception is that they don't match iterators or strings.
- Sequence patterns support extended unpacking:

```
[x, y, *rest]
```

and

```
(x, y,
 *rest)
```

work similar to unpacking assignments. The name after

```
*
```

may also be

```
—
```

, so

```
(x, y, *_)
```

matches a sequence of at least two items without binding the remaining items.

- Mapping patterns:

```
{"bandwidth": b, "latency": l}
```

captures the

```
"bandwidth"
```

and

```
"latency"
```

values from a dictionary. Unlike sequence patterns, extra keys are ignored. An unpacking like

```
**rest
```


is also supported. (But

```
**  
_
```

would be redundant, so it is not allowed.)

- Subpatterns may be captured using the

```
as
```

keyword:

Copy

```
case (Point(x1, y1), Point(x2, y2) as p2): ...
```

will capture the second element of the input as

```
p2
```

(as long as the input is a sequence of two points)

- Most literals are compared by equality, however the singletons

```
True
```

,

```
False
```

and

```
None
```

are compared by identity.

- Patterns may use named constants. These must be dotted names to prevent them from being interpreted as capture variables:

Copy

```

from enum import Enum
class Color(Enum):
    RED = 'red'
    GREEN = 'green'
    BLUE = 'blue'

color = Color(input("Enter your choice of 'red', 'blue' or 'green': "))

match color:
    case Color.RED:
        print("I see red!")
    case Color.GREEN:
        print("Grass is green")
    case Color.BLUE:
        print("I'm feeling the blues :(")

```

For a more detailed explanation and additional examples, you can look into [PEP 636](#) which is written in a tutorial format.

4.8. Defining Functions¶

We can create a function that writes the Fibonacci series to an arbitrary boundary:

Copy

```

>>> def fib(n):    # write Fibonacci series less than n
...     """Print a Fibonacci series less than n."""
...     a, b = 0, 1
...     while a < n:
...         print(a, end=' ')
...         a, b = b, a+b
...     print()
...
>>> # Now call the function we just defined:
>>> fib(2000)
0 1 1 2 3 5 8 13 21 34 55 89 144 233 377 610 987 1597

```

The keyword

def

introduces a function *definition*. It must be followed by the function name and the parenthesized list of formal parameters. The statements that form the body of the function start at the next line, and must be indented.

The first statement of the function body can optionally be a string literal; this string literal is the function's documentation string, or *docstring*. (More about docstrings can be found in the section [Documentation Strings](#).) There are tools which use docstrings to automatically produce online or printed documentation, or to let the user interactively browse through code; it's good practice to include docstrings in code that you write, so make a habit of it.

The *execution* of a function introduces a new symbol table used for the local variables of the function. More precisely, all variable assignments in a function store the value in the local symbol table; whereas variable references first look in the local symbol table, then in the local symbol tables of enclosing functions, then in the global symbol table, and finally in the table of built-in names. Thus, global variables and variables of enclosing functions cannot be directly assigned a value within a function (unless, for global variables, named in a

```
global
```

statement, or, for variables of enclosing functions, named in a

```
nonlocal
```

statement), although they may be referenced.

The actual parameters (arguments) to a function call are introduced in the local symbol table of the called function when it is called; thus, arguments are passed using *call by value* (where the *value* is always an object *reference*, not the value of the object). [1] When a function calls another function, or calls itself recursively, a new local symbol table is created for that call.

A function definition associates the function name with the function object in the current symbol table. The interpreter recognizes the object pointed to by that name as a user-defined function. Other names can also point to that same function object and can also be used to access the function:

Copy

```
>>> fib
<function fib at 10042ed0>
>>> f = fib
>>> f(100)
0 1 1 2 3 5 8 13 21 34 55 89
```

Coming from other languages, you might object that

```
fib
```

is not a function but a procedure since it doesn't return a value. In fact, even functions without a

```
return
```

statement do return a value, albeit a rather boring one. This value is called

```
None
```

(it's a built-in name). Writing the value

```
None
```

is normally suppressed by the interpreter if it would be the only value written. You can see it if you really want to using

```
print()
```

```
:
```

Copy

```
>>> fib(0)
>>> print(fib(0))
None
```

It is simple to write a function that returns a list of the numbers of the Fibonacci series, instead of printing it:

Copy

```
>>> def fib2(n): # return Fibonacci series up to n
...     """Return a list containing the Fibonacci series up to n."""
...     result = []
...     a, b = 0, 1
...     while a < n:
...         result.append(a)    # see below
...         a, b = b, a+b
...     return result
...
>>> f100 = fib2(100)    # call it
>>> f100                # write the result
[0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89]
```

This example, as usual, demonstrates some new Python features:

- The

```
return
```

statement returns with a value from a function.

```
return
```

without an expression argument returns

```
None
```

. Falling off the end of a function also returns

```
None
```

.

- The statement

```
result.append(a)
```

calls a *method* of the list object

```
result
```

. A method is a function that 'belongs' to an object and is named

```
obj.methodname
```

, where

```
obj
```

is some object (this may be an expression), and

```
methodname
```

is the name of a method that is defined by the object's type. Different types define different methods. Methods of different types may have the same name without causing ambiguity. (It is possible to define your own object types and methods, using *classes*, see [Classes](#)) The method

```
append()
```

shown in the example is defined for list objects; it adds a new element at the end of the list. In this example it is equivalent to

```
result = result + [a]
```

, but more efficient.

4.9. More on Defining Functions¶

It is also possible to define functions with a variable number of arguments. There are three forms, which can be combined.

4.9.1. Default Argument Values¶

The most useful form is to specify a default value for one or more arguments. This creates a function that can be called with fewer arguments than it is defined to allow. For example:

Copy

```
def ask_ok(prompt, retries=4, reminder='Please try again!'):
    while True:
        reply = input(prompt)
        if reply in {'y', 'ye', 'yes'}:
            return True
        if reply in {'n', 'no', 'nop', 'nope'}:
            return False
        retries = retries - 1
        if retries < 0:
            raise ValueError('invalid user response')
        print(reminder)
```

This function can be called in several ways:

- giving only the mandatory argument:

```
ask_ok('Do you really want to quit?')
```

- giving one of the optional arguments:

```
ask_ok('OK to overwrite the file?', 2)
```

- or even giving all arguments:

```
ask_ok('OK to overwrite the file?', 2, 'Come on, only yes or no!')
```

This example also introduces the

```
in
```

keyword. This tests whether or not a sequence contains a certain value.

The default values are evaluated at the point of function definition in the *defining* scope, so that

Copy

```
i = 5

def f(arg=i):
    print(arg)

i = 6
f()
```

will print

Important warning: The default value is evaluated only once. This makes a difference when the default is a mutable object such as a list, dictionary, or instances of most classes. For example, the following function accumulates the arguments passed to it on subsequent calls:

Copy

```
def f(a, L=[]):
    L.append(a)
    return L

print(f(1))
print(f(2))
print(f(3))
```

This will print

Copy

```
[1]
[1, 2]
[1, 2, 3]
```

If you don't want the default to be shared between subsequent calls, you can write the function like this instead:

Copy

```
def f(a, L=None):
    if L is None:
        L = []
    L.append(a)
    return L
```

4.9.2. Keyword Arguments¶

Functions can also be called using keyword arguments of the form

```
kwarg=value
```

. For instance, the following function:

Copy

```
def parrot(voltage, state='a stiff', action='vroom', type='Norwegian Blue'):
    print("-- This parrot wouldn't", action, end=' ')
    print("if you put", voltage, "volts through it.")
    print("-- Lovely plumage, the", type)
    print("-- It's", state, "!")
```

accepts one required argument (

voltage

) and three optional arguments (

state

,

action

, and

type

). This function can be called in any of the following ways:

Copy

```
parrot(1000)                # 1 positional argument
parrot(voltage=1000)        # 1 keyword argument
parrot(voltage=1000000, action='V00000M') # 2 keyword arguments
parrot(action='V00000M', voltage=1000000)  # 2 keyword arguments
parrot('a million', 'bereft of life', 'jump') # 3 positional arguments
parrot('a thousand', state='pushing up the daisies') # 1 positional, 1 keyword
```

but all the following calls would be invalid:

Copy

```
parrot()                # required argument missing
parrot(voltage=5.0, 'dead') # non-keyword argument after a keyword argument
parrot(110, voltage=220)   # duplicate value for the same argument
parrot(actor='John Cleese') # unknown keyword argument
```

In a function call, keyword arguments must follow positional arguments. All the keyword arguments passed must match one of the arguments accepted by the function (e.g.

actor

is not a valid argument for the

```
parrot
```

function), and their order is not important. This also includes non-optional arguments (e.g.

```
parrot(voltage=1000)
```

is valid too). No argument may receive a value more than once. Here's an example that fails due to this restriction:

Copy

```
>>> def function(a):
...     pass
...
>>> function(0, a=0)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: function() got multiple values for argument 'a'
```

When a final formal parameter of the form

```
**name
```

is present, it receives a dictionary (see Mapping Types — dict) containing all keyword arguments except for those corresponding to a formal parameter. This may be combined with a formal parameter of the form

```
*name
```

(described in the next subsection) which receives a tuple containing the positional arguments beyond the formal parameter list. (

```
*name
```

must occur before

```
**name
```

.) For example, if we define a function like this:

Copy

```
def cheeseshop(kind, *arguments, **keywords):
    print("-- Do you have any", kind, "?")
    print("-- I'm sorry, we're all out of", kind)
    for arg in arguments:
        print(arg)
    print("-" * 40)
    for kw in keywords:
        print(kw, ":", keywords[kw])
```

It could be called like this:

Copy

```
cheeseshop("Limburger", "It's very runny, sir.",
           "It's really very, VERY runny, sir.",
           shopkeeper="Michael Palin",
           client="John Cleese",
           sketch="Cheese Shop Sketch")
```

and of course it would print:

```
-- Do you have any Limburger ?
-- I'm sorry, we're all out of Limburger
It's very runny, sir.
It's really very, VERY runny, sir.
-----
shopkeeper : Michael Palin
client : John Cleese
sketch : Cheese Shop Sketch
```

Note that the order in which the keyword arguments are printed is guaranteed to match the order in which they were provided in the function call.

4.9.3. Special parameters¶

By default, arguments may be passed to a Python function either by position or explicitly by keyword. For readability and performance, it makes sense to restrict the way arguments can be passed so that a developer need only look at the function definition to determine if items are passed by position, by position or keyword, or by keyword.

A function definition may look like:

```
def f(pos1, pos2, /, pos_or_kwd, *, kwd1, kwd2):
    -----
    |           |           |
    |           | Positional or keyword |
    |           |           |
    |           |           | - Keyword only
    -- Positional only
```

where

and

are optional. If used, these symbols indicate the kind of parameter by how the arguments may be passed to the function: positional-only, positional-or-keyword, and keyword-only. Keyword parameters are also referred to as named parameters.

4.9.3.1. Positional-or-Keyword Arguments

and

are not present in the function definition, arguments may be passed to a function by position or by keyword.

4.9.3.2. Positional-Only Parameters

Looking at this in a bit more detail, it is possible to mark certain parameters as *positional-only*. If *positional-only*, the parameters' order matters, and the parameters cannot be passed by keyword. Positional-only parameters are placed before a

(forward-slash). The

is used to logically separate the positional-only parameters from the rest of the parameters. If there is no

in the function definition, there are no positional-only parameters.

Parameters following the

```
/
```

may be *positional-or-keyword* or *keyword-only*.

4.9.3.3. Keyword-Only Arguments¶

To mark parameters as *keyword-only*, indicating the parameters must be passed by keyword argument, place an

```
*
```

in the arguments list just before the first *keyword-only* parameter.

4.9.3.4. Function Examples¶

Consider the following example function definitions paying close attention to the markers

```
/
```

and

```
*
```

```
:
```

Copy

```
>>> def standard_arg(arg):
...     print(arg)
...
>>> def pos_only_arg(arg, /):
...     print(arg)
...
>>> def kwd_only_arg(*, arg):
...     print(arg)
...
>>> def combined_example(pos_only, /, standard, *, kwd_only):
...     print(pos_only, standard, kwd_only)
```

The first function definition,

```
standard_arg
```

, the most familiar form, places no restrictions on the calling convention and arguments may be passed by position or keyword:

Copy

```
>>> standard_arg(2)
2

>>> standard_arg(arg=2)
2
```

The second function

```
pos_only_arg
```

is restricted to only use positional parameters as there is a

```
/
```

in the function definition:

Copy

```
>>> pos_only_arg(1)
1

>>> pos_only_arg(arg=1)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: pos_only_arg() got some positional-only arguments passed as keyword arguments: 'arg'
```

The third function

```
kwd_only_arg
```

only allows keyword arguments as indicated by a

```
*
```

in the function definition:

Copy

```
>>> kwd_only_arg(3)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: kwd_only_arg() takes 0 positional arguments but 1 was given

>>> kwd_only_arg(arg=3)
3
```

And the last uses all three calling conventions in the same function definition:

Copy

```
>>> combined_example(1, 2, 3)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: combined_example() takes 2 positional arguments but 3 were given

>>> combined_example(1, 2, kwd_only=3)
1 2 3

>>> combined_example(1, standard=2, kwd_only=3)
1 2 3

>>> combined_example(pos_only=1, standard=2, kwd_only=3)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: combined_example() got some positional-only arguments passed as keyword arguments: 'pos_only'
```

Finally, consider this function definition which has a potential collision between the positional argument

```
name
```

and

```
**kwds
```

which has

```
name
```

as a key:

Copy

```
def foo(name, **kwds):
    return 'name' in kwds
```

There is no possible call that will make it return

```
True
```

as the keyword

```
'name'
```

will always bind to the first parameter. For example:

Copy

```
>>> foo(1, **{'name': 2})
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: foo() got multiple values for argument 'name'
>>>
```

But using

/

(positional only arguments), it is possible since it allows

name

as a positional argument and

'name'

as a key in the keyword arguments:

Copy

```
>>> def foo(name, /, **kws):
...     return 'name' in kws
...
>>> foo(1, **{'name': 2})
True
```

In other words, the names of positional-only parameters can be used in

**kws

without ambiguity.

4.9.3.5. Recap¶

The use case will determine which parameters to use in the function definition:

Copy

```
def f(pos1, pos2, /, pos_or_kwd, *, kwd1, kwd2):
```

As guidance:

- Use positional-only if you want the name of the parameters to not be available to the user. This is useful when parameter names have no real meaning, if you want to enforce the order

of the arguments when the function is called or if you need to take some positional parameters and arbitrary keywords.

- Use keyword-only when names have meaning and the function definition is more understandable by being explicit with names or you want to prevent users relying on the position of the argument being passed.
- For an API, use positional-only to prevent breaking API changes if the parameter's name is modified in the future.

4.9.4. Arbitrary Argument Lists¶

Finally, the least frequently used option is to specify that a function can be called with an arbitrary number of arguments. These arguments will be wrapped up in a tuple (see Tuples and Sequences). Before the variable number of arguments, zero or more normal arguments may occur.

Copy

```
def write_multiple_items(file, separator, *args):
    file.write(separator.join(args))
```

Normally, these *variadic* arguments will be last in the list of formal parameters, because they scoop up all remaining input arguments that are passed to the function. Any formal parameters which occur after the

```
*args
```

parameter are 'keyword-only' arguments, meaning that they can only be used as keywords rather than positional arguments.

Copy

```
>>> def concat(*args, sep="/"):
...     return sep.join(args)
...
>>> concat("earth", "mars", "venus")
'earth/mars/venus'
>>> concat("earth", "mars", "venus", sep=".")
'earth.mars.venus'
```

4.9.5. Unpacking Argument Lists¶

The reverse situation occurs when the arguments are already in a list or tuple but need to be unpacked for a function call requiring separate positional arguments. For instance, the built-in

```
range()
```


function expects separate *start* and *stop* arguments. If they are not available separately, write the function call with the

```
*
```

-operator to unpack the arguments out of a list or tuple:

Copy

```
>>> list(range(3, 6))           # normal call with separate arguments
[3, 4, 5]
>>> args = [3, 6]
>>> list(range(*args))         # call with arguments unpacked from a list
[3, 4, 5]
```

In the same fashion, dictionaries can deliver keyword arguments with the

```
**
```

-operator:

Copy

```
>>> def parrot(voltage, state='a stiff', action='vroom'):
...     print("-- This parrot wouldn't", action, end=' ')
...     print("if you put", voltage, "volts through it.", end=' ')
...     print("E's", state, "!")
...
>>> d = {"voltage": "four million", "state": "bleedin' demised", "action": "V
OOM"}
>>> parrot(**d)
-- This parrot wouldn't VOOM if you put four million volts through it. E's bl
eedin' demised !
```

4.9.6. Lambda Expressions ¶

Small anonymous functions can be created with the

```
lambda
```

keyword. This function returns the sum of its two arguments:

```
lambda a, b: a+b
```

. Lambda functions can be used wherever function objects are required. They are syntactically restricted to a single expression. Semantically, they are just syntactic sugar for a normal function definition. Like nested function definitions, lambda functions can reference variables from the containing scope:

Copy

```
>>> def make_incrementor(n):  
...     return lambda x: x + n  
...  
>>> f = make_incrementor(42)  
>>> f(0)  
42  
>>> f(1)  
43
```

The above example uses a lambda expression to return a function. Another use is to pass a small function as an argument. For instance,

```
list.sort()
```

takes a sorting key function *key* which can be a lambda function:

Copy

```
>>> pairs = [(1, 'one'), (2, 'two'), (3, 'three'), (4, 'four')]  
>>> pairs.sort(key=lambda pair: pair[1])  
>>> pairs  
[(4, 'four'), (1, 'one'), (3, 'three'), (2, 'two')]
```

4.9.7. Documentation Strings¶

Here are some conventions about the content and formatting of documentation strings.

The first line should always be a short, concise summary of the object's purpose. For brevity, it should not explicitly state the object's name or type, since these are available by other means (except if the name happens to be a verb describing a function's operation). This line should begin with a capital letter and end with a period.

If there are more lines in the documentation string, the second line should be blank, visually separating the summary from the rest of the description. The following lines should be one or more paragraphs describing the object's calling conventions, its side effects, etc.

The Python parser strips indentation from multi-line string literals when they serve as module, class, or function docstrings.

Here is an example of a multi-line docstring:

Copy

```

>>> def my_function():
...     """Do nothing, but document it.
...
...     No, really, it doesn't do anything:
...
...     >>> my_function()
...     >>>
...     """
...     pass
...
>>> print(my_function.__doc__)
Do nothing, but document it.

No, really, it doesn't do anything:

    >>> my_function()
    >>>

```

4.9.8. Function Annotations¶

Function annotations are completely optional metadata information about the types used by user-defined functions (see [PEP 3107](#) and [PEP 484](#) for more information).

Annotations are stored in the

```
__annotations__
```

attribute of the function as a dictionary and have no effect on any other part of the function.

Parameter annotations are defined by a colon after the parameter name, followed by an expression evaluating to the value of the annotation. Return annotations are defined by a literal

```
->
```

, followed by an expression, between the parameter list and the colon denoting the end of the

```
def
```

statement. The following example has a required argument, an optional argument, and the return value annotated:

Copy

```
>>> def f(ham: str, eggs: str = 'eggs') -> str:
...     print("Annotations:", f.__annotations__)
...     print("Arguments:", ham, eggs)
...     return ham + ' and ' + eggs
...
>>> f('spam')
Annotations: {'ham': <class 'str'>, 'return': <class 'str'>, 'eggs': <class
'str'>}}
Arguments: spam eggs
'spam and eggs'
```

4.10. Intermezzo: Coding Style¶

Now that you are about to write longer, more complex pieces of Python, it is a good time to talk about *coding style*. Most languages can be written (or more concisely, *formatted*) in different styles; some are more readable than others. Making it easy for others to read your code is always a good idea, and adopting a nice coding style helps tremendously for that.

For Python, [PEP 8](#) has emerged as the style guide that most projects adhere to; it promotes a very readable and eye-pleasing coding style. Every Python developer should read it at some point; here are the most important points extracted for you:

- Use 4-space indentation, and no tabs.

4 spaces are a good compromise between small indentation (allows greater nesting depth) and large indentation (easier to read). Tabs introduce confusion, and are best left out.

- Wrap lines so that they don't exceed 79 characters.

This helps users with small displays and makes it possible to have several code files side-by-side on larger displays.

- Use blank lines to separate functions and classes, and larger blocks of code inside functions.
- When possible, put comments on a line of their own.
- Use docstrings.
- Use spaces around operators and after commas, but not directly inside bracketing constructs:

```
a = f(1, 2) + g(3, 4)
```

- Name your classes and functions consistently; the convention is to use

```
UpperCamelCase
```

for classes and

```
lowercase_with_underscores
```

for functions and methods. Always use

```
self
```

as the name for the first method argument (see A First Look at Classes for more on classes and methods).

- Don't use fancy encodings if your code is meant to be used in international environments. Python's default, UTF-8, or even plain ASCII work best in any case.
- Likewise, don't use non-ASCII characters in identifiers if there is only the slightest chance people speaking a different language will read or maintain the code.

Footnotes

[\[1\]](#)

Actually, *call by object reference* would be a better description, since if a mutable object is passed, the caller will see any changes the callee makes to it (items inserted into a list).

5. Data Structures¶

This chapter describes some things you've learned about already in more detail, and adds some new things as well.

5.1. More on Lists¶

The list data type has some more methods. Here are all of the methods of list objects:

`list.append(x)`

Add an item to the end of the list. Similar to

```
a[len(a):] = [x]
```

.

`list.extend(iterable)`

Extend the list by appending all the items from the iterable. Similar to

```
a[len(a):] = iterable
```

.

`list.insert(i, x)`

Insert an item at a given position. The first argument is the index of the element before which to insert, so

```
a.insert(0, x)
```

inserts at the front of the list, and

```
a.insert(len(a), x)
```

is equivalent to

```
a.append(x)
```

.

`list.remove(x)`

Remove the first item from the list whose value is equal to `x`. It raises a

```
ValueError
```

if there is no such item.

`list.pop([i])`

Remove the item at the given position in the list, and return it. If no index is specified,

```
a.pop()
```

removes and returns the last item in the list. It raises an

```
IndexError
```

if the list is empty or the index is outside the list range.

`list.clear()`

Remove all items from the list. Similar to

```
del a[:]
```

.

`list.index(x[, start[, end]])`

Return zero-based index of the first occurrence of *x* in the list. Raises a

```
ValueError
```

if there is no such item.

The optional arguments *start* and *end* are interpreted as in the slice notation and are used to limit the search to a particular subsequence of the list. The returned index is computed relative to the beginning of the full sequence rather than the *start* argument.

`list.count(x)`

Return the number of times *x* appears in the list.

`list.sort(*, key=None, reverse=False)`

Sort the items of the list in place (the arguments can be used for sort customization, see

```
sorted()
```

for their explanation).

`list.reverse()`

Reverse the elements of the list in place.

`list.copy()`

Return a shallow copy of the list. Similar to

```
a[:]
```

.

An example that uses most of the list methods:

Copy

```
>>> fruits = ['orange', 'apple', 'pear', 'banana', 'kiwi', 'apple', 'banana']
>>> fruits.count('apple')
2
>>> fruits.count('tangerine')
0
>>> fruits.index('banana')
3
>>> fruits.index('banana', 4) # Find next banana starting at position 4
6
>>> fruits.reverse()
>>> fruits
['banana', 'apple', 'kiwi', 'banana', 'pear', 'apple', 'orange']
>>> fruits.append('grape')
>>> fruits
['banana', 'apple', 'kiwi', 'banana', 'pear', 'apple', 'orange', 'grape']
>>> fruits.sort()
>>> fruits
['apple', 'apple', 'banana', 'banana', 'grape', 'kiwi', 'orange', 'pear']
>>> fruits.pop()
'pear'
```

You might have noticed that methods like

insert

,

remove

or

sort

that only modify the list have no return value printed – they return the default

None

. [\[1\]](#) This is a design principle for all mutable data structures in Python.

Another thing you might notice is that not all data can be sorted or compared. For instance,

[None, 'hello', 10]

doesn't sort because integers can't be compared to strings and

None

can't be compared to other types. Also, there are some types that don't have a defined ordering relation. For example,

```
3+4j < 5+7j
```

isn't a valid comparison.

5.1.1. Using Lists as Stacks¶

The list methods make it very easy to use a list as a stack, where the last element added is the first element retrieved ("last-in, first-out"). To add an item to the top of the stack, use

```
append()
```

. To retrieve an item from the top of the stack, use

```
pop()
```

without an explicit index. For example:

Copy

```
>>> stack = [3, 4, 5]
>>> stack.append(6)
>>> stack.append(7)
>>> stack
[3, 4, 5, 6, 7]
>>> stack.pop()
7
>>> stack
[3, 4, 5, 6]
>>> stack.pop()
6
>>> stack.pop()
5
>>> stack
[3, 4]
```

5.1.2. Using Lists as Queues¶

It is also possible to use a list as a queue, where the first element added is the first element retrieved ("first-in, first-out"); however, lists are not efficient for this purpose. While appends and pops from the end of list are fast, doing inserts or pops from the beginning of a list is slow (because all of the other elements have to be shifted by one).

To implement a queue, use

```
collections.deque
```

which was designed to have fast appends and pops from both ends. For example:

Copy

```
>>> from collections import deque
>>> queue = deque(["Eric", "John", "Michael"])
>>> queue.append("Terry")           # Terry arrives
>>> queue.append("Graham")         # Graham arrives
>>> queue.popleft()                # The first to arrive now leaves
'Eric'
>>> queue.popleft()                # The second to arrive now leaves
'John'
>>> queue                          # Remaining queue in order of arrival
deque(['Michael', 'Terry', 'Graham'])
```

5.1.3. List Comprehensions¶

List comprehensions provide a concise way to create lists. Common applications are to make new lists where each element is the result of some operations applied to each member of another sequence or iterable, or to create a subsequence of those elements that satisfy a certain condition.

For example, assume we want to create a list of squares, like:

Copy

```
>>> squares = []
>>> for x in range(10):
...     squares.append(x**2)
...
>>> squares
[0, 1, 4, 9, 16, 25, 36, 49, 64, 81]
```

Note that this creates (or overwrites) a variable named

x

that still exists after the loop completes. We can calculate the list of squares without any side effects using:

Copy

```
squares = list(map(lambda x: x**2, range(10)))
```

or, equivalently:

Copy

```
squares = [x**2 for x in range(10)]
```

which is more concise and readable.

A list comprehension consists of brackets containing an expression followed by a

```
for
```

clause, then zero or more

```
for
```

or

```
if
```

clauses. The result will be a new list resulting from evaluating the expression in the context of the

```
for
```

and

```
if
```

clauses which follow it. For example, this listcomp combines the elements of two lists if they are not equal:

Copy

```
>>> [(x, y) for x in [1,2,3] for y in [3,1,4] if x != y]
[(1, 3), (1, 4), (2, 3), (2, 1), (2, 4), (3, 1), (3, 4)]
```

and it's equivalent to:

Copy

```
>>> combs = []
>>> for x in [1,2,3]:
...     for y in [3,1,4]:
...         if x != y:
...             combs.append((x, y))
...
>>> combs
[(1, 3), (1, 4), (2, 3), (2, 1), (2, 4), (3, 1), (3, 4)]
```

Note how the order of the

```
for
```

and

```
if
```

statements is the same in both these snippets.

If the expression is a tuple (e.g. the

```
(x, y)
```

in the previous example), it must be parenthesized.

Copy

```
>>> vec = [-4, -2, 0, 2, 4]
>>> # create a new list with the values doubled
>>> [x*2 for x in vec]
[-8, -4, 0, 4, 8]
>>> # filter the list to exclude negative numbers
>>> [x for x in vec if x >= 0]
[0, 2, 4]
>>> # apply a function to all the elements
>>> [abs(x) for x in vec]
[4, 2, 0, 2, 4]
>>> # call a method on each element
>>> freshfruit = [' banana', ' loganberry ', 'passion fruit ']
>>> [weapon.strip() for weapon in freshfruit]
['banana', 'loganberry', 'passion fruit']
>>> # create a list of 2-tuples like (number, square)
>>> [(x, x**2) for x in range(6)]
[(0, 0), (1, 1), (2, 4), (3, 9), (4, 16), (5, 25)]
>>> # the tuple must be parenthesized, otherwise an error is raised
>>> [x, x**2 for x in range(6)]
File "<stdin>", line 1
    [x, x**2 for x in range(6)]
    ^^^^^^^
SyntaxError: did you forget parentheses around the comprehension target?
>>> # flatten a list using a listcomp with two 'for'
>>> vec = [[1,2,3], [4,5,6], [7,8,9]]
>>> [num for elem in vec for num in elem]
[1, 2, 3, 4, 5, 6, 7, 8, 9]
```

List comprehensions can contain complex expressions and nested functions:

Copy

```
>>> from math import pi
>>> [str(round(pi, i)) for i in range(1, 6)]
['3.1', '3.14', '3.142', '3.1416', '3.14159']
```

5.1.4. Nested List Comprehensions¶

The initial expression in a list comprehension can be any arbitrary expression, including another list comprehension.

Consider the following example of a 3x4 matrix implemented as a list of 3 lists of length 4:

Copy

```
>>> matrix = [
...     [1, 2, 3, 4],
...     [5, 6, 7, 8],
...     [9, 10, 11, 12],
... ]
```

The following list comprehension will transpose rows and columns:

Copy

```
>>> [[row[i] for row in matrix] for i in range(4)]
[[1, 5, 9], [2, 6, 10], [3, 7, 11], [4, 8, 12]]
```

As we saw in the previous section, the inner list comprehension is evaluated in the context of the

for

that follows it, so this example is equivalent to:

Copy

```
>>> transposed = []
>>> for i in range(4):
...     transposed.append([row[i] for row in matrix])
...
>>> transposed
[[1, 5, 9], [2, 6, 10], [3, 7, 11], [4, 8, 12]]
```

which, in turn, is the same as:

Copy

```
>>> transposed = []
>>> for i in range(4):
...     # the following 3 lines implement the nested listcomp
...     transposed_row = []
...     for row in matrix:
...         transposed_row.append(row[i])
...     transposed.append(transposed_row)
...
>>> transposed
[[1, 5, 9], [2, 6, 10], [3, 7, 11], [4, 8, 12]]
```

In the real world, you should prefer built-in functions to complex flow statements. The

```
zip()
```

function would do a great job for this use case:

Copy

```
>>> list(zip(*matrix))
[(1, 5, 9), (2, 6, 10), (3, 7, 11), (4, 8, 12)]
```

See [Unpacking Argument Lists](#) for details on the asterisk in this line.

5.2. The

```
del
```

statement

There is a way to remove an item from a list given its index instead of its value: the

```
del
```

statement. This differs from the

```
pop()
```

method which returns a value. The

```
del
```

statement can also be used to remove slices from a list or clear the entire list (which we did earlier by assignment of an empty list to the slice). For example:

Copy

```
>>> a = [-1, 1, 66.25, 333, 333, 1234.5]
>>> del a[0]
>>> a
[1, 66.25, 333, 333, 1234.5]
>>> del a[2:4]
>>> a
[1, 66.25, 1234.5]
>>> del a[:]
>>> a
[]
```

```
del
```

can also be used to delete entire variables:

Copy

```
>>> del a
```

Referencing the name

```
a
```

hereafter is an error (at least until another value is assigned to it). We'll find other uses for

```
del
```

later.

5.3. Tuples and Sequences¶

We saw that lists and strings have many common properties, such as indexing and slicing operations. They are two examples of *sequence* data types (see Sequence Types — list, tuple, range). Since Python is an evolving language, other sequence data types may be added. There is also another standard sequence data type: the *tuple*.

A tuple consists of a number of values separated by commas, for instance:

Copy

```

>>> t = 12345, 54321, 'hello!'
>>> t[0]
12345
>>> t
(12345, 54321, 'hello!')
>>> # Tuples may be nested:
>>> u = t, (1, 2, 3, 4, 5)
>>> u
((12345, 54321, 'hello!'), (1, 2, 3, 4, 5))
>>> # Tuples are immutable:
>>> t[0] = 88888
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: 'tuple' object does not support item assignment
>>> # but they can contain mutable objects:
>>> v = ([1, 2, 3], [3, 2, 1])
>>> v
([1, 2, 3], [3, 2, 1])

```

As you see, on output tuples are always enclosed in parentheses, so that nested tuples are interpreted correctly; they may be input with or without surrounding parentheses, although often parentheses are necessary anyway (if the tuple is part of a larger expression). It is not possible to assign to the individual items of a tuple, however it is possible to create tuples which contain mutable objects, such as lists.

Though tuples may seem similar to lists, they are often used in different situations and for different purposes. Tuples are immutable, and usually contain a heterogeneous sequence of elements that are accessed via unpacking (see later in this section) or indexing (or even by attribute in the case of

namedtuples

). Lists are mutable, and their elements are usually homogeneous and are accessed by iterating over the list.

A special problem is the construction of tuples containing 0 or 1 items: the syntax has some extra quirks to accommodate these. Empty tuples are constructed by an empty pair of parentheses; a tuple with one item is constructed by following a value with a comma (it is not sufficient to enclose a single value in parentheses). Ugly, but effective. For example:

Copy


```
>>> empty = ()
>>> singleton = 'hello',    # <-- note trailing comma
>>> len(empty)
0
>>> len(singleton)
1
>>> singleton
('hello',)
```

The statement

```
t = 12345, 54321, 'hello!'
```

is an example of *tuple packing*: the values

```
12345
```

```
,
```

```
54321
```

and

```
'hello!'
```

are packed together in a tuple. The reverse operation is also possible:

Copy

```
>>> x, y, z = t
```

This is called, appropriately enough, *sequence unpacking* and works for any sequence on the right-hand side. Sequence unpacking requires that there are as many variables on the left side of the equals sign as there are elements in the sequence. Note that multiple assignment is really just a combination of tuple packing and sequence unpacking.

5.4. Sets ¶

Python also includes a data type for sets. A set is an unordered collection with no duplicate elements. Basic uses include membership testing and eliminating duplicate entries. Set objects also support mathematical operations like union, intersection, difference, and symmetric difference.

Curly braces or the

```
set()
```

function can be used to create sets. Note: to create an empty set you have to use

```
set()
```

, not

```
{}
```

; the latter creates an empty dictionary, a data structure that we discuss in the next section.

Here is a brief demonstration:

Copy

```
>>> basket = {'apple', 'orange', 'apple', 'pear', 'orange', 'banana'}
>>> print(basket)                # show that duplicates have been removed
{'orange', 'banana', 'pear', 'apple'}
>>> 'orange' in basket           # fast membership testing
True
>>> 'crabgrass' in basket
False

>>> # Demonstrate set operations on unique letters from two words
>>>
>>> a = set('abracadabra')
>>> b = set('alacazam')
>>> a                           # unique letters in a
{'a', 'r', 'b', 'c', 'd'}
>>> a - b                       # letters in a but not in b
{'r', 'd', 'b'}
>>> a | b                       # letters in a or b or both
{'a', 'c', 'r', 'd', 'b', 'm', 'z', 'l'}
>>> a & b                       # letters in both a and b
{'a', 'c'}
>>> a ^ b                       # letters in a or b but not both
{'r', 'd', 'b', 'm', 'z', 'l'}
```

Similarly to [list comprehensions](#), set comprehensions are also supported:

Copy

```
>>> a = {x for x in 'abracadabra' if x not in 'abc'}
>>> a
{'r', 'd'}
```

5.5. Dictionaries¶

Another useful data type built into Python is the *dictionary* (see Mapping Types — dict). Dictionaries are sometimes found in other languages as “associative memories” or “associative arrays”. Unlike sequences, which are indexed by a range of numbers, dictionaries are indexed by *keys*, which can be any immutable type; strings and numbers can always be keys. Tuples can be used as keys if they contain only strings, numbers, or tuples; if a tuple contains any mutable object either directly or indirectly, it cannot be used as a key. You can’t use lists as keys, since lists can be modified in place using index assignments, slice assignments, or methods like

```
append()
```

and

```
extend()
```

.

It is best to think of a dictionary as a set of *key: value* pairs, with the requirement that the keys are unique (within one dictionary). A pair of braces creates an empty dictionary:

```
{}
```

. Placing a comma-separated list of key:value pairs within the braces adds initial key:value pairs to the dictionary; this is also the way dictionaries are written on output.

The main operations on a dictionary are storing a value with some key and extracting the value given the key. It is also possible to delete a key:value pair with

```
del
```

. If you store using a key that is already in use, the old value associated with that key is forgotten.

Extracting a value for a non-existent key by subscripting (

```
d[key]
```

) raises a

```
KeyError
```

. To avoid getting this error when trying to access a possibly non-existent key, use the

```
get()
```

method instead, which returns

```
None
```

(or a specified default value) if the key is not in the dictionary.

Performing

```
list(d)
```

on a dictionary returns a list of all the keys used in the dictionary, in insertion order (if you want it sorted, just use

```
sorted(d)
```

instead). To check whether a single key is in the dictionary, use the

```
in
```

keyword.

Here is a small example using a dictionary:

Copy

```
>>> tel = {'jack': 4098, 'sape': 4139}
>>> tel['guido'] = 4127
>>> tel
{'jack': 4098, 'sape': 4139, 'guido': 4127}
>>> tel['jack']
4098
>>> tel['irv']
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
KeyError: 'irv'
>>> print(tel.get('irv'))
None
>>> del tel['sape']
>>> tel['irv'] = 4127
>>> tel
{'jack': 4098, 'guido': 4127, 'irv': 4127}
>>> list(tel)
['jack', 'guido', 'irv']
>>> sorted(tel)
['guido', 'irv', 'jack']
>>> 'guido' in tel
True
>>> 'jack' not in tel
False
```

The

```
dict()
```

constructor builds dictionaries directly from sequences of key-value pairs:

Copy

```
>>> dict([('sape', 4139), ('guido', 4127), ('jack', 4098)])
{'sape': 4139, 'guido': 4127, 'jack': 4098}
```

In addition, dict comprehensions can be used to create dictionaries from arbitrary key and value expressions:

Copy

```
>>> {x: x**2 for x in (2, 4, 6)}
{2: 4, 4: 16, 6: 36}
```

When the keys are simple strings, it is sometimes easier to specify pairs using keyword arguments:

Copy

```
>>> dict(sape=4139, guido=4127, jack=4098)
{'sape': 4139, 'guido': 4127, 'jack': 4098}
```

5.6. Looping Techniques¶

When looping through dictionaries, the key and corresponding value can be retrieved at the same time using the

```
items()
```

method.

Copy

```
>>> knights = {'gallahad': 'the pure', 'robin': 'the brave'}
>>> for k, v in knights.items():
...     print(k, v)
...
gallahad the pure
robin the brave
```

When looping through a sequence, the position index and corresponding value can be retrieved at the same time using the

```
enumerate()
```

function.

Copy

```
>>> for i, v in enumerate(['tic', 'tac', 'toe']):
...     print(i, v)
...
0 tic
1 tac
2 toe
```

To loop over two or more sequences at the same time, the entries can be paired with the

```
zip()
```

function.

Copy

```
>>> questions = ['name', 'quest', 'favorite color']
>>> answers = ['lancelot', 'the holy grail', 'blue']
>>> for q, a in zip(questions, answers):
...     print('What is your {0}? It is {1}.'.format(q, a))
...
What is your name? It is lancelot.
What is your quest? It is the holy grail.
What is your favorite color? It is blue.
```

To loop over a sequence in reverse, first specify the sequence in a forward direction and then call the

```
reversed()
```

function.

Copy

```
>>> for i in reversed(range(1, 10, 2)):
...     print(i)
...
9
7
5
3
1
```

To loop over a sequence in sorted order, use the

```
sorted()
```

function which returns a new sorted list while leaving the source unaltered.

Copy

```
>>> basket = ['apple', 'orange', 'apple', 'pear', 'orange', 'banana']
>>> for i in sorted(basket):
...     print(i)
...
apple
apple
banana
orange
orange
pear
```

Using

```
set()
```

on a sequence eliminates duplicate elements. The use of

```
sorted()
```

in combination with

```
set()
```

over a sequence is an idiomatic way to loop over unique elements of the sequence in sorted order.

Copy

```
>>> basket = ['apple', 'orange', 'apple', 'pear', 'orange', 'banana']
>>> for f in sorted(set(basket)):
...     print(f)
...
apple
banana
orange
pear
```

It is sometimes tempting to change a list while you are looping over it; however, it is often simpler and safer to create a new list instead.

Copy

```
>>> import math
>>> raw_data = [56.2, float('NaN'), 51.7, 55.3, 52.5, float('NaN'), 47.8]
>>> filtered_data = []
>>> for value in raw_data:
...     if not math.isnan(value):
...         filtered_data.append(value)
...
>>> filtered_data
[56.2, 51.7, 55.3, 52.5, 47.8]
```

5.7. More on Conditions¶

The conditions used in

```
while
```

and

```
if
```

statements can contain any operators, not just comparisons.

The comparison operators

```
in
```

and

```
not in
```

are membership tests that determine whether a value is in (or not in) a container. The operators

```
is
```

and

```
is not
```

compare whether two objects are really the same object. All comparison operators have the same priority, which is lower than that of all numerical operators.

Comparisons can be chained. For example,

```
a < b == c
```

tests whether

```
a
```

is less than

```
b
```


and moreover

b

equals

c

.

Comparisons may be combined using the Boolean operators

and

and

or

, and the outcome of a comparison (or of any other Boolean expression) may be negated with

not

. These have lower priorities than comparison operators; between them,

not

has the highest priority and

or

the lowest, so that

A and

not B or C

is equivalent to

(A and (not B)) or C

. As always, parentheses can be used to express the desired composition.

The Boolean operators

and

and

or

are so-called *short-circuit* operators: their arguments are evaluated from left to right, and evaluation stops as soon as the outcome is determined. For example, if

A

and

C

are true but

B

is false,

A and B and C

does not evaluate the expression

C

. When used as a general value and not as a Boolean, the return value of a short-circuit operator is the last evaluated argument.

It is possible to assign the result of a comparison or other Boolean expression to a variable. For example,

Copy

```
>>> string1, string2, string3 = '', 'Trondheim', 'Hammer Dance'
>>> non_null = string1 or string2 or string3
>>> non_null
'Trondheim'
```

Note that in Python, unlike C, assignment inside expressions must be done explicitly with the walrus operator

:=

. This avoids a common class of problems encountered in C programs: typing

=

in an expression when

==

was intended.

5.8. Comparing Sequences and Other Types¶

Sequence objects typically may be compared to other objects with the same sequence type. The comparison uses *lexicographical* ordering: first the first two items are compared, and if they differ this determines the outcome of the comparison; if they are equal, the next two items are compared, and so on, until either sequence is exhausted. If two items to be compared are themselves sequences of the same type, the lexicographical comparison is carried out recursively. If all items of two sequences compare equal, the sequences are considered equal. If one sequence is an initial sub-sequence of the other, the shorter sequence is the smaller (lesser) one. Lexicographical ordering for strings uses the Unicode code point number to order individual characters. Some examples of comparisons between sequences of the same type:

Copy

```
(1, 2, 3) < (1, 2, 4)
[1, 2, 3] < [1, 2, 4]
'ABC' < 'C' < 'Pascal' < 'Python'
(1, 2, 3, 4) < (1, 2, 4)
(1, 2) < (1, 2, -1)
(1, 2, 3) == (1.0, 2.0, 3.0)
(1, 2, ('aa', 'ab')) < (1, 2, ('abc', 'a'), 4)
```

Note that comparing objects of different types with

<

or

>

is legal provided that the objects have appropriate comparison methods. For example, mixed numeric types are compared according to their numeric value, so 0 equals 0.0, etc. Otherwise, rather than providing an arbitrary ordering, the interpreter will raise a

TypeError

exception.

Footnotes

[1]

Other languages may return the mutated object, which allows method chaining, such as

```
d->insert("a")->remove("b")->sort();
```

6. Modules¶

If you quit from the Python interpreter and enter it again, the definitions you have made (functions and variables) are lost. Therefore, if you want to write a somewhat longer program, you are better off using a text editor to prepare the input for the interpreter and running it with that file as input instead. This is known as creating a *script*. As your program gets longer, you may want to split it into several files for easier maintenance. You may also want to use a handy function that you've written in several programs without copying its definition into each program.

To support this, Python has a way to put definitions in a file and use them in a script or in an interactive instance of the interpreter. Such a file is called a *module*; definitions from a module can be *imported* into other modules or into the *main* module (the collection of variables that you have access to in a script executed at the top level and in calculator mode).

A module is a file containing Python definitions and statements. The file name is the module name with the suffix

```
.py
```

appended. Within a module, the module's name (as a string) is available as the value of the global variable

```
__name__
```

. For instance, use your favorite text editor to create a file called

```
fibonacci.py
```

in the current directory with the following contents:

```
Copy
```

```
# Fibonacci numbers module

def fib(n):
    """Write Fibonacci series up to n."""
    a, b = 0, 1
    while a < n:
        print(a, end=' ')
        a, b = b, a+b
    print()

def fib2(n):
    """Return Fibonacci series up to n."""
    result = []
    a, b = 0, 1
    while a < n:
        result.append(a)
        a, b = b, a+b
    return result
```

Now enter the Python interpreter and import this module with the following command:

Copy

```
>>> import fibo
```

This does not add the names of the functions defined in

```
fibo
```

directly to the current namespace (see Python Scopes and Namespaces for more details); it only adds the module name

```
fibo
```

there. Using the module name you can access the functions:

Copy

```
>>> fibo.fib(1000)
0 1 1 2 3 5 8 13 21 34 55 89 144 233 377 610 987
>>> fibo.fib2(100)
[0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89]
>>> fibo.__name__
'fibo'
```

If you intend to use a function often you can assign it to a local name:

Copy

```
>>> fib = fibo.fib
>>> fib(500)
0 1 1 2 3 5 8 13 21 34 55 89 144 233 377
```

6.1. More on Modules¶

A module can contain executable statements as well as function definitions. These statements are intended to initialize the module. They are executed only the *first* time the module name is encountered in an import statement. [\[1\]](#) (They are also run if the file is executed as a script.)

Each module has its own private namespace, which is used as the global namespace by all functions defined in the module. Thus, the author of a module can use global variables in the module without worrying about accidental clashes with a user's global variables. On the other hand, if you know what you are doing you can touch a module's global variables with the same notation used to refer to its functions,

```
modname.itemname
```

.

Modules can import other modules. It is customary but not required to place all

```
import
```

statements at the beginning of a module (or script, for that matter). The imported module names, if placed at the top level of a module (outside any functions or classes), are added to the module's global namespace.

There is a variant of the

```
import
```

statement that imports names from a module directly into the importing module's namespace. For example:

Copy

```
>>> from fibo import fib, fib2
>>> fib(500)
0 1 1 2 3 5 8 13 21 34 55 89 144 233 377
```

This does not introduce the module name from which the imports are taken in the local namespace (so in the example,

```
fibo
```

is not defined).

There is even a variant to import all names that a module defines:

Copy

```
>>> from fibo import *
>>> fib(500)
0 1 1 2 3 5 8 13 21 34 55 89 144 233 377
```

This imports all names except those beginning with an underscore (

—

). In most cases Python programmers do not use this facility since it introduces an unknown set of names into the interpreter, possibly hiding some things you have already defined.

Note that in general the practice of importing

*

from a module or package is frowned upon, since it often causes poorly readable code. However, it is okay to use it to save typing in interactive sessions.

If the module name is followed by

as

, then the name following

as

is bound directly to the imported module.

Copy

```
>>> import fibo as fib
>>> fib.fib(500)
0 1 1 2 3 5 8 13 21 34 55 89 144 233 377
```

This is effectively importing the module in the same way that

import fibo

will do, with the only difference of it being available as

fib

It can also be used when utilising

```
from
```

with similar effects:

Copy

```
>>> from fibo import fib as fibonacci
>>> fibonacci(500)
0 1 1 2 3 5 8 13 21 34 55 89 144 233 377
```

Note

For efficiency reasons, each module is only imported once per interpreter session. Therefore, if you change your modules, you must restart the interpreter – or, if it's just one module you want to test interactively, use

```
importlib.reload()
```

, e.g.

```
import importlib;
importlib.reload(modulename)
```

6.1.1. Executing modules as scripts

When you run a Python module with

Copy

```
python fibo.py <arguments>
```

the code in the module will be executed, just as if you imported it, but with the

```
__name__
```

set to

```
"__main__"
```

. That means that by adding this code at the end of your module:

Copy

```
if __name__ == "__main__":
    import sys
    fib(int(sys.argv[1]))
```


you can make the file usable as a script as well as an importable module, because the code that parses the command line only runs if the module is executed as the “main” file:

Copy

```
$ python fibo.py 50
0 1 1 2 3 5 8 13 21 34
```

If the module is imported, the code is not run:

Copy

```
>>> import fibo
>>>
```

This is often used either to provide a convenient user interface to a module, or for testing purposes (running the module as a script executes a test suite).

6.1.2. The Module Search Path

When a module named

```
spam
```

is imported, the interpreter first searches for a built-in module with that name. These module names are listed in

```
sys.builtin_module_names
```

. If not found, it then searches for a file named

```
spam.py
```

in a list of directories given by the variable

```
sys.path
```

.

```
sys.path
```

is initialized from these locations:

- The directory containing the input script (or the current directory when no file is specified).

- PYTHONPATH

(a list of directory names, with the same syntax as the shell variable

```
PATH
```

).

- The installation-dependent default (by convention including a

```
site-packages
```

directory, handled by the

```
site
```

module).

More details are at [The initialization of the sys.path module search path](#).

Note

On file systems which support symlinks, the directory containing the input script is calculated after the symlink is followed. In other words the directory containing the symlink is **not** added to the module search path.

After initialization, Python programs can modify

```
sys.path
```

. The directory containing the script being run is placed at the beginning of the search path, ahead of the standard library path. This means that scripts in that directory will be loaded instead of modules of the same name in the library directory. This is an error unless the replacement is intended. See section [Standard Modules](#) for more information.

6.1.3. “Compiled” Python files¶

To speed up loading modules, Python caches the compiled version of each module in the

```
__pycache__
```

directory under the name

```
module.version.pyc
```

, where the version encodes the format of the compiled file; it generally contains the Python version number. For example, in CPython release 3.3 the compiled version of spam.py would be cached as

```
__pycache__/spam.cpython-33.pyc
```

. This naming convention allows compiled modules from different releases and different versions of Python to coexist.

Python checks the modification date of the source against the compiled version to see if it's out of date and needs to be recompiled. This is a completely automatic process. Also, the compiled modules are platform-independent, so the same library can be shared among systems with different architectures.

Python does not check the cache in two circumstances. First, it always recompiles and does not store the result for the module that's loaded directly from the command line. Second, it does not

check the cache if there is no source module. To support a non-source (compiled only) distribution, the compiled module must be in the source directory, and there must not be a source module.

Some tips for experts:

- You can use the

```
-O
```

or

```
-OO
```

switches on the Python command to reduce the size of a compiled module. The

```
-O
```

switch removes assert statements, the

```
-OO
```

switch removes both assert statements and `__doc__` strings. Since some programs may rely on having these available, you should only use this option if you know what you're doing.

“Optimized” modules have an

```
opt-
```

tag and are usually smaller. Future releases may change the effects of optimization.

- A program doesn't run any faster when it is read from a

```
.pyc
```

file than when it is read from a

```
.py
```

file; the only thing that's faster about

```
.pyc
```

files is the speed with which they are loaded.

- The module

```
compileall
```

can create .pyc files for all modules in a directory.

- There is more detail on this process, including a flow chart of the decisions, in [PEP 3147](#).

6.2. Standard Modules¶

Python comes with a library of standard modules, described in a separate document, the Python Library Reference (“Library Reference” hereafter). Some modules are built into the interpreter;

these provide access to operations that are not part of the core of the language but are nevertheless built in, either for efficiency or to provide access to operating system primitives such as system calls. The set of such modules is a configuration option which also depends on the underlying platform. For example, the

```
winreg
```

module is only provided on Windows systems. One particular module deserves some attention:

```
sys
```

, which is built into every Python interpreter. The variables

```
sys.ps1
```

and

```
sys.ps2
```

define the strings used as primary and secondary prompts:

Copy

```
>>> import sys
>>> sys.ps1
'>>> '
>>> sys.ps2
'... '
>>> sys.ps1 = 'C> '
C> print('Yuck!')
Yuck!
C>
```

These two variables are only defined if the interpreter is in interactive mode.

The variable

```
sys.path
```

is a list of strings that determines the interpreter's search path for modules. It is initialized to a default path taken from the environment variable

```
PYTHONPATH
```

, or from a built-in default if

```
PYTHONPATH
```

is not set. You can modify it using standard list operations:

Copy

```
>>> import sys
>>> sys.path.append('/ufs/guido/lib/python')
```

6.3. The

dir()

Function¶

The built-in function

dir()

is used to find out which names a module defines. It returns a sorted list of strings:

Copy

```
>>> import fibo, sys
>>> dir(fibo)
['__name__', 'fib', 'fib2']
>>> dir(sys)
['__breakpointhook__', '__displayhook__', '__doc__', '__excepthook__',
 '__interactivehook__', '__loader__', '__name__', '__package__', '__spec__',
 '__stderr__', '__stdin__', '__stdout__', '__unraisablehook__',
 '_clear_type_cache', '_current_frames', '_debugmallocstats', '_framework',
 '_getframe', '_git', '_home', '_xoptions', 'abiflags', 'addaudithook',
 'api_version', 'argv', 'audit', 'base_exec_prefix', 'base_prefix',
 'breakpointhook', 'builtin_module_names', 'byteorder', 'call_tracing',
 'callstats', 'copyright', 'displayhook', 'dont_write_bytecode', 'exc_info',
 'excepthook', 'exec_prefix', 'executable', 'exit', 'flags', 'float_info',
 'float_repr_style', 'get_asyncgen_hooks', 'get_coroutine_origin_tracking_dep
th',
 'getallocatedblocks', 'getdefaultencoding', 'getdlopenflags',
 'getfilesystemencodeerrors', 'getfilesystemencoding', 'getprofile',
 'getrecursionlimit', 'getrefcount', 'getsizeof', 'getswitchinterval',
 'gettrace', 'hash_info', 'hexversion', 'implementation', 'int_info',
 'intern', 'is_finalizing', 'last_traceback', 'last_type', 'last_value',
 'maxsize', 'maxunicode', 'meta_path', 'modules', 'path', 'path_hooks',
 'path_importer_cache', 'platform', 'prefix', 'ps1', 'ps2', 'pycache_prefix',
 'set_asyncgen_hooks', 'set_coroutine_origin_tracking_depth', 'setdlopenflag
s',
 'setprofile', 'setrecursionlimit', 'setswitchinterval', 'settrace', 'stderr',
 'stdin', 'stdout', 'thread_info', 'unraisablehook', 'version', 'version_info',
 'warnoptions']
```

Without arguments,

```
dir()
```

lists the names you have defined currently:

Copy

```
>>> a = [1, 2, 3, 4, 5]
>>> import fibo
>>> fib = fibo.fib
>>> dir()
['__builtins__', '__name__', 'a', 'fib', 'fibo', 'sys']
```

Note that it lists all types of names: variables, modules, functions, etc.

```
dir()
```

does not list the names of built-in functions and variables. If you want a list of those, they are defined in the standard module

```
builtins
```

```
:
```

Copy

```
>>> import builtins
>>> dir(builtins)
['ArithmeticError', 'AssertionError', 'AttributeError', 'BaseException',
 'BlockingIOError', 'BrokenPipeError', 'BufferError', 'BytesWarning',
 'ChildProcessError', 'ConnectionAbortedError', 'ConnectionError',
 'ConnectionRefusedError', 'ConnectionResetError', 'DeprecationWarning',
 'EOFError', 'Ellipsis', 'EnvironmentError', 'Exception', 'False',
 'FileExistsError', 'FileNotFoundError', 'FloatingPointError',
 'FutureWarning', 'GeneratorExit', 'IOError', 'ImportError',
 'ImportWarning', 'IndentationError', 'IndexError', 'InterruptedError',
 'IsADirectoryError', 'KeyError', 'KeyboardInterrupt', 'LookupError',
 'MemoryError', 'NameError', 'None', 'NotADirectoryError', 'NotImplemented',
 'NotImplementedError', 'OSError', 'OverflowError',
 'PendingDeprecationWarning', 'PermissionError', 'ProcessLookupError',
 'ReferenceError', 'ResourceWarning', 'RuntimeError', 'RuntimeWarning',
 'StopIteration', 'SyntaxError', 'SyntaxWarning', 'SystemError',
 'SystemExit', 'TabError', 'TimeoutError', 'True', 'TypeError',
 'UnboundLocalError', 'UnicodeDecodeError', 'UnicodeEncodeError',
 'UnicodeError', 'UnicodeTranslateError', 'UnicodeWarning', 'UserWarning',
 'ValueError', 'Warning', 'ZeroDivisionError', '_', '__build_class__',
 '__debug__', '__doc__', '__import__', '__name__', '__package__', 'abs',
 'all', 'any', 'ascii', 'bin', 'bool', 'bytearray', 'bytes', 'callable',
 'chr', 'classmethod', 'compile', 'complex', 'copyright', 'credits',
 'delattr', 'dict', 'dir', 'divmod', 'enumerate', 'eval', 'exec', 'exit',
 'filter', 'float', 'format', 'frozenset', 'getattr', 'globals', 'hasattr',
 'hash', 'help', 'hex', 'id', 'input', 'int', 'isinstance', 'issubclass',
 'iter', 'len', 'license', 'list', 'locals', 'map', 'max', 'memoryview',
 'min', 'next', 'object', 'oct', 'open', 'ord', 'pow', 'print', 'property',
 'quit', 'range', 'repr', 'reversed', 'round', 'set', 'setattr', 'slice',
 'sorted', 'staticmethod', 'str', 'sum', 'super', 'tuple', 'type', 'vars',
 'zip']
```

6.4. Packages¶

Packages are a way of structuring Python’s module namespace by using “dotted module names”. For example, the module name

```
A.B
```

designates a submodule named

```
B
```

in a package named

```
A
```

. Just like the use of modules saves the authors of different modules from having to worry about each other’s global variable names, the use of dotted module names saves the authors of multi-module packages like NumPy or Pillow from having to worry about each other’s module names.

Suppose you want to design a collection of modules (a “package”) for the uniform handling of sound files and sound data. There are many different sound file formats (usually recognized by their extension, for example:

```
.wav
```

```
,
```

```
.aiff
```

```
,
```

```
.au
```

), so you may need to create and maintain a growing collection of modules for the conversion between the various file formats. There are also many different operations you might want to perform on sound data (such as mixing, adding echo, applying an equalizer function, creating an artificial stereo effect), so in addition you will be writing a never-ending stream of modules to perform these operations. Here’s a possible structure for your package (expressed in terms of a hierarchical filesystem):

sound/	Top-level package
__init__.py	Initialize the sound package
formats/	Subpackage for file format conversions
__init__.py	
wavread.py	
wavwrite.py	
aiffread.py	
aiffwrite.py	
auread.py	
auwrite.py	
...	
effects/	Subpackage for sound effects
__init__.py	
echo.py	
surround.py	
reverse.py	
...	
filters/	Subpackage for filters
__init__.py	
equalizer.py	
vocoder.py	
karaoke.py	
...	

When importing the package, Python searches through the directories on

```
sys.path
```

looking for the package subdirectory.

The

```
__init__.py
```

files are required to make Python treat directories containing the file as packages (unless using a namespace package, a relatively advanced feature). This prevents directories with a common name, such as

```
string
```

, from unintentionally hiding valid modules that occur later on the module search path. In the simplest case,

```
__init__.py
```

can just be an empty file, but it can also execute initialization code for the package or set the

```
__all__
```

variable, described later.

Users of the package can import individual modules from the package, for example:

Copy

```
import sound.effects.echo
```

This loads the submodule

```
sound.effects.echo
```

. It must be referenced with its full name.

Copy

```
sound.effects.echo.echofilter(input, output, delay=0.7, atten=4)
```

An alternative way of importing the submodule is:

Copy

```
from sound.effects import echo
```

This also loads the submodule

```
echo
```

, and makes it available without its package prefix, so it can be used as follows:

Copy

```
echo.echofilter(input, output, delay=0.7, atten=4)
```

Yet another variation is to import the desired function or variable directly:

Copy

```
from sound.effects.echo import echofilter
```

Again, this loads the submodule

```
echo
```

, but this makes its function

```
echofilter()
```

directly available:

Copy

```
echofilter(input, output, delay=0.7, atten=4)
```

Note that when using

```
from package import item
```

, the item can be either a submodule (or subpackage) of the package, or some other name defined in the package, like a function, class or variable. The

```
import
```

statement first tests whether the item is defined in the package; if not, it assumes it is a module and attempts to load it. If it fails to find it, an

```
ImportError
```

exception is raised.

Contrarily, when using syntax like

```
import item.subitem.subsubitem
```

, each item except for the last must be a package; the last item can be a module or a package but can't be a class or function or variable defined in the previous item.

6.4.1. Importing * From a Package¶

Now what happens when the user writes

```
from sound.effects import *
```

? Ideally, one would hope that this somehow goes out to the filesystem, finds which submodules are present in the package, and imports them all. This could take a long time and importing submodules might have unwanted side-effects that should only happen when the sub-module is explicitly imported.

The only solution is for the package author to provide an explicit index of the package. The

```
import
```

statement uses the following convention: if a package's

```
__init__.py
```

code defines a list named

```
__all__
```

, it is taken to be the list of module names that should be imported when

```
from package import *
```

is encountered. It is up to the package author to keep this list up-to-date when a new version of the package is released. Package authors may also decide not to support it, if they don't see a use for importing `*` from their package. For example, the file

```
sound/effects/__init__.py
```

could contain the following code:

Copy

```
__all__ = ["echo", "surround", "reverse"]
```

This would mean that

```
from sound.effects import *
```

would import the three named submodules of the

```
sound.effects
```

package.

Be aware that submodules might become shadowed by locally defined names. For example, if you added a

```
reverse
```

function to the

```
sound/effects/__init__.py
```

file, the

```
from sound.effects import *
```

would only import the two submodules

```
echo
```

and

```
surround
```

, but *not* the

```
reverse
```

submodule, because it is shadowed by the locally defined

```
reverse
```

function:

Copy

```
__all__ = [
    "echo",      # refers to the 'echo.py' file
    "surround",  # refers to the 'surround.py' file
    "reverse",   # !!! refers to the 'reverse' function now !!!
]

def reverse(msg: str): # <-- this name shadows the 'reverse.py' submodule
    return msg[::-1]   #      in the case of a 'from sound.effects import *'
```

If

```
__all__
```

is not defined, the statement

```
from sound.effects import *
```

does *not* import all submodules from the package

```
sound.effects
```

into the current namespace; it only ensures that the package

```
sound.effects
```

has been imported (possibly running any initialization code in

```
__init__.py
```

) and then imports whatever names are defined in the package. This includes any names defined (and submodules explicitly loaded) by

```
__init__.py
```

. It also includes any submodules of the package that were explicitly loaded by previous

```
import
```

statements. Consider this code:

Copy

```
import sound.effects.echo
import sound.effects.surround
from sound.effects import *
```

In this example, the

```
echo
```

and

```
surround
```

modules are imported in the current namespace because they are defined in the

```
sound.effects
```

package when the

```
from...import
```

statement is executed. (This also works when

```
__all__
```

is defined.)

Although certain modules are designed to export only names that follow certain patterns when you use

```
import *
```

, it is still considered bad practice in production code.

Remember, there is nothing wrong with using

```
from package import  
specific_submodule
```

! In fact, this is the recommended notation unless the importing module needs to use submodules with the same name from different packages.

6.4.2. Intra-package References¶

When packages are structured into subpackages (as with the

```
sound
```

package in the example), you can use absolute imports to refer to submodules of siblings packages. For example, if the module

```
sound.filters.vocoder
```

needs to use the

```
echo
```

module in the

```
sound.effects
```

package, it can use

```
from  
sound.effects import echo
```

.

You can also write relative imports, with the

```
from module import name
```

form of import statement. These imports use leading dots to indicate the current and parent packages involved in the relative import. From the

```
surround
```

module for example, you might use:

Copy

```
from . import echo
from .. import formats
from ..filters import equalizer
```

Note that relative imports are based on the name of the current module's package. Since the main module does not have a package, modules intended for use as the main module of a Python application must always use absolute imports.

6.4.3. Packages in Multiple Directories^[1]

Packages support one more special attribute,

```
__path__
```

. This is initialized to be a sequence of strings containing the name of the directory holding the package's

```
__init__.py
```

before the code in that file is executed. This variable can be modified; doing so affects future searches for modules and subpackages contained in the package.

While this feature is not often needed, it can be used to extend the set of modules found in a package.

Footnotes

[\[1\]](#)

In fact function definitions are also 'statements' that are 'executed'; the execution of a module-level function definition adds the function name to the module's global namespace.

7. Input and Output¶

There are several ways to present the output of a program; data can be printed in a human-readable form, or written to a file for future use. This chapter will discuss some of the possibilities.

7.1. Fancier Output Formatting¶

So far we've encountered two ways of writing values: *expression statements* and the

```
print()
```

function. (A third way is using the

```
write()
```

method of file objects; the standard output file can be referenced as

```
sys.stdout
```

. See the Library Reference for more information on this.)

Often you'll want more control over the formatting of your output than simply printing space-separated values. There are several ways to format output.

- To use [formatted string literals](#), begin a string with

```
f
```

or

```
F
```

before the opening quotation mark or triple quotation mark. Inside this string, you can write a Python expression between

```
{
```

and

```
}
```

characters that can refer to variables or literal values.

Copy

```
>>> year = 2016
>>> event = 'Referendum'
>>> f'Results of the {year} {event}'
'Results of the 2016 Referendum'
```


- The

```
str.format()
```

method of strings requires more manual effort. You'll still use

```
{
```

and

```
}
```

to mark where a variable will be substituted and can provide detailed formatting directives, but you'll also need to provide the information to be formatted. In the following code block there are two examples of how to format variables:

Copy

```
>>> yes_votes = 42_572_654
>>> total_votes = 85_705_149
>>> percentage = yes_votes / total_votes
>>> '{:-9} YES votes  {:.2%}'.format(yes_votes, percentage)
' 42572654 YES votes  49.67%'
```

Notice how the

```
yes_votes
```

are padded with spaces and a negative sign only for negative numbers. The example also prints

```
percentage
```

multiplied by 100, with 2 decimal places and followed by a percent sign (see Format Specification Mini-Language for details).

- Finally, you can do all the string handling yourself by using string slicing and concatenation operations to create any layout you can imagine. The string type has some methods that perform useful operations for padding strings to a given column width.

When you don't need fancy output but just want a quick display of some variables for debugging purposes, you can convert any value to a string with the

```
repr()
```

or

```
str()
```

functions.

The

```
str()
```

function is meant to return representations of values which are fairly human-readable, while

```
repr()
```

is meant to generate representations which can be read by the interpreter (or will force a

```
SyntaxError
```

if there is no equivalent syntax). For objects which don't have a particular representation for human consumption,

```
str()
```

will return the same value as

```
repr()
```

. Many values, such as numbers or structures like lists and dictionaries, have the same representation using either function. Strings, in particular, have two distinct representations.

Some examples:

Copy

```
>>> s = 'Hello, world.'
>>> str(s)
'Hello, world.'
>>> repr(s)
"'Hello, world.'"
>>> str(1/7)
'0.14285714285714285'
>>> x = 10 * 3.25
>>> y = 200 * 200
>>> s = 'The value of x is ' + repr(x) + ', and y is ' + repr(y) + '...'
>>> print(s)
The value of x is 32.5, and y is 40000...
>>> # The repr() of a string adds string quotes and backslashes:
>>> hello = 'hello, world\n'
>>> hellos = repr(hello)
>>> print(hellos)
'hello, world\n'
>>> # The argument to repr() may be any Python object:
>>> repr((x, y, ('spam', 'eggs'))))
"(32.5, 40000, ('spam', 'eggs'))"
```

The

```
string
```

module contains support for a simple templating approach based upon regular expressions, via

```
string.Template
```

. This offers yet another way to substitute values into strings, using placeholders like

```
$x
```

and replacing them with values from a dictionary. This syntax is easy to use, although it offers much less control for formatting.

7.1.1. Formatted String Literals¶

Formatted string literals (also called f-strings for short) let you include the value of Python expressions inside a string by prefixing the string with

```
f
```

or

```
F
```

and writing expressions as

```
{expression}
```

.

An optional format specifier can follow the expression. This allows greater control over how the value is formatted. The following example rounds pi to three places after the decimal:

Copy

```
>>> import math
>>> print(f'The value of pi is approximately {math.pi:.3f}.')
The value of pi is approximately 3.142.
```

Passing an integer after the

```
':'
```

will cause that field to be a minimum number of characters wide. This is useful for making columns line up.

Copy

```
>>> table = {'Sjoerd': 4127, 'Jack': 4098, 'Dcab': 7678}
>>> for name, phone in table.items():
...     print(f'{name:10} ==> {phone:10d}')
...
Sjoerd      ==>      4127
Jack        ==>      4098
Dcab        ==>      7678
```

Other modifiers can be used to convert the value before it is formatted.

```
'!a'
```

applies

```
ascii()
```

```
,
```

```
'!s'
```

applies

```
str()
```

, and

```
'!r'
```

applies

```
repr()
```

```
:
```

Copy

```
>>> animals = 'eels'
>>> print(f'My hovercraft is full of {animals}.')
My hovercraft is full of eels.
>>> print(f'My hovercraft is full of {animals!r}.')
My hovercraft is full of 'eels'.
```

The

```
=
```

specifier can be used to expand an expression to the text of the expression, an equal sign, then the representation of the evaluated expression:

Copy

```
>>> bugs = 'roaches'
>>> count = 13
>>> area = 'living room'
>>> print(f'Debugging {bugs=} {count=} {area=}')
Debugging bugs='roaches' count=13 area='living room'
```

See self-documenting expressions for more information on the

```
=
```

specifier. For a reference on these format specifications, see the reference guide for the Format Specification Mini-Language.

7.1.2. The String `format()` Method

Basic usage of the

```
str.format()
```

method looks like this:

Copy

```
>>> print('We are the {} who say "{}!"'.format('knights', 'Ni'))
We are the knights who say "Ni!"
```

The brackets and characters within them (called format fields) are replaced with the objects passed into the

```
str.format()
```

method. A number in the brackets can be used to refer to the position of the object passed into the

```
str.format()
```

method.

Copy

```
>>> print('{0} and {1}'.format('spam', 'eggs'))
spam and eggs
>>> print('{1} and {0}'.format('spam', 'eggs'))
eggs and spam
```

If keyword arguments are used in the

```
str.format()
```

method, their values are referred to by using the name of the argument.

Copy

```
>>> print('This {food} is {adjective}.'.format(
...     food='spam', adjective='absolutely horrible'))
This spam is absolutely horrible.
```

Positional and keyword arguments can be arbitrarily combined:

Copy

```
>>> print('The story of {0}, {1}, and {other}.'.format('Bill', 'Manfred',
...                                                    other='Georg'))
The story of Bill, Manfred, and Georg.
```

If you have a really long format string that you don't want to split up, it would be nice if you could reference the variables to be formatted by name instead of by position. This can be done by simply passing the dict and using square brackets

```
'[]'
```

to access the keys.

Copy

```
>>> table = {'Sjoerd': 4127, 'Jack': 4098, 'Dcab': 8637678}
>>> print('Jack: {0[Jack]:d}; Sjoerd: {0[Sjoerd]:d}; '
...       'Dcab: {0[Dcab]:d}'.format(table))
Jack: 4098; Sjoerd: 4127; Dcab: 8637678
```

This could also be done by passing the

```
table
```

dictionary as keyword arguments with the

```
**
```

notation.

Copy

```
>>> table = {'Sjoerd': 4127, 'Jack': 4098, 'Dcab': 8637678}
>>> print('Jack: {Jack:d}; Sjoerd: {Sjoerd:d}; Dcab: {Dcab:d}'.format(**table))
Jack: 4098; Sjoerd: 4127; Dcab: 8637678
```

This is particularly useful in combination with the built-in function

```
vars()
```

, which returns a dictionary containing all local variables:

Copy

```
>>> table = {k: str(v) for k, v in vars().items()}
>>> message = " ".join([f'{k}: ' + '{' + k + '}';' for k in table.keys()])
>>> print(message.format(**table))
__name__: __main__; __doc__: None; __package__: None; __loader__: ...
```

As an example, the following lines produce a tidily aligned set of columns giving integers and their squares and cubes:

Copy

```
>>> for x in range(1, 11):
...     print('{0:2d} {1:3d} {2:4d}'.format(x, x*x, x*x*x))
...
1   1   1
2   4   8
3   9  27
4  16  64
5  25 125
6  36 216
7  49 343
8  64 512
9  81 729
10 100 1000
```

For a complete overview of string formatting with

```
str.format()
```

, see Format String Syntax.

7.1.3. Manual String Formatting¶

Here's the same table of squares and cubes, formatted manually:

Copy

```
>>> for x in range(1, 11):
...     print(repr(x).rjust(2), repr(x*x).rjust(3), end=' ')
...     # Note use of 'end' on previous line
...     print(repr(x*x*x).rjust(4))
...
1   1   1
2   4   8
3   9  27
4  16  64
5  25 125
6  36 216
7  49 343
8  64 512
9  81 729
10 100 1000
```

(Note that the one space between each column was added by the way

```
print()
```

works: it always adds spaces between its arguments.)

The

```
str.rjust()
```

method of string objects right-justifies a string in a field of a given width by padding it with spaces on the left. There are similar methods

```
str.ljust()
```

and

```
str.center()
```

. These methods do not write anything, they just return a new string. If the input string is too long, they don't truncate it, but return it unchanged; this will mess up your column lay-out but that's usually better than the alternative, which would be lying about a value. (If you really want truncation you can always add a slice operation, as in

```
x.ljust(n)[:n]
```

.)

There is another method,

```
str.zfill()
```

, which pads a numeric string on the left with zeros. It understands about plus and minus signs:

Copy


```
>>> '12'.zfill(5)
'00012'
>>> '-3.14'.zfill(7)
'-003.14'
>>> '3.14159265359'.zfill(5)
'3.14159265359'
```

7.1.4. Old string formatting¶

The % operator (modulo) can also be used for string formatting. Given

```
format % values
```

(where *format* is a string),

```
%
```

conversion specifications in *format* are replaced with zero or more elements of *values*. This operation is commonly known as string interpolation. For example:

Copy

```
>>> import math
>>> print('The value of pi is approximately %5.3f.' % math.pi)
The value of pi is approximately 3.142.
```

More information can be found in the printf-style String Formatting section.

7.2. Reading and Writing Files¶

```
open()
```

returns a file object, and is most commonly used with two positional arguments and one keyword argument:

```
open(filename, mode, encoding=None)
```

Copy

```
>>> f = open('workfile', 'w', encoding="utf-8")
```

The first argument is a string containing the filename. The second argument is another string containing a few characters describing the way in which the file will be used. *mode* can be

```
'r'
```

when the file will only be read,

```
'w'
```

for only writing (an existing file with the same name will be erased), and

```
'a'
```

opens the file for appending; any data written to the file is automatically added to the end.

```
'r+'
```

opens the file for both reading and writing. The *mode* argument is optional;

```
'r'
```

will be assumed if it's omitted.

Normally, files are opened in *text mode*, that means, you read and write strings from and to the file, which are encoded in a specific *encoding*. If *encoding* is not specified, the default is platform dependent (see

```
open()
```

). Because UTF-8 is the modern de-facto standard,

```
encoding="utf-8"
```

is recommended unless you know that you need to use a different encoding. Appending a

```
'b'
```

to the mode opens the file in *binary mode*. Binary mode data is read and written as

```
bytes
```

objects. You can not specify *encoding* when opening file in binary mode.

In text mode, the default when reading is to convert platform-specific line endings (

```
\n
```

on Unix,

```
\r\n
```

on Windows) to just

```
\n
```

. When writing in text mode, the default is to convert occurrences of

```
\n
```

back to platform-specific line endings. This behind-the-scenes modification to file data is fine for text files, but will corrupt binary data like that in

```
JPEG
```

or

EXE

files. Be very careful to use binary mode when reading and writing such files.

It is good practice to use the

`with`

keyword when dealing with file objects. The advantage is that the file is properly closed after its suite finishes, even if an exception is raised at some point. Using

`with`

is also much shorter than writing equivalent

`try`

-

`finally`

blocks:

Copy

```
>>> with open('workfile', encoding="utf-8") as f:
...     read_data = f.read()

>>> # We can check that the file has been automatically closed.
>>> f.closed
True
```

If you're not using the

`with`

keyword, then you should call

`f.close()`

to close the file and immediately free up any system resources used by it.

Warning

Calling

`f.write()`

without using the

`with`

keyword or calling

```
f.close()
```

might result in the arguments of

```
f.write()
```

not being completely written to the disk, even if the program exits successfully.

After a file object is closed, either by a

```
with
```

statement or by calling

```
f.close()
```

, attempts to use the file object will automatically fail.

Copy

```
>>> f.close()
>>> f.read()
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: I/O operation on closed file.
```

7.2.1. Methods of File Objects

The rest of the examples in this section will assume that a file object called

```
f
```

has already been created.

To read a file's contents, call

```
f.read(size)
```

, which reads some quantity of data and returns it as a string (in text mode) or bytes object (in binary mode). *size* is an optional numeric argument. When *size* is omitted or negative, the entire contents of the file will be read and returned; it's your problem if the file is twice as large as your machine's memory. Otherwise, at most *size* characters (in text mode) or *size* bytes (in binary mode) are read and returned. If the end of the file has been reached,

```
f.read()
```

will return an empty string (

```
''
```

).

Copy

```
>>> f.read()
'This is the entire file.\n'
>>> f.read()
''
```

```
f.readline()
```

reads a single line from the file; a newline character (

```
\n
```

) is left at the end of the string, and is only omitted on the last line of the file if the file doesn't end in a newline. This makes the return value unambiguous; if

```
f.readline()
```

returns an empty string, the end of the file has been reached, while a blank line is represented by

```
'\n'
```

, a string containing only a single newline.

Copy

```
>>> f.readline()
'This is the first line of the file.\n'
>>> f.readline()
'Second line of the file\n'
>>> f.readline()
''
```

For reading lines from a file, you can loop over the file object. This is memory efficient, fast, and leads to simple code:

Copy

```
>>> for line in f:
...     print(line, end='')
...
This is the first line of the file.
Second line of the file
```

If you want to read all the lines of a file in a list you can also use

```
list(f)
```

or

```
f.readlines()
```

.

```
f.write(string)
```

writes the contents of *string* to the file, returning the number of characters written.

Copy

```
>>> f.write('This is a test\n')
15
```

Other types of objects need to be converted – either to a string (in text mode) or a bytes object (in binary mode) – before writing them:

Copy

```
>>> value = ('the answer', 42)
>>> s = str(value) # convert the tuple to string
>>> f.write(s)
18
```

```
f.tell()
```

returns an integer giving the file object's current position in the file represented as number of bytes from the beginning of the file when in binary mode and an opaque number when in text mode.

To change the file object's position, use

```
f.seek(offset, whence)
```

. The position is computed from adding *offset* to a reference point; the reference point is selected by the *whence* argument. A *whence* value of 0 measures from the beginning of the file, 1 uses the current file position, and 2 uses the end of the file as the reference point. *whence* can be omitted and defaults to 0, using the beginning of the file as the reference point.

Copy

```

>>> f = open('workfile', 'rb+')
>>> f.write(b'0123456789abcdef')
16
>>> f.seek(5)      # Go to the 6th byte in the file
5
>>> f.read(1)
b'5'
>>> f.seek(-3, 2)  # Go to the 3rd byte before the end
13
>>> f.read(1)
b'd'

```

In text files (those opened without a

```
b
```

in the mode string), only seeks relative to the beginning of the file are allowed (the exception being seeking to the very file end with

```
seek(0, 2)
```

) and the only valid *offset* values are those returned from the

```
f.tell()
```

, or zero. Any other *offset* value produces undefined behaviour.

File objects have some additional methods, such as

```
isatty()
```

and

```
truncate()
```

which are less frequently used; consult the Library Reference for a complete guide to file objects.

7.2.2. Saving structured data with

```
json
```



Strings can easily be written to and read from a file. Numbers take a bit more effort, since the

```
read()
```

method only returns strings, which will have to be passed to a function like

```
int()
```

, which takes a string like

```
'123'
```

and returns its numeric value 123. When you want to save more complex data types like nested lists and dictionaries, parsing and serializing by hand becomes complicated.

Rather than having users constantly writing and debugging code to save complicated data types to files, Python allows you to use the popular data interchange format called [JSON \(JavaScript Object Notation\)](#). The standard module called

```
json
```

can take Python data hierarchies, and convert them to string representations; this process is called *serializing*. Reconstructing the data from the string representation is called *deserializing*. Between serializing and deserializing, the string representing the object may have been stored in a file or data, or sent over a network connection to some distant machine.

Note

The JSON format is commonly used by modern applications to allow for data exchange. Many programmers are already familiar with it, which makes it a good choice for interoperability.

If you have an object

```
x
```

, you can view its JSON string representation with a simple line of code:

Copy

```
>>> import json
>>> x = [1, 'simple', 'list']
>>> json.dumps(x)
'[1, "simple", "list"]'
```

Another variant of the

```
dumps()
```

function, called

```
dump()
```

, simply serializes the object to a text file. So if

```
f
```

is a text file object opened for writing, we can do this:

Copy

```
json.dump(x, f)
```


To decode the object again, if

```
f
```

is a binary file or text file object which has been opened for reading:

Copy

```
x = json.load(f)
```

Note

JSON files must be encoded in UTF-8. Use

```
encoding="utf-8"
```

when opening JSON file as a text file for both of reading and writing.

This simple serialization technique can handle lists and dictionaries, but serializing arbitrary class instances in JSON requires a bit of extra effort. The reference for the

```
json
```

module contains an explanation of this.

See also

```
pickle
```

- the pickle module

Contrary to [JSON](#), *pickle* is a protocol which allows the serialization of arbitrarily complex Python objects. As such, it is specific to Python and cannot be used to communicate with applications written in other languages. It is also insecure by default: deserializing pickle data coming from an untrusted source can execute arbitrary code, if the data was crafted by a skilled attacker.

8. Errors and Exceptions¶

Until now error messages haven't been more than mentioned, but if you have tried out the examples you have probably seen some. There are (at least) two distinguishable kinds of errors: *syntax errors* and *exceptions*.

8.1. Syntax Errors¶

Syntax errors, also known as parsing errors, are perhaps the most common kind of complaint you get while you are still learning Python:

Copy

```
>>> while True print('Hello world')
      File "<stdin>", line 1
        while True print('Hello world')
                        ^^^^^
SyntaxError: invalid syntax
```

The parser repeats the offending line and displays little arrows pointing at the place where the error was detected. Note that this is not always the place that needs to be fixed. In the example, the error is detected at the function

```
print()
```

, since a colon (

```
':'
```

) is missing just before it.

The file name (

```
<stdin>
```

in our example) and line number are printed so you know where to look in case the input came from a file.

8.2. Exceptions¶

Even if a statement or expression is syntactically correct, it may cause an error when an attempt is made to execute it. Errors detected during execution are called *exceptions* and are not unconditionally fatal: you will soon learn how to handle them in Python programs. Most exceptions are not handled by programs, however, and result in error messages as shown here:

Copy

```

>>> 10 * (1/0)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
    10 * (1/0)
    ~^^
ZeroDivisionError: division by zero
>>> 4 + spam*3
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
    4 + spam*3
    ^^^^
NameError: name 'spam' is not defined
>>> '2' + 2
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
    '2' + 2
    ~~~~^~~
TypeError: can only concatenate str (not "int") to str

```

The last line of the error message indicates what happened. Exceptions come in different types, and the type is printed as part of the message: the types in the example are

```
ZeroDivisionError
```

,

```
NameError
```

and

```
TypeError
```

. The string printed as the exception type is the name of the built-in exception that occurred. This is true for all built-in exceptions, but need not be true for user-defined exceptions (although it is a useful convention). Standard exception names are built-in identifiers (not reserved keywords).

The rest of the line provides detail based on the type of exception and what caused it.

The preceding part of the error message shows the context where the exception occurred, in the form of a stack traceback. In general it contains a stack traceback listing source lines; however, it will not display lines read from standard input.

Built-in Exceptions lists the built-in exceptions and their meanings.

8.3. Handling Exceptions¶

It is possible to write programs that handle selected exceptions. Look at the following example, which asks the user for input until a valid integer has been entered, but allows the user to interrupt the program (using Control-C or whatever the operating system supports); note that a user-generated interruption is signalled by raising the

KeyboardInterrupt

exception.

Copy

```
>>> while True:
...     try:
...         x = int(input("Please enter a number: "))
...         break
...     except ValueError:
...         print("Oops! That was no valid number. Try again...")
... 
```

The

try

statement works as follows.

- First, the *try clause* (the statement(s) between the

try

and

except

keywords) is executed.

- If no exception occurs, the *except clause* is skipped and execution of the

try

statement is finished.

- If an exception occurs during execution of the

try

clause, the rest of the clause is skipped. Then, if its type matches the exception named after the

except

keyword, the *except clause* is executed, and then execution continues after the try/except block.

- If an exception occurs which does not match the exception named in the *except clause*, it is passed on to outer

```
try
```

statements; if no handler is found, it is an *unhandled exception* and execution stops with an error message.

A

```
try
```

statement may have more than one *except clause*, to specify handlers for different exceptions. At most one handler will be executed. Handlers only handle exceptions that occur in the corresponding *try clause*, not in other handlers of the same

```
try
```

statement. An *except clause* may name multiple exceptions as a parenthesized tuple, for example:

Copy

```
... except (RuntimeError, TypeError, NameError):  
...     pass
```

A class in an

```
except
```

clause matches exceptions which are instances of the class itself or one of its derived classes (but not the other way around — an *except clause* listing a derived class does not match instances of its base classes). For example, the following code will print B, C, D in that order:

Copy

```

class B(Exception):
    pass

class C(B):
    pass

class D(C):
    pass

for cls in [B, C, D]:
    try:
        raise cls()
    except D:
        print("D")
    except C:
        print("C")
    except B:
        print("B")

```

Note that if the *except clauses* were reversed (with

```
except B
```

first), it would have printed B, B, B — the first matching *except clause* is triggered.

When an exception occurs, it may have associated values, also known as the exception's *arguments*. The presence and types of the arguments depend on the exception type.

The *except clause* may specify a variable after the exception name. The variable is bound to the exception instance which typically has an

```
args
```

attribute that stores the arguments. For convenience, builtin exception types define

```
__str__()
```

to print all the arguments without explicitly accessing

```
.args
```

.

Copy

```
>>> try:
...     raise Exception('spam', 'eggs')
... except Exception as inst:
...     print(type(inst))    # the exception type
...     print(inst.args)    # arguments stored in .args
...     print(inst)         # __str__ allows args to be printed directly,
...                           # but may be overridden in exception subclasses
...     x, y = inst.args    # unpack args
...     print('x =', x)
...     print('y =', y)
...
<class 'Exception'>
('spam', 'eggs')
('spam', 'eggs')
x = spam
y = eggs
```

The exception's

```
__str__()
```

output is printed as the last part ('detail') of the message for unhandled exceptions.

```
BaseException
```

is the common base class of all exceptions. One of its subclasses,

```
Exception
```

, is the base class of all the non-fatal exceptions. Exceptions which are not subclasses of

```
Exception
```

are not typically handled, because they are used to indicate that the program should terminate.

They include

```
SystemExit
```

which is raised by

```
sys.exit()
```

and

```
KeyboardInterrupt
```

which is raised when a user wishes to interrupt the program.

```
Exception
```

can be used as a wildcard that catches (almost) everything. However, it is good practice to be as

specific as possible with the types of exceptions that we intend to handle, and to allow any unexpected exceptions to propagate on.

The most common pattern for handling

Exception

is to print or log the exception and then re-raise it (allowing a caller to handle the exception as well):

Copy

```
import sys

try:
    f = open('myfile.txt')
    s = f.readline()
    i = int(s.strip())
except OSError as err:
    print("OS error:", err)
except ValueError:
    print("Could not convert data to an integer.")
except Exception as err:
    print(f"Unexpected {err=}, {type(err)=}")
    raise
```

The

try

...

except

statement has an optional *else clause*, which, when present, must follow all *except clauses*. It is useful for code that must be executed if the *try clause* does not raise an exception. For example:

Copy

```
for arg in sys.argv[1:]:
    try:
        f = open(arg, 'r')
    except OSError:
        print('cannot open', arg)
    else:
        print(arg, 'has', len(f.readlines()), 'lines')
        f.close()
```

The use of the


```
else
```

clause is better than adding additional code to the

```
try
```

clause because it avoids accidentally catching an exception that wasn't raised by the code being protected by the

```
try
```

```
...
```

```
except
```

statement.

Exception handlers do not handle only exceptions that occur immediately in the *try clause*, but also those that occur inside functions that are called (even indirectly) in the *try clause*. For example:

Copy

```
>>> def this_fails():
...     x = 1/0
...
>>> try:
...     this_fails()
... except ZeroDivisionError as err:
...     print('Handling run-time error:', err)
...
Handling run-time error: division by zero
```

8.4. Raising Exceptions¶

The

```
raise
```

statement allows the programmer to force a specified exception to occur. For example:

Copy

```
>>> raise NameError('HiThere')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
    raise NameError('HiThere')
NameError: HiThere
```

The sole argument to

```
raise
```

indicates the exception to be raised. This must be either an exception instance or an exception class (a class that derives from

```
BaseException
```

, such as

```
Exception
```

or one of its subclasses). If an exception class is passed, it will be implicitly instantiated by calling its constructor with no arguments:

Copy

```
raise ValueError # shorthand for 'raise ValueError()'
```

If you need to determine whether an exception was raised but don't intend to handle it, a simpler form of the

```
raise
```

statement allows you to re-raise the exception:

Copy

```
>>> try:
...     raise NameError('HiThere')
... except NameError:
...     print('An exception flew by!')
...     raise
...
An exception flew by!
Traceback (most recent call last):
  File "<stdin>", line 2, in <module>
    raise NameError('HiThere')
NameError: HiThere
```

8.5. Exception Chaining

If an unhandled exception occurs inside an

```
except
```

section, it will have the exception being handled attached to it and included in the error message:

Copy

```
>>> try:
...     open("database.sqlite")
... except OSError:
...     raise RuntimeError("unable to handle error")
...
Traceback (most recent call last):
  File "<stdin>", line 2, in <module>
    open("database.sqlite")
    ~~~~^~~~~~
FileNotFoundError: [Errno 2] No such file or directory: 'database.sqlite'

During handling of the above exception, another exception occurred:

Traceback (most recent call last):
  File "<stdin>", line 4, in <module>
    raise RuntimeError("unable to handle error")
RuntimeError: unable to handle error
```

To indicate that an exception is a direct consequence of another, the

`raise`

statement allows an optional

`from`

clause:

Copy

```
# exc must be exception instance or None.
raise RuntimeError from exc
```

This can be useful when you are transforming exceptions. For example:

Copy

```
>>> def func():
...     raise ConnectionError
...
>>> try:
...     func()
... except ConnectionError as exc:
...     raise RuntimeError('Failed to open database') from exc
...
Traceback (most recent call last):
  File "<stdin>", line 2, in <module>
    func()
    ~~~~^^
  File "<stdin>", line 2, in func
ConnectionError
```

The above exception was the direct cause of the following exception:

```
Traceback (most recent call last):
  File "<stdin>", line 4, in <module>
    raise RuntimeError('Failed to open database') from exc
RuntimeError: Failed to open database
```

It also allows disabling automatic exception chaining using the

```
from None
```

idiom:

Copy

```
>>> try:
...     open('database.sqlite')
... except OSError:
...     raise RuntimeError from None
...
Traceback (most recent call last):
  File "<stdin>", line 4, in <module>
    raise RuntimeError from None
RuntimeError
```

For more information about chaining mechanics, see [Built-in Exceptions](#).

8.6. User-defined Exceptions¶

Programs may name their own exceptions by creating a new exception class (see [Classes](#) for more about Python classes). Exceptions should typically be derived from the

```
Exception
```

class, either directly or indirectly.

Exception classes can be defined which do anything any other class can do, but are usually kept simple, often only offering a number of attributes that allow information about the error to be extracted by handlers for the exception.

Most exceptions are defined with names that end in “Error”, similar to the naming of the standard exceptions.

Many standard modules define their own exceptions to report errors that may occur in functions they define.

8.7. Defining Clean-up Actions¶

The

```
try
```

statement has another optional clause which is intended to define clean-up actions that must be executed under all circumstances. For example:

Copy

```
>>> try:
...     raise KeyboardInterrupt
... finally:
...     print('Goodbye, world!')
...
Goodbye, world!
Traceback (most recent call last):
  File "<stdin>", line 2, in <module>
    raise KeyboardInterrupt
KeyboardInterrupt
```

If a

```
finally
```

clause is present, the

```
finally
```

clause will execute as the last task before the

```
try
```

statement completes. The

```
finally
```

clause runs whether or not the

```
try
```

statement produces an exception. The following points discuss more complex cases when an exception occurs:

- If an exception occurs during execution of the

```
try
```

clause, the exception may be handled by an

```
except
```

clause. If the exception is not handled by an

```
except
```

clause, the exception is re-raised after the

```
finally
```

clause has been executed.

- An exception could occur during execution of an

```
except
```

or

```
else
```

clause. Again, the exception is re-raised after the

```
finally
```

clause has been executed.

- If the

```
finally
```

clause executes a

```
break
```

,

```
continue
```

or

```
return
```

statement, exceptions are not re-raised. This can be confusing and is therefore discouraged. From version 3.14 the compiler emits a

```
SyntaxWarning
```

for it (see [PEP 765](#)).

- If the

```
try
```

statement reaches a

```
break
```

,

```
continue
```

or

```
return
```

statement, the

```
finally
```

clause will execute just prior to the

```
break
```

,

```
continue
```

or

```
return
```

statement's execution.

- If a

```
finally
```

clause includes a

```
return
```

statement, the returned value will be the one from the

```
finally
```

clause's

```
return
```

statement, not the value from the

```
try
```

clause's

```
return
```

statement. This can be confusing and is therefore discouraged. From version 3.14 the compiler emits a

```
SyntaxWarning
```

for it (see [PEP 765](#)).

For example:

Copy

```
>>> def bool_return():
...     try:
...         return True
...     finally:
...         return False
...
>>> bool_return()
False
```

A more complicated example:

Copy


```

>>> def divide(x, y):
...     try:
...         result = x / y
...     except ZeroDivisionError:
...         print("division by zero!")
...     else:
...         print("result is", result)
...     finally:
...         print("executing finally clause")
...
>>> divide(2, 1)
result is 2.0
executing finally clause
>>> divide(2, 0)
division by zero!
executing finally clause
>>> divide("2", "1")
executing finally clause
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
    divide("2", "1")
    ~~~~~^~~~~~
  File "<stdin>", line 3, in divide
    result = x / y
            ~^^~
TypeError: unsupported operand type(s) for /: 'str' and 'str'

```

As you can see, the

```
finally
```

clause is executed in any event. The

```
TypeError
```

raised by dividing two strings is not handled by the

```
except
```

clause and therefore re-raised after the

```
finally
```

clause has been executed.

In real world applications, the

```
finally
```

clause is useful for releasing external resources (such as files or network connections), regardless of whether the use of the resource was successful.

8.8. Predefined Clean-up Actions¶

Some objects define standard clean-up actions to be undertaken when the object is no longer needed, regardless of whether or not the operation using the object succeeded or failed. Look at the following example, which tries to open a file and print its contents to the screen.

Copy

```
for line in open("myfile.txt"):
    print(line, end="")
```

The problem with this code is that it leaves the file open for an indeterminate amount of time after this part of the code has finished executing. This is not an issue in simple scripts, but can be a problem for larger applications. The

with

statement allows objects like files to be used in a way that ensures they are always cleaned up promptly and correctly.

Copy

```
with open("myfile.txt") as f:
    for line in f:
        print(line, end="")
```

After the statement is executed, the file *f* is always closed, even if a problem was encountered while processing the lines. Objects which, like files, provide predefined clean-up actions will indicate this in their documentation.

8.9. Raising and Handling Multiple Unrelated Exceptions¶

There are situations where it is necessary to report several exceptions that have occurred. This is often the case in concurrency frameworks, when several tasks may have failed in parallel, but there are also other use cases where it is desirable to continue execution and collect multiple errors rather than raise the first exception.

The builtin

ExceptionGroup

wraps a list of exception instances so that they can be raised together. It is an exception itself, so it can be caught like any other exception.

Copy

```
>>> def f():
...     excs = [OSError('error 1'), SystemError('error 2')]
...     raise ExceptionGroup('there were problems', excs)
...
>>> f()
+ Exception Group Traceback (most recent call last):
|   File "<stdin>", line 1, in <module>
|     f()
|     ~^^
|   File "<stdin>", line 3, in f
|     raise ExceptionGroup('there were problems', excs)
| ExceptionGroup: there were problems (2 sub-exceptions)
+--+----- 1 -----+
|   OSError: error 1
+----- 2 -----+
|   SystemError: error 2
+-----+
>>> try:
...     f()
... except Exception as e:
...     print(f'caught {type(e)}: e')
...
caught <class 'ExceptionGroup'>: e
>>>
```

By using

```
except*
```

instead of

```
except
```

, we can selectively handle only the exceptions in the group that match a certain type. In the following example, which shows a nested exception group, each

```
except*
```

clause extracts from the group exceptions of a certain type while letting all other exceptions propagate to other clauses and eventually to be reraised.

Copy

```

>>> def f():
...     raise ExceptionGroup(
...         "group1",
...         [
...             OSError(1),
...             SystemError(2),
...             ExceptionGroup(
...                 "group2",
...                 [
...                     OSError(3),
...                     RecursionError(4)
...                 ]
...             )
...         ]
...     )
...
>>> try:
...     f()
... except* OSError as e:
...     print("There were OSErrors")
... except* SystemError as e:
...     print("There were SystemErrors")
...
There were OSErrors
There were SystemErrors
+ Exception Group Traceback (most recent call last):
|   File "<stdin>", line 2, in <module>
|       f()
|       ~^^
|   File "<stdin>", line 2, in f
|       raise ExceptionGroup(
|         ...<12 lines>...
|     )
| ExceptionGroup: group1 (1 sub-exception)
+-+----- 1 -----
| ExceptionGroup: group2 (1 sub-exception)
+-+----- 1 -----
| RecursionError: 4
+-----
>>>

```

Note that the exceptions nested in an exception group must be instances, not types. This is because in practice the exceptions would typically be ones that have already been raised and

caught by the program, along the following pattern:

Copy

```
>>> excs = []
... for test in tests:
...     try:
...         test.run()
...     except Exception as e:
...         excs.append(e)
...
>>> if excs:
...     raise ExceptionGroup("Test Failures", excs)
...
```

8.10. Enriching Exceptions with Notes ¶

When an exception is created in order to be raised, it is usually initialized with information that describes the error that has occurred. There are cases where it is useful to add information after the exception was caught. For this purpose, exceptions have a method

```
add_note(note)
```

that accepts a string and adds it to the exception's notes list. The standard traceback rendering includes all notes, in the order they were added, after the exception.

Copy

```
>>> try:
...     raise TypeError('bad type')
... except Exception as e:
...     e.add_note('Add some information')
...     e.add_note('Add some more information')
...     raise
...
Traceback (most recent call last):
  File "<stdin>", line 2, in <module>
    raise TypeError('bad type')
TypeError: bad type
Add some information
Add some more information
>>>
```

For example, when collecting exceptions into an exception group, we may want to add context information for the individual errors. In the following each exception in the group has a note

indicating when this error has occurred.

Copy

```

>>> def f():
...     raise OSError('operation failed')
...
>>> excs = []
>>> for i in range(3):
...     try:
...         f()
...     except Exception as e:
...         e.add_note(f'Happened in Iteration {i+1}')
...         excs.append(e)
...
>>> raise ExceptionGroup('We have some problems', excs)
+ Exception Group Traceback (most recent call last):
|   File "<stdin>", line 1, in <module>
|       raise ExceptionGroup('We have some problems', excs)
| ExceptionGroup: We have some problems (3 sub-exceptions)
+----- 1 -----
| Traceback (most recent call last):
|   File "<stdin>", line 3, in <module>
|       f()
|       ~^^
|   File "<stdin>", line 2, in f
|       raise OSError('operation failed')
| OSError: operation failed
| Happened in Iteration 1
+----- 2 -----
| Traceback (most recent call last):
|   File "<stdin>", line 3, in <module>
|       f()
|       ~^^
|   File "<stdin>", line 2, in f
|       raise OSError('operation failed')
| OSError: operation failed
| Happened in Iteration 2
+----- 3 -----
| Traceback (most recent call last):
|   File "<stdin>", line 3, in <module>
|       f()
|       ~^^
|   File "<stdin>", line 2, in f
|       raise OSError('operation failed')
| OSError: operation failed
| Happened in Iteration 3

```

+-----

>>>

9. Classes¶

Classes provide a means of bundling data and functionality together. Creating a new class creates a new *type* of object, allowing new *instances* of that type to be made. Each class instance can have attributes attached to it for maintaining its state. Class instances can also have methods (defined by its class) for modifying its state.

Compared with other programming languages, Python's class mechanism adds classes with a minimum of new syntax and semantics. It is a mixture of the class mechanisms found in C++ and Modula-3. Python classes provide all the standard features of Object Oriented Programming: the class inheritance mechanism allows multiple base classes, a derived class can override any methods of its base class or classes, and a method can call the method of a base class with the same name. Objects can contain arbitrary amounts and kinds of data. As is true for modules, classes partake of the dynamic nature of Python: they are created at runtime, and can be modified further after creation.

In C++ terminology, normally class members (including the data members) are *public* (except see below [Private Variables](#)), and all member functions are *virtual*. As in Modula-3, there are no shorthands for referencing the object's members from its methods: the method function is declared with an explicit first argument representing the object, which is provided implicitly by the call. As in Smalltalk, classes themselves are objects. This provides semantics for importing and renaming. Unlike C++ and Modula-3, built-in types can be used as base classes for extension by the user. Also, like in C++, most built-in operators with special syntax (arithmetic operators, subscripting etc.) can be redefined for class instances.

(Lacking universally accepted terminology to talk about classes, I will make occasional use of Smalltalk and C++ terms. I would use Modula-3 terms, since its object-oriented semantics are closer to those of Python than C++, but I expect that few readers have heard of it.)

9.1. A Word About Names and Objects¶

Objects have individuality, and multiple names (in multiple scopes) can be bound to the same object. This is known as aliasing in other languages. This is usually not appreciated on a first glance at Python, and can be safely ignored when dealing with immutable basic types (numbers, strings, tuples). However, aliasing has a possibly surprising effect on the semantics of Python code involving mutable objects such as lists, dictionaries, and most other types. This is usually used to the benefit of the program, since aliases behave like pointers in some respects. For example, passing an object is cheap since only a pointer is passed by the implementation; and if a function modifies an object passed as an argument, the caller will see the change — this eliminates the need for two different argument passing mechanisms as in Pascal.

9.2. Python Scopes and Namespaces¶

Before introducing classes, I first have to tell you something about Python's scope rules. Class definitions play some neat tricks with namespaces, and you need to know how scopes and namespaces work to fully understand what's going on. Incidentally, knowledge about this subject is useful for any advanced Python programmer.

Let's begin with some definitions.

A *namespace* is a mapping from names to objects. Most namespaces are currently implemented as Python dictionaries, but that's normally not noticeable in any way (except for performance), and it may change in the future. Examples of namespaces are: the set of built-in names (containing functions such as

```
abs()
```

, and built-in exception names); the global names in a module; and the local names in a function invocation. In a sense the set of attributes of an object also form a namespace. The important thing to know about namespaces is that there is absolutely no relation between names in different namespaces; for instance, two different modules may both define a function

```
maximize
```

without confusion — users of the modules must prefix it with the module name.

By the way, I use the word *attribute* for any name following a dot — for example, in the expression

```
z.real
```

,

```
real
```

is an attribute of the object

```
z
```

. Strictly speaking, references to names in modules are attribute references: in the expression

```
modname.funcname
```

,

```
modname
```

is a module object and

```
funcname
```

is an attribute of it. In this case there happens to be a straightforward mapping between the module's attributes and the global names defined in the module: they share the same namespace!

[1]

Attributes may be read-only or writable. In the latter case, assignment to attributes is possible.

Module attributes are writable: you can write

```
modname.the_answer = 42
```

. Writable attributes may also be deleted with the

```
del
```

statement. For example,

```
del modname.the_answer
```

will remove the attribute

```
the_answer
```

from the object named by

```
modname
```

.

Namespaces are created at different moments and have different lifetimes. The namespace containing the built-in names is created when the Python interpreter starts up, and is never deleted. The global namespace for a module is created when the module definition is read in; normally, module namespaces also last until the interpreter quits. The statements executed by the top-level invocation of the interpreter, either read from a script file or interactively, are considered part of a module called

```
__main__
```

, so they have their own global namespace. (The built-in names actually also live in a module; this is called

```
builtins
```

.)

The local namespace for a function is created when the function is called, and deleted when the function returns or raises an exception that is not handled within the function. (Actually, forgetting would be a better way to describe what actually happens.) Of course, recursive invocations each have their own local namespace.

A *scope* is a textual region of a Python program where a namespace is directly accessible. “Directly accessible” here means that an unqualified reference to a name attempts to find the name in the namespace.

Although scopes are determined statically, they are used dynamically. At any time during execution, there are 3 or 4 nested scopes whose namespaces are directly accessible:

- the innermost scope, which is searched first, contains the local names

- the scopes of any enclosing functions, which are searched starting with the nearest enclosing scope, contain non-local, but also non-global names
- the next-to-last scope contains the current module's global names
- the outermost scope (searched last) is the namespace containing built-in names

If a name is declared global, then all references and assignments go directly to the next-to-last scope containing the module's global names. To rebind variables found outside of the innermost scope, the

```
nonlocal
```

statement can be used; if not declared nonlocal, those variables are read-only (an attempt to write to such a variable will simply create a *new* local variable in the innermost scope, leaving the identically named outer variable unchanged).

Usually, the local scope references the local names of the (textually) current function. Outside functions, the local scope references the same namespace as the global scope: the module's namespace. Class definitions place yet another namespace in the local scope.

It is important to realize that scopes are determined textually: the global scope of a function defined in a module is that module's namespace, no matter from where or by what alias the function is called. On the other hand, the actual search for names is done dynamically, at run time — however, the language definition is evolving towards static name resolution, at “compile” time, so don't rely on dynamic name resolution! (In fact, local variables are already determined statically.)

A special quirk of Python is that – if no

```
global
```

or

```
nonlocal
```

statement is in effect – assignments to names always go into the innermost scope. Assignments do not copy data — they just bind names to objects. The same is true for deletions: the statement

```
del x
```

removes the binding of

```
x
```

from the namespace referenced by the local scope. In fact, all operations that introduce new names use the local scope: in particular,

```
import
```

statements and function definitions bind the module or function name in the local scope.

The

```
global
```

statement can be used to indicate that particular variables live in the global scope and should be rebound there; the

```
nonlocal
```

statement indicates that particular variables live in an enclosing scope and should be rebound there.

9.2.1. Scopes and Namespaces Example¶

This is an example demonstrating how to reference the different scopes and namespaces, and how

```
global
```

and

```
nonlocal
```

affect variable binding:

Copy

```
def scope_test():
    def do_local():
        spam = "local spam"

    def do_nonlocal():
        nonlocal spam
        spam = "nonlocal spam"

    def do_global():
        global spam
        spam = "global spam"

    spam = "test spam"
    do_local()
    print("After local assignment:", spam)
    do_nonlocal()
    print("After nonlocal assignment:", spam)
    do_global()
    print("After global assignment:", spam)

scope_test()
print("In global scope:", spam)
```

The output of the example code is:

```
After local assignment: test spam
After nonlocal assignment: nonlocal spam
After global assignment: nonlocal spam
In global scope: global spam
```

Note how the *local* assignment (which is default) didn't change *scope_test*'s binding of *spam*. The

```
nonlocal
```

assignment changed *scope_test*'s binding of *spam*, and the

```
global
```

assignment changed the module-level binding.

You can also see that there was no previous binding for *spam* before the

```
global
```

assignment.

9.3. A First Look at Classes¶

Classes introduce a little bit of new syntax, three new object types, and some new semantics.

9.3.1. Class Definition Syntax¶

The simplest form of class definition looks like this:

Copy

```
class ClassName:
    <statement-1>
    .
    .
    .
    <statement-N>
```

Class definitions, like function definitions (

```
def
```

statements) must be executed before they have any effect. (You could conceivably place a class definition in a branch of an

```
if
```

statement, or inside a function.)

In practice, the statements inside a class definition will usually be function definitions, but other statements are allowed, and sometimes useful — we'll come back to this later. The function definitions inside a class normally have a peculiar form of argument list, dictated by the calling conventions for methods — again, this is explained later.

When a class definition is entered, a new namespace is created, and used as the local scope — thus, all assignments to local variables go into this new namespace. In particular, function definitions bind the name of the new function here.

When a class definition is left normally (via the end), a *class object* is created. This is basically a wrapper around the contents of the namespace created by the class definition; we'll learn more about class objects in the next section. The original local scope (the one in effect just before the class definition was entered) is reinstated, and the class object is bound here to the class name given in the class definition header (

```
ClassName
```

in the example).

9.3.2. Class Objects¶

Class objects support two kinds of operations: attribute references and instantiation.

Attribute references use the standard syntax used for all attribute references in Python:

```
obj.name
```

. Valid attribute names are all the names that were in the class's namespace when the class object was created. So, if the class definition looked like this:

Copy

```
class MyClass:
    """A simple example class"""
    i = 12345

    def f(self):
        return 'hello world'
```

then

```
MyClass.i
```

and

```
MyClass.f
```

are valid attribute references, returning an integer and a function object, respectively. Class attributes can also be assigned to, so you can change the value of

```
MyClass.i
```

by assignment.

```
__doc__
```

is also a valid attribute, returning the docstring belonging to the class:

```
"A simple example class"
```

.

Class *instantiation* uses function notation. Just pretend that the class object is a parameterless function that returns a new instance of the class. For example (assuming the above class):

Copy

```
x = MyClass()
```

creates a new *instance* of the class and assigns this object to the local variable

```
x
```

.

The instantiation operation (“calling” a class object) creates an empty object. Many classes like to create objects with instances customized to a specific initial state. Therefore a class may define a special method named

```
__init__()
```

, like this:

Copy

```
def __init__(self):  
    self.data = []
```

When a class defines an

```
__init__()
```

method, class instantiation automatically invokes

```
__init__()
```

for the newly created class instance. So in this example, a new, initialized instance can be obtained by:

Copy


```
x = MyClass()
```

Of course, the

```
__init__()
```

method may have arguments for greater flexibility. In that case, arguments given to the class instantiation operator are passed on to

```
__init__()
```

. For example,

Copy

```
>>> class Complex:
...     def __init__(self, realpart, imagpart):
...         self.r = realpart
...         self.i = imagpart
...
>>> x = Complex(3.0, -4.5)
>>> x.r, x.i
(3.0, -4.5)
```

9.3.3. Instance Objects¶

Now what can we do with instance objects? The only operations understood by instance objects are attribute references. There are two kinds of valid attribute names: data attributes and methods.

Data attributes correspond to “instance variables” in Smalltalk, and to “data members” in C++. Data attributes need not be declared; like local variables, they spring into existence when they are first assigned to. For example, if

```
x
```

is the instance of

```
MyClass
```

created above, the following piece of code will print the value

```
16
```

, without leaving a trace:

Copy

```
x.counter = 1
while x.counter < 10:
    x.counter = x.counter * 2
print(x.counter)
del x.counter
```

The other kind of instance attribute reference is a *method*. A method is a function that “belongs to” an object.

Valid method names of an instance object depend on its class. By definition, all attributes of a class that are function objects define corresponding methods of its instances. So in our example,

```
x.f
```

is a valid method reference, since

```
MyClass.f
```

is a function, but

```
x.i
```

is not, since

```
MyClass.i
```

is not. But

```
x.f
```

is not the same thing as

```
MyClass.f
```

— it is a *method object*, not a function object.

9.3.4. Method Objects¶

Usually, a method is called right after it is bound:

Copy

```
x.f()
```

If

```
x = MyClass()
```

, as above, this will return the string

```
'hello world'
```

. However, it is not necessary to call a method right away:

```
x.f
```

is a method object, and can be stored away and called at a later time. For example:

Copy

```
xf = x.f
while True:
    print(xf())
```

will continue to print

```
hello world
```

until the end of time.

What exactly happens when a method is called? You may have noticed that

```
x.f()
```

was called without an argument above, even though the function definition for

```
f()
```

specified an argument. What happened to the argument? Surely Python raises an exception when a function that requires an argument is called without any — even if the argument isn't actually used...

Actually, you may have guessed the answer: the special thing about methods is that the instance object is passed as the first argument of the function. In our example, the call

```
x.f()
```

is exactly equivalent to

```
MyClass.f(x)
```

. In general, calling a method with a list of n arguments is equivalent to calling the corresponding function with an argument list that is created by inserting the method's instance object before the first argument.

In general, methods work as follows. When a non-data attribute of an instance is referenced, the instance's class is searched. If the name denotes a valid class attribute that is a function object, references to both the instance object and the function object are packed into a method object. When the method object is called with an argument list, a new argument list is constructed from the instance object and the argument list, and the function object is called with this new argument list.

9.3.5. Class and Instance Variables¶

Generally speaking, instance variables are for data unique to each instance and class variables are for attributes and methods shared by all instances of the class:

Copy

```
class Dog:

    kind = 'canine'          # class variable shared by all instances

    def __init__(self, name):
        self.name = name    # instance variable unique to each instance

>>> d = Dog('Fido')
>>> e = Dog('Buddy')
>>> d.kind                # shared by all dogs
'canine'
>>> e.kind                # shared by all dogs
'canine'
>>> d.name                # unique to d
'Fido'
>>> e.name                # unique to e
'Buddy'
```

As discussed in [A Word About Names and Objects](#), shared data can have possibly surprising effects involving mutable objects such as lists and dictionaries. For example, the *tricks* list in the following code should not be used as a class variable because just a single list would be shared by all *Dog* instances:

Copy

```
class Dog:

    tricks = []             # mistaken use of a class variable

    def __init__(self, name):
        self.name = name

    def add_trick(self, trick):
        self.tricks.append(trick)

>>> d = Dog('Fido')
>>> e = Dog('Buddy')
>>> d.add_trick('roll over')
>>> e.add_trick('play dead')
>>> d.tricks                # unexpectedly shared by all dogs
['roll over', 'play dead']
```

Correct design of the class should use an instance variable instead:

Copy

```
class Dog:

    def __init__(self, name):
        self.name = name
        self.tricks = []    # creates a new empty list for each dog

    def add_trick(self, trick):
        self.tricks.append(trick)

>>> d = Dog('Fido')
>>> e = Dog('Buddy')
>>> d.add_trick('roll over')
>>> e.add_trick('play dead')
>>> d.tricks
['roll over']
>>> e.tricks
['play dead']
```

9.4. Random Remarks ¶

If the same attribute name occurs in both an instance and in a class, then attribute lookup prioritizes the instance:

Copy

```
>>> class Warehouse:
...     purpose = 'storage'
...     region = 'west'
...
>>> w1 = Warehouse()
>>> print(w1.purpose, w1.region)
storage west
>>> w2 = Warehouse()
>>> w2.region = 'east'
>>> print(w2.purpose, w2.region)
storage east
```

Data attributes may be referenced by methods as well as by ordinary users (“clients”) of an object. In other words, classes are not usable to implement pure abstract data types. In fact, nothing in

Python makes it possible to enforce data hiding — it is all based upon convention. (On the other hand, the Python implementation, written in C, can completely hide implementation details and control access to an object if necessary; this can be used by extensions to Python written in C.)

Clients should use data attributes with care — clients may mess up invariants maintained by the methods by stamping on their data attributes. Note that clients may add data attributes of their own to an instance object without affecting the validity of the methods, as long as name conflicts are avoided — again, a naming convention can save a lot of headaches here.

There is no shorthand for referencing data attributes (or other methods!) from within methods. I find that this actually increases the readability of methods: there is no chance of confusing local variables and instance variables when glancing through a method.

Often, the first argument of a method is called

```
self
```

. This is nothing more than a convention: the name

```
self
```

has absolutely no special meaning to Python. Note, however, that by not following the convention your code may be less readable to other Python programmers, and it is also conceivable that a *class browser* program might be written that relies upon such a convention.

Any function object that is a class attribute defines a method for instances of that class. It is not necessary that the function definition is textually enclosed in the class definition: assigning a function object to a local variable in the class is also ok. For example:

Copy

```
# Function defined outside the class
def f1(self, x, y):
    return min(x, x+y)

class C:
    f = f1

    def g(self):
        return 'hello world'

    h = g
```

Now

```
f
```

,

g

and

h

are all attributes of class

C

that refer to function objects, and consequently they are all methods of instances of

C

—

h

being exactly equivalent to

g

. Note that this practice usually only serves to confuse the reader of a program.

Methods may call other methods by using method attributes of the

self

argument:

Copy

```
class Bag:
    def __init__(self):
        self.data = []

    def add(self, x):
        self.data.append(x)

    def addtwice(self, x):
        self.add(x)
        self.add(x)
```

Methods may reference global names in the same way as ordinary functions. The global scope associated with a method is the module containing its definition. (A class is never used as a global scope.) While one rarely encounters a good reason for using global data in a method, there are many legitimate uses of the global scope: for one thing, functions and modules imported into the global scope can be used by methods, as well as functions and classes defined in it. Usually, the class containing the method is itself defined in this global scope, and in the next section we'll find some good reasons why a method would want to reference its own class.

Each value is an object, and therefore has a *class* (also called its *type*). It is stored as

```
object.__class__
```

9.5. Inheritance¶

Of course, a language feature would not be worthy of the name “class” without supporting inheritance. The syntax for a derived class definition looks like this:

Copy

```
class DerivedClassName(BaseClassName):  
    <statement-1>  
    .  
    .  
    .  
    <statement-N>
```

The name

BaseClassName

must be defined in a namespace accessible from the scope containing the derived class definition. In place of a base class name, other arbitrary expressions are also allowed. This can be useful, for example, when the base class is defined in another module:

Copy

```
class DerivedClassName(modname.BaseClassName):
```

Execution of a derived class definition proceeds the same as for a base class. When the class object is constructed, the base class is remembered. This is used for resolving attribute references: if a requested attribute is not found in the class, the search proceeds to look in the base class. This rule is applied recursively if the base class itself is derived from some other class.

There’s nothing special about instantiation of derived classes:

```
DerivedClassName()
```

creates a new instance of the class. Method references are resolved as follows: the corresponding class attribute is searched, descending down the chain of base classes if necessary, and the method reference is valid if this yields a function object.

Derived classes may override methods of their base classes. Because methods have no special privileges when calling other methods of the same object, a method of a base class that calls

another method defined in the same base class may end up calling a method of a derived class that overrides it. (For C++ programmers: all methods in Python are effectively

```
virtual
```

.)

An overriding method in a derived class may in fact want to extend rather than simply replace the base class method of the same name. There is a simple way to call the base class method directly: just call

```
BaseClassName.methodname(self,  
arguments)
```

. This is occasionally useful to clients as well. (Note that this only works if the base class is accessible as

```
BaseClassName
```

in the global scope.)

Python has two built-in functions that work with inheritance:

- Use

```
isinstance()
```

to check an instance's type:

```
isinstance(obj, int)
```

will be

```
True
```

only if

```
obj.__class__
```

is

```
int
```

or some class derived from

```
int
```

.

- Use

```
issubclass()
```

to check class inheritance:

```
issubclass(bool, int)
```

is

True

since

bool

is a subclass of

int

. However,

```
issubclass(float, int)
```

is

False

since

float

is not a subclass of

int

.

9.5.1. Multiple Inheritance¶

Python supports a form of multiple inheritance as well. A class definition with multiple base classes looks like this:

Copy

```
class DerivedClassName(Base1, Base2, Base3):
    <statement-1>
    .
    .
    .
    <statement-N>
```

For most purposes, in the simplest cases, you can think of the search for attributes inherited from a parent class as depth-first, left-to-right, not searching twice in the same class where there is an overlap in the hierarchy. Thus, if an attribute is not found in

DerivedClassName

, it is searched for in

Base1

, then (recursively) in the base classes of

```
Base1
```

, and if it was not found there, it was searched for in

```
Base2
```

, and so on.

In fact, it is slightly more complex than that; the method resolution order changes dynamically to support cooperative calls to

```
super()
```

. This approach is known in some other multiple-inheritance languages as call-next-method and is more powerful than the super call found in single-inheritance languages.

Dynamic ordering is necessary because all cases of multiple inheritance exhibit one or more diamond relationships (where at least one of the parent classes can be accessed through multiple paths from the bottommost class). For example, all classes inherit from

```
object
```

, so any case of multiple inheritance provides more than one path to reach

```
object
```

. To keep the base classes from being accessed more than once, the dynamic algorithm linearizes the search order in a way that preserves the left-to-right ordering specified in each class, that calls each parent only once, and that is monotonic (meaning that a class can be subclassed without affecting the precedence order of its parents). Taken together, these properties make it possible to design reliable and extensible classes with multiple inheritance. For more detail, see The Python 2.3 Method Resolution Order.

9.6. Private Variables¶

“Private” instance variables that cannot be accessed except from inside an object don’t exist in Python. However, there is a convention that is followed by most Python code: a name prefixed with an underscore (e.g.

```
_spam
```

) should be treated as a non-public part of the API (whether it is a function, a method or a data member). It should be considered an implementation detail and subject to change without notice.

Since there is a valid use-case for class-private members (namely to avoid name clashes of names with names defined by subclasses), there is limited support for such a mechanism, called *name mangling*. Any identifier of the form

```
__spam
```

(at least two leading underscores, at most one trailing underscore) is textually replaced with

```
_classname__spam
```

, where

```
classname
```

is the current class name with leading underscore(s) stripped. This mangling is done without regard to the syntactic position of the identifier, as long as it occurs within the definition of a class.

See also

The private name mangling specifications for details and special cases.

Name mangling is helpful for letting subclasses override methods without breaking intraclass method calls. For example:

Copy

```
class Mapping:
    def __init__(self, iterable):
        self.items_list = []
        self.__update(iterable)

    def update(self, iterable):
        for item in iterable:
            self.items_list.append(item)

    __update = update    # private copy of original update() method

class MappingSubclass(Mapping):

    def update(self, keys, values):
        # provides new signature for update()
        # but does not break __init__()
        for item in zip(keys, values):
            self.items_list.append(item)
```

The above example would work even if

```
MappingSubclass
```

were to introduce a

```
__update
```

identifier since it is replaced with

```
_Mapping__update
```

in the

```
Mapping
```

class and

```
_MappingSubclass__update
```

in the

```
MappingSubclass
```

class respectively.

Note that the mangling rules are designed mostly to avoid accidents; it still is possible to access or modify a variable that is considered private. This can even be useful in special circumstances, such as in the debugger.

Notice that code passed to

```
exec()
```

or

```
eval()
```

does not consider the classname of the invoking class to be the current class; this is similar to the effect of the

```
global
```

statement, the effect of which is likewise restricted to code that is byte-compiled together. The same restriction applies to

```
getattr()
```

,

```
setattr()
```

and

```
delattr()
```

, as well as when referencing

```
__dict__
```

directly.

9.7. Odds and Ends¶

Sometimes it is useful to have a data type similar to the Pascal “record” or C “struct”, bundling together a few named data items. The idiomatic approach is to use

```
dataclasses
```

for this purpose:

Copy

```
from dataclasses import dataclass

@dataclass
class Employee:
    name: str
    dept: str
    salary: int
```

Copy

```
>>> john = Employee('john', 'computer lab', 1000)
>>> john.dept
'computer lab'
>>> john.salary
1000
```

A piece of Python code that expects a particular abstract data type can often be passed a class that emulates the methods of that data type instead. For instance, if you have a function that formats some data from a file object, you can define a class with methods

```
read()
```

and

```
readline()
```

that get the data from a string buffer instead, and pass it as an argument.

Instance method objects have attributes, too:

```
m.__self__
```

is the instance object with the method

```
m()
```

, and

```
m.__func__
```

is the function object corresponding to the method.

9.8. Iterators¶

By now you have probably noticed that most container objects can be looped over using a

```
for
```

statement:

Copy

```
for element in [1, 2, 3]:
    print(element)
for element in (1, 2, 3):
    print(element)
for key in {'one':1, 'two':2}:
    print(key)
for char in "123":
    print(char)
for line in open("myfile.txt"):
    print(line, end='')
```

This style of access is clear, concise, and convenient. The use of iterators pervades and unifies Python. Behind the scenes, the

```
for
```

statement calls

```
iter()
```

on the container object. The function returns an iterator object that defines the method

```
__next__()
```

which accesses elements in the container one at a time. When there are no more elements,

```
__next__()
```

raises a

```
StopIteration
```

exception which tells the

```
for
```

loop to terminate. You can call the

```
__next__()
```

method using the

```
next()
```

built-in function; this example shows how it all works:

Copy

```
>>> s = 'abc'
>>> it = iter(s)
>>> it
<str_iterator object at 0x10c90e650>
>>> next(it)
'a'
>>> next(it)
'b'
>>> next(it)
'c'
>>> next(it)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
    next(it)
StopIteration
```

Having seen the mechanics behind the iterator protocol, it is easy to add iterator behavior to your classes. Define an

```
__iter__()
```

method which returns an object with a

```
__next__()
```

method. If the class defines

```
__next__()
```

, then

```
__iter__()
```

can just return

```
self
```

:

Copy


```

class Reverse:
    """Iterator for looping over a sequence backwards."""
    def __init__(self, data):
        self.data = data
        self.index = len(data)

    def __iter__(self):
        return self

    def __next__(self):
        if self.index == 0:
            raise StopIteration
        self.index = self.index - 1
        return self.data[self.index]

```

Copy

```

>>> rev = Reverse('spam')
>>> iter(rev)
<__main__.Reverse object at 0x00A1DB50>
>>> for char in rev:
...     print(char)
...
m
a
p
s

```

9.9. Generators¶

Generators are a simple and powerful tool for creating iterators. They are written like regular functions but use the

`yield`

statement whenever they want to return data. Each time

`next()`

is called on it, the generator resumes where it left off (it remembers all the data values and which statement was last executed). An example shows that generators can be trivially easy to create:

Copy

```
def reverse(data):
    for index in range(len(data)-1, -1, -1):
        yield data[index]
```

Copy

```
>>> for char in reverse('golf'):
...     print(char)
...
f
l
o
g
```

Anything that can be done with generators can also be done with class-based iterators as described in the previous section. What makes generators so compact is that the

```
__iter__()
```

and

```
__next__()
```

methods are created automatically.

Another key feature is that the local variables and execution state are automatically saved between calls. This made the function easier to write and much more clear than an approach using instance variables like

```
self.index
```

and

```
self.data
```

.

In addition to automatic method creation and saving program state, when generators terminate, they automatically raise

```
StopIteration
```

. In combination, these features make it easy to create iterators with no more effort than writing a regular function.

9.10. Generator Expressions ¶

Some simple generators can be coded succinctly as expressions using a syntax similar to list comprehensions but with parentheses instead of square brackets. These expressions are designed for situations where the generator is used right away by an enclosing function. Generator

expressions are more compact but less versatile than full generator definitions and tend to be more memory friendly than equivalent list comprehensions.

Examples:

Copy

```
>>> sum(i*i for i in range(10))           # sum of squares
285

>>> xvec = [10, 20, 30]
>>> yvec = [7, 5, 3]
>>> sum(x*y for x,y in zip(xvec, yvec))   # dot product
260

>>> unique_words = set(word for line in page for word in line.split())

>>> valedictorian = max((student.gpa, student.name) for student in graduates)

>>> data = 'golf'
>>> list(data[i] for i in range(len(data)-1, -1, -1))
['f', 'l', 'o', 'g']
```

Footnotes

[\[1\]](#)

Except for one thing. Module objects have a secret read-only attribute called

```
__dict__
```

which returns the dictionary used to implement the module's namespace; the name

```
__dict__
```

is an attribute but not a global name. Obviously, using this violates the abstraction of namespace implementation, and should be restricted to things like post-mortem debuggers.

10. Brief Tour of the Standard Library¶

10.1. Operating System Interface¶

The

```
os
```

module provides dozens of functions for interacting with the operating system:

Copy

```
>>> import os
>>> os.getcwd()          # Return the current working directory
'C:\Python314'
>>> os.chdir('/server/accesslogs')  # Change current working directory
>>> os.system('mkdir today')  # Run the command mkdir in the system shell
0
```

Be sure to use the

```
import os
```

style instead of

```
from os import *
```

. This will keep

```
os.open()
```

from shadowing the built-in

```
open()
```

function which operates much differently.

The built-in

```
dir()
```

and

```
help()
```

functions are useful as interactive aids for working with large modules like

```
os
```

:

Copy

```
>>> import os
>>> dir(os)
<returns a list of all module functions>
>>> help(os)
<returns an extensive manual page created from the module's docstrings>
```

For daily file and directory management tasks, the

shutil

module provides a higher level interface that is easier to use:

Copy

```
>>> import shutil
>>> shutil.copyfile('data.db', 'archive.db')
'archive.db'
>>> shutil.move('/build/executables', 'installdir')
'installdir'
```

10.2. File Wildcards¶

The

glob

module provides a function for making file lists from directory wildcard searches:

Copy

```
>>> import glob
>>> glob.glob('*.py')
['primes.py', 'random.py', 'quote.py']
```

10.3. Command Line Arguments¶

Common utility scripts often need to process command line arguments. These arguments are stored in the

sys

module's *argv* attribute as a list. For instance, let's take the following

demo.py

file:

Copy

```
# File demo.py
import sys
print(sys.argv)
```

Here is the output from running

```
python demo.py one two three
```

at the command line:

Copy

```
['demo.py', 'one', 'two', 'three']
```

The

argparse

module provides a more sophisticated mechanism to process command line arguments. The following script extracts one or more filenames and an optional number of lines to be displayed:

Copy

```
import argparse

parser = argparse.ArgumentParser(
    prog='top',
    description='Show top lines from each file')
parser.add_argument('filenames', nargs='+')
parser.add_argument('-l', '--lines', type=int, default=10)
args = parser.parse_args()
print(args)
```

When run at the command line with

```
python top.py --lines=5 alpha.txt
beta.txt
```

, the script sets

```
args.lines
```

to

```
5
```

and

```
args filenames
```

to

```
['alpha.txt', 'beta.txt']
```

10.4. Error Output Redirection and Program Termination¶

The

```
sys
```

module also has attributes for *stdin*, *stdout*, and *stderr*. The latter is useful for emitting warnings and error messages to make them visible even when *stdout* has been redirected:

Copy

```
>>> sys.stderr.write('Warning, log file not found starting a new one\n')
Warning, log file not found starting a new one
```

The most direct way to terminate a script is to use

```
sys.exit()
```

10.5. String Pattern Matching¶

The

```
re
```

module provides regular expression tools for advanced string processing. For complex matching and manipulation, regular expressions offer succinct, optimized solutions:

Copy

```
>>> import re
>>> re.findall(r'\b[a-z]*', 'which foot or hand fell fastest')
['foot', 'fell', 'fastest']
>>> re.sub(r'(\b[a-z]+) \1', r'\1', 'cat in the the hat')
'cat in the hat'
```

When only simple capabilities are needed, string methods are preferred because they are easier to read and debug:

Copy

```
>>> 'tea for too'.replace('too', 'two')
'tea for two'
```

10.6. Mathematics¶

The

`math`

module gives access to the underlying C library functions for floating-point math:

Copy

```
>>> import math
>>> math.cos(math.pi / 4)
0.70710678118654757
>>> math.log(1024, 2)
10.0
```

The

`random`

module provides tools for making random selections:

Copy

```
>>> import random
>>> random.choice(['apple', 'pear', 'banana'])
'apple'
>>> random.sample(range(100), 10)    # sampling without replacement
[30, 83, 16, 4, 8, 81, 41, 50, 18, 33]
>>> random.random()                  # random float from the interval [0.0, 1.0)
0.17970987693706186
>>> random.randrange(6)              # random integer chosen from range(6)
4
```

The

`statistics`

module calculates basic statistical properties (the mean, median, variance, etc.) of numeric data:

Copy


```
>>> import statistics
>>> data = [2.75, 1.75, 1.25, 0.25, 0.5, 1.25, 3.5]
>>> statistics.mean(data)
1.6071428571428572
>>> statistics.median(data)
1.25
>>> statistics.variance(data)
1.3720238095238095
```

The SciPy project <<https://scipy.org>> has many other modules for numerical computations.

10.7. Internet Access¶

There are a number of modules for accessing the internet and processing internet protocols. Two of the simplest are

```
urllib.request
```

for retrieving data from URLs and

```
smtplib
```

for sending mail:

Copy

```
>>> from urllib.request import urlopen
>>> with urlopen('https://docs.python.org/3/') as response:
...     for line in response:
...         line = line.decode()           # Convert bytes to a str
...         if 'updated' in line:
...             print(line.rstrip())       # Remove trailing newline
...
    Last updated on Nov 11, 2025 (20:11 UTC).

>>> import smtplib
>>> server = smtplib.SMTP('localhost')
>>> server.sendmail('soothsayer@example.org', 'jcaesar@example.org',
... """To: jcaesar@example.org
... From: soothsayer@example.org
...
... Beware the Ides of March.
... """)
>>> server.quit()
```

(Note that the second example needs a mailserver running on localhost.)

10.8. Dates and Times¶

The

`datetime`

module supplies classes for manipulating dates and times in both simple and complex ways. While date and time arithmetic is supported, the focus of the implementation is on efficient member extraction for output formatting and manipulation. The module also supports objects that are timezone aware.

Copy

```
>>> # dates are easily constructed and formatted
>>> from datetime import date
>>> now = date.today()
>>> now
datetime.date(2003, 12, 2)
>>> now.strftime("%m-%d-%y. %d %b %Y is a %A on the %d day of %B.")
'12-02-03. 02 Dec 2003 is a Tuesday on the 02 day of December.'

>>> # dates support calendar arithmetic
>>> birthday = date(1964, 7, 31)
>>> age = now - birthday
>>> age.days
14368
```

10.9. Data Compression¶

Common data archiving and compression formats are directly supported by modules including:

`zlib`

,

`gzip`

,

`bz2`

,

`lzma`

```
zipfile
```

and

```
tarfile
```

Copy

```
>>> import zlib
>>> s = b'witch which has which witches wrist watch'
>>> len(s)
41
>>> t = zlib.compress(s)
>>> len(t)
37
>>> zlib.decompress(t)
b'witch which has which witches wrist watch'
>>> zlib.crc32(s)
226805979
```

10.10. Performance Measurement ¶

Some Python users develop a deep interest in knowing the relative performance of different approaches to the same problem. Python provides a measurement tool that answers those questions immediately.

For example, it may be tempting to use the tuple packing and unpacking feature instead of the traditional approach to swapping arguments. The

```
timeit
```

module quickly demonstrates a modest performance advantage:

Copy

```
>>> from timeit import Timer
>>> Timer('t=a; a=b; b=t', 'a=1; b=2').timeit()
0.57535828626024577
>>> Timer('a,b = b,a', 'a=1; b=2').timeit()
0.54962537085770791
```

In contrast to

```
timeit
```

's fine level of granularity, the

```
profile
```

and

```
pstats
```

modules provide tools for identifying time critical sections in larger blocks of code.

10.11. Quality Control¶

One approach for developing high quality software is to write tests for each function as it is developed and to run those tests frequently during the development process.

The

```
doctest
```

module provides a tool for scanning a module and validating tests embedded in a program's docstrings. Test construction is as simple as cutting-and-pasting a typical call along with its results into the docstring. This improves the documentation by providing the user with an example and it allows the doctest module to make sure the code remains true to the documentation:

Copy

```
def average(values):
    """Computes the arithmetic mean of a list of numbers.

    >>> print(average([20, 30, 70]))
    40.0
    """
    return sum(values) / len(values)

import doctest
doctest.testmod()  # automatically validate the embedded tests
```

The

```
unittest
```

module is not as effortless as the

```
doctest
```

module, but it allows a more comprehensive set of tests to be maintained in a separate file:

Copy

```
import unittest

class TestStatisticalFunctions(unittest.TestCase):

    def test_average(self):
        self.assertEqual(average([20, 30, 70]), 40.0)
        self.assertEqual(round(average([1, 5, 7]), 1), 4.3)
        with self.assertRaises(ZeroDivisionError):
            average([])
        with self.assertRaises(TypeError):
            average(20, 30, 70)

unittest.main() # Calling from the command line invokes all tests
```

10.12. Batteries Included¶

Python has a “batteries included” philosophy. This is best seen through the sophisticated and robust capabilities of its larger packages. For example:

- The

```
xmlrpc.client
```

and

```
xmlrpc.server
```

modules make implementing remote procedure calls into an almost trivial task. Despite the modules’ names, no direct knowledge or handling of XML is needed.

- The

```
email
```

package is a library for managing email messages, including MIME and other [RFC 5322](#)-based message documents. Unlike

```
smtpplib
```

and

```
poplib
```

which actually send and receive messages, the email package has a complete toolset for building or decoding complex message structures (including attachments) and for implementing internet encoding and header protocols.

- The

`json`

package provides robust support for parsing this popular data interchange format. The

`csv`

module supports direct reading and writing of files in Comma-Separated Value format, commonly supported by databases and spreadsheets. XML processing is supported by the

`xml.etree.ElementTree`

,

`xml.dom`

and

`xml.sax`

packages. Together, these modules and packages greatly simplify data interchange between Python applications and other tools.

- The

`sqlite3`

module is a wrapper for the SQLite database library, providing a persistent database that can be updated and accessed using slightly nonstandard SQL syntax.

- Internationalization is supported by a number of modules including

`gettext`

,

`locale`

, and the

`codecs`

package.

11. Brief Tour of the Standard Library — Part



This second tour covers more advanced modules that support professional programming needs. These modules rarely occur in small scripts.

11.1. Output Formatting

The

```
reprlib
```

module provides a version of

```
repr()
```

customized for abbreviated displays of large or deeply nested containers:

Copy

```
>>> import reprlib
>>> reprlib.repr(set('supercalifragilisticexpialidocious'))
"{'a', 'c', 'd', 'e', 'f', 'g', ...}"
```

The

```
pprint
```

module offers more sophisticated control over printing both built-in and user defined objects in a way that is readable by the interpreter. When the result is longer than one line, the “pretty printer” adds line breaks and indentation to more clearly reveal data structure:

Copy

```
>>> import pprint
>>> t = [[['black', 'cyan'], 'white', ['green', 'red']], [['magenta',
...      'yellow'], 'blue']]
...
>>> pprint.pprint(t, width=30)
[[['black', 'cyan'],
  'white',
  ['green', 'red']],
 [['magenta', 'yellow'],
  'blue']]
```

The

textwrap

module formats paragraphs of text to fit a given screen width:

Copy

```
>>> import textwrap
>>> doc = """The wrap() method is just like fill() except that it returns
... a list of strings instead of one big string with newlines to separate
... the wrapped lines."""
...
>>> print(textwrap.fill(doc, width=40))
The wrap() method is just like fill()
except that it returns a list of strings
instead of one big string with newlines
to separate the wrapped lines.
```

The

locale

module accesses a database of culture specific data formats. The grouping attribute of locale's format function provides a direct way of formatting numbers with group separators:

Copy

```
>>> import locale
>>> locale.setlocale(locale.LC_ALL, 'English_United States.1252')
'English_United States.1252'
>>> conv = locale.localeconv()          # get a mapping of conventions
>>> x = 1234567.8
>>> locale.format_string("%d", x, grouping=True)
'1,234,567'
>>> locale.format_string("%s%.*f", (conv['currency_symbol'],
...                               conv['frac_digits'], x), grouping=True)
'$1,234,567.80'
```

11.2. Templating

The

string

module includes a versatile

Template

class with a simplified syntax suitable for editing by end-users. This allows users to customize their applications without having to alter the application.

The format uses placeholder names formed by

\$

with valid Python identifiers (alphanumeric characters and underscores). Surrounding the placeholder with braces allows it to be followed by more alphanumeric letters with no intervening spaces. Writing

\$\$

creates a single escaped

\$

:

Copy

```
>>> from string import Template
>>> t = Template('${village}folk send $$10 to $cause.')
>>> t.substitute(village='Nottingham', cause='the ditch fund')
'Nottinghamfolk send $10 to the ditch fund.'
```

The

substitute()

method raises a

KeyError

when a placeholder is not supplied in a dictionary or a keyword argument. For mail-merge style applications, user supplied data may be incomplete and the

safe_substitute()

method may be more appropriate — it will leave placeholders unchanged if data is missing:

Copy

```
>>> t = Template('Return the $item to $owner.')
>>> d = dict(item='unladen swallow')
>>> t.substitute(d)
Traceback (most recent call last):
...
KeyError: 'owner'
>>> t.safe_substitute(d)
'Return the unladen swallow to $owner.'
```

Template subclasses can specify a custom delimiter. For example, a batch renaming utility for a photo browser may elect to use percent signs for placeholders such as the current date, image sequence number, or file format:

Copy

```
>>> import time, os.path
>>> photofiles = ['img_1074.jpg', 'img_1076.jpg', 'img_1077.jpg']
>>> class BatchRename(Template):
...     delimiter = '%'
...
>>> fmt = input('Enter rename style (%d-date %n-seqnum %f-format): ')
Enter rename style (%d-date %n-seqnum %f-format): Ashley_%n%f

>>> t = BatchRename(fmt)
>>> date = time.strftime('%d%b%y')
>>> for i, filename in enumerate(photofiles):
...     base, ext = os.path.splitext(filename)
...     newname = t.substitute(d=date, n=i, f=ext)
...     print('{0} --> {1}'.format(filename, newname))

img_1074.jpg --> Ashley_0.jpg
img_1076.jpg --> Ashley_1.jpg
img_1077.jpg --> Ashley_2.jpg
```

Another application for templating is separating program logic from the details of multiple output formats. This makes it possible to substitute custom templates for XML files, plain text reports, and HTML web reports.

11.3. Working with Binary Data Record Layouts ¶

The

```
struct
```

module provides

```
pack()
```

and

```
unpack()
```

functions for working with variable length binary record formats. The following example shows how to loop through header information in a ZIP file without using the

```
zipfile
```

module. Pack codes

```
"H"
```

and

```
"I"
```

represent two and four byte unsigned numbers respectively. The

```
"<"
```

indicates that they are standard size and in little-endian byte order:

Copy

```
import struct

with open('myfile.zip', 'rb') as f:
    data = f.read()

start = 0
for i in range(3):
    # show the first 3 file headers
    start += 14
    fields = struct.unpack('<IIHH', data[start:start+16])
    crc32, comp_size, uncomp_size, filenamesize, extra_size = fields

    start += 16
    filename = data[start:start+filenamesize]
    start += filenamesize
    extra = data[start:start+extra_size]
    print(filename, hex(crc32), comp_size, uncomp_size)

    start += extra_size + comp_size    # skip to the next header
```

11.4. Multi-threading¶

Threading is a technique for decoupling tasks which are not sequentially dependent. Threads can be used to improve the responsiveness of applications that accept user input while other tasks run in the background. A related use case is running I/O in parallel with computations in another thread.

The following code shows how the high level

```
threading
```

module can run tasks in background while the main program continues to run:

Copy

```
import threading, zipfile

class AsyncZip(threading.Thread):
    def __init__(self, infile, outfile):
        threading.Thread.__init__(self)
        self.infile = infile
        self.outfile = outfile

    def run(self):
        f = zipfile.ZipFile(self.outfile, 'w', zipfile.ZIP_DEFLATED)
        f.write(self.infile)
        f.close()
        print('Finished background zip of:', self.infile)

background = AsyncZip('mydata.txt', 'myarchive.zip')
background.start()
print('The main program continues to run in foreground.')

background.join()    # Wait for the background task to finish
print('Main program waited until background was done.')
```

The principal challenge of multi-threaded applications is coordinating threads that share data or other resources. To that end, the threading module provides a number of synchronization primitives including locks, events, condition variables, and semaphores.

While those tools are powerful, minor design errors can result in problems that are difficult to reproduce. So, the preferred approach to task coordination is to concentrate all access to a resource in a single thread and then use the

```
queue
```

module to feed that thread with requests from other threads. Applications using

```
Queue
```

objects for inter-thread communication and coordination are easier to design, more readable, and more reliable.

11.5. Logging¶

The

```
logging
```

module offers a full featured and flexible logging system. At its simplest, log messages are sent to a file or to

```
sys.stderr
```

:

Copy

```
import logging
logging.debug('Debugging information')
logging.info('Informational message')
logging.warning('Warning:config file %s not found', 'server.conf')
logging.error('Error occurred')
logging.critical('Critical error -- shutting down')
```

This produces the following output:

```
WARNING:root:Warning:config file server.conf not found
ERROR:root:Error occurred
CRITICAL:root:Critical error -- shutting down
```

By default, informational and debugging messages are suppressed and the output is sent to standard error. Other output options include routing messages through email, datagrams, sockets, or to an HTTP Server. New filters can select different routing based on message priority:

```
DEBUG
```

,

```
INFO
```

,

```
WARNING
```

```
,  
ERROR
```

```
, and
```

```
CRITICAL
```

The logging system can be configured directly from Python or can be loaded from a user editable configuration file for customized logging without altering the application.

11.6. Weak References¶

Python does automatic memory management (reference counting for most objects and garbage collection to eliminate cycles). The memory is freed shortly after the last reference to it has been eliminated.

This approach works fine for most applications but occasionally there is a need to track objects only as long as they are being used by something else. Unfortunately, just tracking them creates a reference that makes them permanent. The

```
weakref
```

module provides tools for tracking objects without creating a reference. When the object is no longer needed, it is automatically removed from a weakref table and a callback is triggered for weakref objects. Typical applications include caching objects that are expensive to create:

Copy

```

>>> import weakref, gc
>>> class A:
...     def __init__(self, value):
...         self.value = value
...     def __repr__(self):
...         return str(self.value)
...
>>> a = A(10)                # create a reference
>>> d = weakref.WeakValueDictionary()
>>> d['primary'] = a          # does not create a reference
>>> d['primary']              # fetch the object if it is still alive
10
>>> del a                    # remove the one reference
>>> gc.collect()             # run garbage collection right away
0
>>> d['primary']              # entry was automatically removed
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
    d['primary']              # entry was automatically removed
  File "C:/python314/lib/weakref.py", line 46, in __getitem__
    o = self.data[key]()
KeyError: 'primary'

```

11.7. Tools for Working with Lists¶

Many data structure needs can be met with the built-in list type. However, sometimes there is a need for alternative implementations with different performance trade-offs.

The

array

module provides an

array

object that is like a list that stores only homogeneous data and stores it more compactly. The following example shows an array of numbers stored as two byte unsigned binary numbers (typecode

"H"

) rather than the usual 16 bytes per entry for regular lists of Python int objects:

Copy

```
>>> from array import array
>>> a = array('H', [4000, 10, 700, 22222])
>>> sum(a)
26932
>>> a[1:3]
array('H', [10, 700])
```

The

collections

module provides a

deque

object that is like a list with faster appends and pops from the left side but slower lookups in the middle. These objects are well suited for implementing queues and breadth first tree searches:

Copy

```
>>> from collections import deque
>>> d = deque(["task1", "task2", "task3"])
>>> d.append("task4")
>>> print("Handling", d.popleft())
Handling task1
```

Copy

```
unsearched = deque([starting_node])
def breadth_first_search(unsearched):
    node = unsearched.popleft()
    for m in gen_moves(node):
        if is_goal(m):
            return m
        unsearched.append(m)
```

In addition to alternative list implementations, the library also offers other tools such as the

bisect

module with functions for manipulating sorted lists:

Copy


```
>>> import bisect
>>> scores = [(100, 'perl'), (200, 'tcl'), (400, 'lua'), (500, 'python')]
>>> bisect.insort(scores, (300, 'ruby'))
>>> scores
[(100, 'perl'), (200, 'tcl'), (300, 'ruby'), (400, 'lua'), (500, 'python')]
```

The

heapq

module provides functions for implementing heaps based on regular lists. The lowest valued entry is always kept at position zero. This is useful for applications which repeatedly access the smallest element but do not want to run a full list sort:

Copy

```
>>> from heapq import heapify, heappop, heappush
>>> data = [1, 3, 5, 7, 9, 2, 4, 6, 8, 0]
>>> heapify(data) # rearrange the list into heap order
>>> heappush(data, -5) # add a new entry
>>> [heappop(data) for i in range(3)] # fetch the three smallest entries
[-5, 0, 1]
```

11.8. Decimal Floating-Point Arithmetic¶

The

decimal

module offers a

Decimal

datatype for decimal floating-point arithmetic. Compared to the built-in

float

implementation of binary floating point, the class is especially helpful for

- financial applications and other uses which require exact decimal representation,
- control over precision,
- control over rounding to meet legal or regulatory requirements,
- tracking of significant decimal places, or
- applications where the user expects the results to match calculations done by hand.

For example, calculating a 5% tax on a 70 cent phone charge gives different results in decimal floating point and binary floating point. The difference becomes significant if the results are rounded to the nearest cent:

Copy

```
>>> from decimal import *
>>> round(Decimal('0.70') * Decimal('1.05'), 2)
Decimal('0.74')
>>> round(.70 * 1.05, 2)
0.73
```

The

Decimal

result keeps a trailing zero, automatically inferring four place significance from multiplicands with two place significance. Decimal reproduces mathematics as done by hand and avoids issues that can arise when binary floating point cannot exactly represent decimal quantities.

Exact representation enables the

Decimal

class to perform modulo calculations and equality tests that are unsuitable for binary floating point:

Copy

```
>>> Decimal('1.00') % Decimal('.10')
Decimal('0.00')
>>> 1.00 % 0.10
0.09999999999999995

>>> sum([Decimal('0.1')]*10) == Decimal('1.0')
True
>>> 0.1 + 0.1 + 0.1 + 0.1 + 0.1 + 0.1 + 0.1 + 0.1 + 0.1 + 0.1 == 1.0
False
```

The

decimal

module provides arithmetic with as much precision as needed:

Copy

```
>>> getcontext().prec = 36
>>> Decimal(1) / Decimal(7)
Decimal('0.142857142857142857142857142857142857')
```

12. Virtual Environments and Packages¶

12.1. Introduction¶

Python applications will often use packages and modules that don't come as part of the standard library. Applications will sometimes need a specific version of a library, because the application may require that a particular bug has been fixed or the application may be written using an obsolete version of the library's interface.

This means it may not be possible for one Python installation to meet the requirements of every application. If application A needs version 1.0 of a particular module but application B needs version 2.0, then the requirements are in conflict and installing either version 1.0 or 2.0 will leave one application unable to run.

The solution for this problem is to create a virtual environment, a self-contained directory tree that contains a Python installation for a particular version of Python, plus a number of additional packages.

Different applications can then use different virtual environments. To resolve the earlier example of conflicting requirements, application A can have its own virtual environment with version 1.0 installed while application B has another virtual environment with version 2.0. If application B requires a library be upgraded to version 3.0, this will not affect application A's environment.

12.2. Creating Virtual Environments¶

The module used to create and manage virtual environments is called

```
venv
```

.

```
venv
```

will install the Python version from which the command was run (as reported by the

```
--version
```

option). For instance, executing the command with

```
python3.12
```

will install version 3.12.

To create a virtual environment, decide upon a directory where you want to place it, and run the

```
venv
```

module as a script with the directory path:

Copy

```
python -m venv tutorial-env
```

This will create the

```
tutorial-env
```

directory if it doesn't exist, and also create directories inside it containing a copy of the Python interpreter and various supporting files.

A common directory location for a virtual environment is

```
.venv
```

. This name keeps the directory typically hidden in your shell and thus out of the way while giving it a name that explains why the directory exists. It also prevents clashing with

```
.env
```

environment variable definition files that some tooling supports.

Once you've created a virtual environment, you may activate it.

On Windows, run:

Copy

```
tutorial-env\Scripts\activate
```

On Unix or MacOS, run:

Copy

```
source tutorial-env/bin/activate
```

(This script is written for the bash shell. If you use the **csh** or **fish** shells, there are alternate

```
activate.csh
```

and

```
activate.fish
```

scripts you should use instead.)

Activating the virtual environment will change your shell's prompt to show what virtual environment you're using, and modify the environment so that running

```
python
```

will get you that particular version and installation of Python. For example:

Copy

```
$ source ~/envs/tutorial-env/bin/activate
(tutorial-env) $ python
Python 3.5.1 (default, May  6 2016, 10:59:36)
...
>>> import sys
>>> sys.path
['', '/usr/local/lib/python35.zip', ...,
 '~/.envs/tutorial-env/lib/python3.5/site-packages']
>>>
```

To deactivate a virtual environment, type:

Copy

```
deactivate
```

into the terminal.

12.3. Managing Packages with `pip`

You can install, upgrade, and remove packages using a program called **pip**. By default

```
pip
```

will install packages from the [Python Package Index](#). You can browse the Python Package Index by going to it in your web browser.

```
pip
```

has a number of subcommands: “install”, “uninstall”, “freeze”, etc. (Consult the Installing Python Modules guide for complete documentation for

```
pip
```

.)

You can install the latest version of a package by specifying a package’s name:

Copy

```
(tutorial-env) $ python -m pip install novas
Collecting novas
  Downloading novas-3.1.1.3.tar.gz (136kB)
Installing collected packages: novas
  Running setup.py install for novas
Successfully installed novas-3.1.1.3
```

You can also install a specific version of a package by giving the package name followed by

```
==
```

and the version number:

Copy

```
(tutorial-env) $ python -m pip install requests==2.6.0
Collecting requests==2.6.0
  Using cached requests-2.6.0-py2.py3-none-any.whl
Installing collected packages: requests
Successfully installed requests-2.6.0
```

If you re-run this command,

```
pip
```

will notice that the requested version is already installed and do nothing. You can supply a different version number to get that version, or you can run

```
python
-m pip install --upgrade
```

to upgrade the package to the latest version:

Copy

```
(tutorial-env) $ python -m pip install --upgrade requests
Collecting requests
Installing collected packages: requests
  Found existing installation: requests 2.6.0
  Uninstalling requests-2.6.0:
    Successfully uninstalled requests-2.6.0
Successfully installed requests-2.7.0
```

```
python -m pip uninstall
```

followed by one or more package names will remove the packages from the virtual environment.

```
python -m pip show
```

will display information about a particular package:

Copy

```
(tutorial-env) $ python -m pip show requests
---
Metadata-Version: 2.0
Name: requests
Version: 2.7.0
Summary: Python HTTP for Humans.
Home-page: http://python-requests.org
Author: Kenneth Reitz
Author-email: me@kennethreitz.com
License: Apache 2.0
Location: /Users/akuchling/envs/tutorial-env/lib/python3.4/site-packages
Requires:
```

```
python -m pip list
```

will display all of the packages installed in the virtual environment:

Copy

```
(tutorial-env) $ python -m pip list
novas (3.1.1.3)
numpy (1.9.2)
pip (7.0.3)
requests (2.7.0)
setuptools (16.0)
```

```
python -m pip freeze
```

will produce a similar list of the installed packages, but the output uses the format that

```
python -m pip install
```

expects. A common convention is to put this list in a

```
requirements.txt
```

file:

Copy


```
(tutorial-env) $ python -m pip freeze > requirements.txt
(tutorial-env) $ cat requirements.txt
novas==3.1.1.3
numpy==1.9.2
requests==2.7.0
```

The

```
requirements.txt
```

can then be committed to version control and shipped as part of an application. Users can then install all the necessary packages with

```
install -r
```

:

Copy

```
(tutorial-env) $ python -m pip install -r requirements.txt
Collecting novas==3.1.1.3 (from -r requirements.txt (line 1))
...
Collecting numpy==1.9.2 (from -r requirements.txt (line 2))
...
Collecting requests==2.7.0 (from -r requirements.txt (line 3))
...
Installing collected packages: novas, numpy, requests
  Running setup.py install for novas
Successfully installed novas-3.1.1.3 numpy-1.9.2 requests-2.7.0
```

pip

has many more options. Consult the [Installing Python Modules](#) guide for complete documentation for

pip

. When you've written a package and want to make it available on the Python Package Index, consult the [Python packaging user guide](#).

13. What Now?¶

Reading this tutorial has probably reinforced your interest in using Python — you should be eager to apply Python to solving your real-world problems. Where should you go to learn more?

This tutorial is part of Python's documentation set. Some other documents in the set are:

- The Python Standard Library:

You should browse through this manual, which gives complete (though terse) reference material about types, functions, and the modules in the standard library. The standard Python distribution includes a *lot* of additional code. There are modules to read Unix mailboxes, retrieve documents via HTTP, generate random numbers, parse command-line options, compress data, and many other tasks. Skimming through the Library Reference will give you an idea of what's available.

- Installing Python Modules explains how to install additional modules written by other Python users.
- The Python Language Reference: A detailed explanation of Python's syntax and semantics. It's heavy reading, but is useful as a complete guide to the language itself.

More Python resources:

- <https://www.python.org>: The major Python website. It contains code, documentation, and pointers to Python-related pages around the web.
- <https://docs.python.org>: Fast access to Python's documentation.
- <https://pypi.org>: The Python Package Index, previously also nicknamed the Cheese Shop [1], is an index of user-created Python modules that are available for download. Once you begin releasing code, you can register it here so that others can find it.
- <https://code.activestate.com/recipes/langs/python/>: The Python Cookbook is a sizable collection of code examples, larger modules, and useful scripts. Particularly notable contributions are collected in a book also titled Python Cookbook (O'Reilly & Associates, ISBN 0-596-00797-3.)
- <https://pyvideo.org> collects links to Python-related videos from conferences and user-group meetings.
- <https://scipy.org>: The Scientific Python project includes modules for fast array computations and manipulations plus a host of packages for such things as linear algebra, Fourier transforms, non-linear solvers, random number distributions, statistical analysis and the like.

For Python-related questions and problem reports, you can post to the newsgroup *comp.lang.python*, or send them to the mailing list at python-list@python.org. The newsgroup and mailing list are gatewayed, so messages posted to one will automatically be forwarded to the other. There are hundreds of postings a day, asking (and answering) questions, suggesting new features, and announcing new modules. Mailing list archives are available at <https://mail.python.org/pipermail/>.

Before posting, be sure to check the list of Frequently Asked Questions (also called the FAQ). The FAQ answers many of the questions that come up again and again, and may already contain the solution for your problem.

Footnotes

[1]

“Cheese Shop” is a Monty Python sketch: a customer enters a cheese shop, but whatever cheese he asks for, the clerk says it’s missing.

14. Interactive Input Editing and History Substitution¶

Some versions of the Python interpreter support editing of the current input line and history substitution, similar to facilities found in the Korn shell and the GNU Bash shell. This is implemented using the [GNU Readline](#) library, which supports various styles of editing. This library has its own documentation which we won't duplicate here.

14.1. Tab Completion and History Editing¶

Completion of variable and module names is automatically enabled at interpreter startup so that the Tab key invokes the completion function; it looks at Python statement names, the current local variables, and the available module names. For dotted expressions such as

```
string.a
```

, it will evaluate the expression up to the final

```
'.'
```

and then suggest completions from the attributes of the resulting object. Note that this may execute application-defined code if an object with a

```
__getattr__()
```

method is part of the expression. The default configuration also saves your history into a file named

```
.python_history
```

in your user directory. The history will be available again during the next interactive interpreter session.

14.2. Alternatives to the Interactive Interpreter¶

This facility is an enormous step forward compared to earlier versions of the interpreter; however, some wishes are left: It would be nice if the proper indentation were suggested on continuation lines (the parser knows if an

```
INDENT
```

token is required next). The completion mechanism might use the interpreter's symbol table. A command to check (or even suggest) matching parentheses, quotes, etc., would also be useful.

One alternative enhanced interactive interpreter that has been around for quite some time is [IPython](#), which features tab completion, object exploration and advanced history management. It

can also be thoroughly customized and embedded into other applications. Another similar enhanced interactive environment is [bpython](#).

15. Floating-Point Arithmetic: Issues and Limitations¶

Floating-point numbers are represented in computer hardware as base 2 (binary) fractions. For example, the **decimal** fraction

0.625

has value $6/10 + 2/100 + 5/1000$, and in the same way the **binary** fraction

0.101

has value $1/2 + 0/4 + 1/8$. These two fractions have identical values, the only real difference being that the first is written in base 10 fractional notation, and the second in base 2.

Unfortunately, most decimal fractions cannot be represented exactly as binary fractions. A consequence is that, in general, the decimal floating-point numbers you enter are only approximated by the binary floating-point numbers actually stored in the machine.

The problem is easier to understand at first in base 10. Consider the fraction $1/3$. You can approximate that as a base 10 fraction:

Copy

0.3

or, better,

Copy

0.33

or, better,

Copy

0.333

and so on. No matter how many digits you're willing to write down, the result will never be exactly $1/3$, but will be an increasingly better approximation of $1/3$.

In the same way, no matter how many base 2 digits you're willing to use, the decimal value 0.1 cannot be represented exactly as a base 2 fraction. In base 2, $1/10$ is the infinitely repeating fraction

Copy

```
0.00011001100110011001100110011001100110011001100110011...
```

Stop at any finite number of bits, and you get an approximation. On most machines today, floats are approximated using a binary fraction with the numerator using the first 53 bits starting with the most significant bit and with the denominator as a power of two. In the case of $1/10$, the binary fraction is

```
3602879701896397 / 2 ** 55
```

which is close to but not exactly equal to the true value of $1/10$.

Many users are not aware of the approximation because of the way values are displayed. Python only prints a decimal approximation to the true decimal value of the binary approximation stored by the machine. On most machines, if Python were to print the true decimal value of the binary approximation stored for 0.1 , it would have to display:

Copy

```
>>> 0.1
0.1000000000000000055511151231257827021181583404541015625
```

That is more digits than most people find useful, so Python keeps the number of digits manageable by displaying a rounded value instead:

Copy

```
>>> 1 / 10
0.1
```

Just remember, even though the printed result looks like the exact value of $1/10$, the actual stored value is the nearest representable binary fraction.

Interestingly, there are many different decimal numbers that share the same nearest approximate binary fraction. For example, the numbers

```
0.1
```

and

```
0.10000000000000001
```

and

```
0.1000000000000000055511151231257827021181583404541015625
```

are all approximated by

```
3602879701896397 / 2 ** 55
```

. Since all of these decimal values share the same approximation, any one of them could be displayed while still preserving the invariant

```
eval(repr(x)) == x
```

.

Historically, the Python prompt and built-in

```
repr()
```

function would choose the one with 17 significant digits,

```
0.10000000000000001
```

. Starting with Python 3.1, Python (on most systems) is now able to choose the shortest of these and simply display

```
0.1
```

.

Note that this is in the very nature of binary floating point: this is not a bug in Python, and it is not a bug in your code either. You'll see the same kind of thing in all languages that support your hardware's floating-point arithmetic (although some languages may not *display* the difference by default, or in all output modes).

For more pleasant output, you may wish to use string formatting to produce a limited number of significant digits:

Copy

```
>>> format(math.pi, '.12g') # give 12 significant digits
'3.14159265359'

>>> format(math.pi, '.2f')   # give 2 digits after the point
'3.14'

>>> repr(math.pi)
'3.141592653589793'
```

It's important to realize that this is, in a real sense, an illusion: you're simply rounding the *display* of the true machine value.

One illusion may beget another. For example, since 0.1 is not exactly 1/10, summing three values of 0.1 may not yield exactly 0.3, either:

Copy


```
>>> 0.1 + 0.1 + 0.1 == 0.3
False
```

Also, since the 0.1 cannot get any closer to the exact value of 1/10 and 0.3 cannot get any closer to the exact value of 3/10, then pre-rounding with

```
round()
```

function cannot help:

Copy

```
>>> round(0.1, 1) + round(0.1, 1) + round(0.1, 1) == round(0.3, 1)
False
```

Though the numbers cannot be made closer to their intended exact values, the

```
math.isclose()
```

function can be useful for comparing inexact values:

Copy

```
>>> math.isclose(0.1 + 0.1 + 0.1, 0.3)
True
```

Alternatively, the

```
round()
```

function can be used to compare rough approximations:

Copy

```
>>> round(math.pi, ndigits=2) == round(22 / 7, ndigits=2)
True
```

Binary floating-point arithmetic holds many surprises like this. The problem with “0.1” is explained in precise detail below, in the “Representation Error” section. See [Examples of Floating Point Problems](#) for a pleasant summary of how binary floating point works and the kinds of problems commonly encountered in practice. Also see [The Perils of Floating Point](#) for a more complete account of other common surprises.

As that says near the end, “there are no easy answers.” Still, don’t be unduly wary of floating point! The errors in Python float operations are inherited from the floating-point hardware, and on most machines are on the order of no more than 1 part in 2^{53} per operation. That’s more than

adequate for most tasks, but you do need to keep in mind that it's not decimal arithmetic and that every float operation can suffer a new rounding error.

While pathological cases do exist, for most casual use of floating-point arithmetic you'll see the result you expect in the end if you simply round the display of your final results to the number of decimal digits you expect.

```
str()
```

usually suffices, and for finer control see the

```
str.format()
```

method's format specifiers in Format String Syntax.

For use cases which require exact decimal representation, try using the

```
decimal
```

module which implements decimal arithmetic suitable for accounting applications and high-precision applications.

Another form of exact arithmetic is supported by the

```
fractions
```

module which implements arithmetic based on rational numbers (so the numbers like 1/3 can be represented exactly).

If you are a heavy user of floating-point operations you should take a look at the NumPy package and many other packages for mathematical and statistical operations supplied by the SciPy project. See <<https://scipy.org>>.

Python provides tools that may help on those rare occasions when you really *do* want to know the exact value of a float. The

```
float.as_integer_ratio()
```

method expresses the value of a float as a fraction:

Copy

```
>>> x = 3.14159
>>> x.as_integer_ratio()
(3537115888337719, 1125899906842624)
```

Since the ratio is exact, it can be used to losslessly recreate the original value:

Copy

```
>>> x == 3537115888337719 / 1125899906842624
True
```

The

```
float.hex()
```

method expresses a float in hexadecimal (base 16), again giving the exact value stored by your computer:

Copy

```
>>> x.hex()
'0x1.921f9f01b866ep+1'
```

This precise hexadecimal representation can be used to reconstruct the float value exactly:

Copy

```
>>> x == float.fromhex('0x1.921f9f01b866ep+1')
True
```

Since the representation is exact, it is useful for reliably porting values across different versions of Python (platform independence) and exchanging data with other languages that support the same format (such as Java and C99).

Another helpful tool is the

```
sum()
```

function which helps mitigate loss-of-precision during summation. It uses extended precision for intermediate rounding steps as values are added onto a running total. That can make a difference in overall accuracy so that the errors do not accumulate to the point where they affect the final total:

Copy

```
>>> 0.1 + 0.1 + 0.1 + 0.1 + 0.1 + 0.1 + 0.1 + 0.1 + 0.1 + 0.1 == 1.0
False
>>> sum([0.1] * 10) == 1.0
True
```

The

```
math.fsum()
```

goes further and tracks all of the “lost digits” as values are added onto a running total so that the

result has only a single rounding. This is slower than

```
sum()
```

but will be more accurate in uncommon cases where large magnitude inputs mostly cancel each other out leaving a final sum near zero:

Copy

```
>>> arr = [-0.10430216751806065, -266310978.67179024, 143401161448607.16,
...        -143401161400469.7, 266262841.31058735, -0.003244936839808227]
>>> float(sum(map(Fraction, arr)))    # Exact summation with single rounding
8.042173697819788e-13
>>> math.fsum(arr)                   # Single rounding
8.042173697819788e-13
>>> sum(arr)                         # Multiple roundings in extended precision
8.042178034628478e-13
>>> total = 0.0
>>> for x in arr:
...     total += x                   # Multiple roundings in standard precision
...
>>> total                             # Straight addition has no correct digits!
-0.0051575902860057365
```

15.1. Representation Error¶

This section explains the “0.1” example in detail, and shows how you can perform an exact analysis of cases like this yourself. Basic familiarity with binary floating-point representation is assumed.

Representation error refers to the fact that some (most, actually) decimal fractions cannot be represented exactly as binary (base 2) fractions. This is the chief reason why Python (or Perl, C, C++, Java, Fortran, and many others) often won’t display the exact decimal number you expect.

Why is that? $1/10$ is not exactly representable as a binary fraction. Since at least 2000, almost all machines use IEEE 754 binary floating-point arithmetic, and almost all platforms map Python floats to IEEE 754 binary64 “double precision” values. IEEE 754 binary64 values contain 53 bits of precision, so on input the computer strives to convert 0.1 to the closest fraction it can of the form $J/2^{**N}$ where J is an integer containing exactly 53 bits. Rewriting

Copy

```
1 / 10 ~= J / (2**N)
```

as

Copy

```
J ~= 2**N / 10
```

and recalling that J has exactly 53 bits (is

```
>= 2**52
```

but

```
< 2**53
```

), the best value for N is 56:

Copy

```
>>> 2**52 <= 2**56 // 10 < 2**53
True
```

That is, 56 is the only value for N that leaves J with exactly 53 bits. The best possible value for J is then that quotient rounded:

Copy

```
>>> q, r = divmod(2**56, 10)
>>> r
6
```

Since the remainder is more than half of 10, the best approximation is obtained by rounding up:

Copy

```
>>> q+1
7205759403792794
```

Therefore the best possible approximation to $1/10$ in IEEE 754 double precision is:

Copy

```
7205759403792794 / 2 ** 56
```

Dividing both the numerator and denominator by two reduces the fraction to:

Copy

```
3602879701896397 / 2 ** 55
```

Note that since we rounded up, this is actually a little bit larger than $1/10$; if we had not rounded up, the quotient would have been a little bit smaller than $1/10$. But in no case can it be *exactly* $1/10$!

So the computer never “sees” $1/10$: what it sees is the exact fraction given above, the best IEEE 754 double approximation it can get:

Copy

```
>>> 0.1 * 2 ** 55
3602879701896397.0
```

If we multiply that fraction by 10^{55} , we can see the value out to 55 decimal digits:

Copy

```
>>> 3602879701896397 * 10 ** 55 // 2 ** 55
1000000000000000055511151231257827021181583404541015625
```

meaning that the exact number stored in the computer is equal to the decimal value $0.1000000000000000055511151231257827021181583404541015625$. Instead of displaying the full decimal value, many languages (including older versions of Python), round the result to 17 significant digits:

Copy

```
>>> format(0.1, '.17f')
'0.10000000000000001'
```

The

fractions

and

decimal

modules make these calculations easy:

Copy

```
>>> from decimal import Decimal
>>> from fractions import Fraction

>>> Fraction.from_float(0.1)
Fraction(3602879701896397, 36028797018963968)

>>> (0.1).as_integer_ratio()
(3602879701896397, 36028797018963968)

>>> Decimal.from_float(0.1)
Decimal('0.1000000000000000055511151231257827021181583404541015625')

>>> format(Decimal.from_float(0.1), '.17')
'0.10000000000000001'
```

16. Appendix¶

16.1. Interactive Mode¶

There are two variants of the interactive REPL. The classic basic interpreter is supported on all platforms with minimal line control capabilities.

Since Python 3.13, a new interactive shell is used by default. This one supports color, multiline editing, history browsing, and paste mode. To disable color, see [Controlling color](#) for details. Function keys provide some additional functionality. F1 enters the interactive help browser

```
pydoc
```

. F2 allows for browsing command-line history with neither output nor the `>>>` and `...` prompts. F3 enters “paste mode”, which makes pasting larger blocks of code easier. Press F3 to return to the regular prompt.

When using the new interactive shell, exit the shell by typing `exit` or `quit`. Adding call parentheses after those commands is not required.

If the new interactive shell is not desired, it can be disabled via the

```
PYTHON_BASIC_REPL
```

environment variable.

16.1.1. Error Handling¶

When an error occurs, the interpreter prints an error message and a stack trace. In interactive mode, it then returns to the primary prompt; when input came from a file, it exits with a nonzero exit status after printing the stack trace. (Exceptions handled by an

```
except
```

clause in a

```
try
```

statement are not errors in this context.) Some errors are unconditionally fatal and cause an exit with a nonzero exit status; this applies to internal inconsistencies and some cases of running out of memory. All error messages are written to the standard error stream; normal output from executed commands is written to standard output.

Typing the interrupt character (usually `Control-C` or `Delete`) to the primary or secondary prompt cancels the input and returns to the primary prompt. [\[1\]](#) Typing an interrupt while a command is

executing raises the

```
KeyboardInterrupt
```

exception, which may be handled by a

```
try
```

statement.

16.1.2. Executable Python Scripts¶

On BSD'ish Unix systems, Python scripts can be made directly executable, like shell scripts, by putting the line

Copy

```
#!/usr/bin/env python3
```

(assuming that the interpreter is on the user's

```
PATH
```

) at the beginning of the script and giving the file an executable mode. The

```
#!
```

must be the first two characters of the file. On some platforms, this first line must end with a Unix-style line ending (

```
'\n'
```

), not a Windows (

```
'\r\n'
```

) line ending. Note that the hash, or pound, character,

```
'#'
```

, is used to start a comment in Python.

The script can be given an executable mode, or permission, using the **chmod** command.

Copy

```
$ chmod +x myscript.py
```

On Windows systems, there is no notion of an “executable mode”. The Python installer automatically associates

```
.py
```

files with

```
python.exe
```

so that a double-click on a Python file will run it as a script. The extension can also be

```
.pyw
```

, in that case, the console window that normally appears is suppressed.

16.1.3. The Interactive Startup File ¶

When you use Python interactively, it is frequently handy to have some standard commands executed every time the interpreter is started. You can do this by setting an environment variable named

```
PYTHONSTARTUP
```

to the name of a file containing your start-up commands. This is similar to the

```
.profile
```

feature of the Unix shells.

This file is only read in interactive sessions, not when Python reads commands from a script, and not when

```
/dev/tty
```

is given as the explicit source of commands (which otherwise behaves like an interactive session). It is executed in the same namespace where interactive commands are executed, so that objects that it defines or imports can be used without qualification in the interactive session. You can also change the prompts

```
sys.ps1
```

and

```
sys.ps2
```

in this file.

If you want to read an additional start-up file from the current directory, you can program this in the global start-up file using code like

```
if
os.path.isfile('.pythonrc.py'): exec(open('.pythonrc.py').read())
```

. If you want to use the startup file in a script, you must do this explicitly in the script:

Copy

```
import os
filename = os.environ.get('PYTHONSTARTUP')
if filename and os.path.isfile(filename):
    with open(filename) as fobj:
        startup_file = fobj.read()
        exec(startup_file)
```

16.1.4. The Customization Modules^[1]

Python provides two hooks to let you customize it: `sitecustomize` and `usercustomize`. To see how it works, you need first to find the location of your user site-packages directory. Start Python and run this code:

Copy

```
>>> import site
>>> site.getusersitepackages()
'/home/user/.local/lib/python3.x/site-packages'
```

Now you can create a file named

`usercustomize.py`

in that directory and put anything you want in it. It will affect every invocation of Python, unless it is started with the

`-S`

option to disable the automatic import.

`sitecustomize` works in the same way, but is typically created by an administrator of the computer in the global site-packages directory, and is imported before `usercustomize`. See the documentation of the

`site`

module for more details.

Footnotes

[\[1\]](#)

A problem with the GNU Readline package may prevent this.